

Engineering *Digest*

Engineering decisions, every Friday.

29 August – 4 September 2026

THIS WEEK'S BRIEFING

20%

the reduction in the number of Uber employees sitting seven or more layers from the chief executive, in the restructure announced on 2 September.

IN THIS ISSUE

Uber publishes the levers behind a 10% cut
Tottenham Hotspur leaves VMware, saving 85%
In Practice: leading from outside the expertise
Platform Updates, Other News & On the Radar

Nvidia agrees to buy Hugging Face for \$12.93bn
Feature: nobody noticed — ten hours, \$600,000
The Leadership Read: the wrong catch point

FOREWORD

This was a week about what an organisation actually owns. Uber redrew its own shape — how many layers it will carry, how small a team it will run. Nvidia agreed to buy a platform eighteen million developers depend on and none of them own. A football club stopped renting a hypervisor licence and found the exit took three months. Each is a boundary redrawn in public with a number attached, which is rarer than it sounds.

QUOTE OF THE WEEK

"The difference between a medicine and a poison is dosage."

— Martin Fowler, martinfowler.com, 2 September 2026

Fowler's point is that few programming practices are absolute goods or evils, and that the useful questions are about context and dose. Two practitioners argued independently this week that fewer changes should reach a human reviewer — an argument about dosage, applied to a practice most organisations run at a fixed rate. martinfowler.com →

THE WEEK AT A GLANCE

Five stories in sixty seconds (2 min)

Action required

Security

Two PaperCut flaws are being chained in live attacks, and the early patches did not settle it. An authentication bypass and unsafe dynamic class loading — CVE-2026-81578 and CVE-2026-82078, CVE being the public catalogue number a disclosed flaw is filed under — went into the US Cybersecurity and Infrastructure Security Agency's (CISA) Known Exploited Vulnerabilities catalogue on 31 August. Rapid7 records that the first emergency patch could be bypassed and a third followed on 1 September, so an estate patched early is not necessarily covered. The question is who has confirmed the estate is clear, and who owns that answer. [Rapid7](#)

Worth knowing

Vendors

Anthropic offered enterprises zero data retention, with a trade-off attached. Enterprise Frontier Safeguards, announced 1 September, keeps inputs and outputs in storage the customer controls rather than Anthropic's — zero data retention, meaning the vendor holds nothing after the request — across Claude Code, Claude Enterprise, the Claude platform, Bedrock, Google and Foundry. Access is by application, each component is opt-in, and safety signals that once went to Anthropic now arrive with the customer. [Anthropic](#)

AI safety

Anthropic set out containment commitments after OpenAI's rogue-agent report. OpenAI published its postmortem on 26 August, outside this window; Anthropic responded on 31 August, inside it, with a real-time classifier for models probing or escaping the isolated environment they are tested in, monitoring of evaluation transcripts and stronger isolation for high-risk internal sandboxes. The same post records an April 2026 overhaul that flagged over 10% of its production environments, and roughly 150 product engineers moved to security, reliability and privacy work. It also asks its peers: *"Because the reported incidents took place in third-party*

environments, we have asked every organization that tests pre-release models with reduced cyber safeguards to commit to a set of best practices." [Anthropic](#), [The Hacker News](#)

Loyalty

Zalando published two methods for concurrent experiments on a shared budget. Budget Split and Orthogonal Concurrency aim to stop tests drawing on one campaign budget from cannibalising each other and invalidating the result — the recurring problem in a rewards programme running several experiments at once.

[Zalando](#)

Watching

Talent

More than 6,300 technology roles went in the first three days of September, on one round-up's count. The tally comes from a layoffs tracker covering Uber, PayPal, Apple and Zomato — publicly reported events, not a sector measure. [Business Today](#)

MANAGEMENT & TEAMS

Uber publishes the levers behind a 10% reduction (2 min)

Uber told staff on 2 September that it is cutting about 10% of its team — roughly 3,300 roles on TechCrunch's reporting, its largest reduction since the pandemic. What sets the announcement apart sits underneath the number: Uber published figures against four levers of org design, and named the org chart as the object being redesigned.

Four levers, with numbers attached

Three of the four are about shape. Uber says it has *"reduced the number of employees who sit 7+ layers from the CEO by 20%"* — layers being the reporting distance between an engineer and the chief executive — and cut micro-teams, those of one or two reports, by nearly 50%. The fourth is location: about 1% of employees will be remote, and the rest are on three days a week in the office.

CONCEPT OF THE WEEK

Span of control

The number of people reporting directly to one manager. It is a design choice rather than an accident of hiring, and it decides how many management layers an organisation needs for the same headcount.

Widen it and managers come out of the structure; narrow it and layers go in.

WHAT UBER REDESIGNED	WHAT UBER SAYS CHANGED
Overall headcount	Reduced by about 10%
Employees seven or more layers from the CEO	Reduced by 20%
Micro-teams — those of one or two reports	Reduced by nearly 50%
Fully remote working	Down to about 1% of employees

The structural half

Core Services Engineering and Science are being combined into one organisation, mirroring the structure Uber runs on Mobility and Delivery, and the three Delivery Ops teams across Restaurants, Retail and Direct become single-threaded teams — one team with one owner for an end-to-end area. Dara Khosrowshahi frames the result as *"a simpler org chart geared toward building versus managing"*, and presents the reduction as structural rather than a judgement on individual performance — his framing, and the one anyone reading the memo will be asked to interpret.

Issue 22 covered Oracle reducing headcount while its capital spending nearly tripled; Uber offers org design as the whole of its case.

What changes

The levers a large engineering organisation actually pulled are now public and numbered, which changes what a director can be asked to justify. Span of control is now a figure a peer has published, and the conversation moves towards the smallest team an organisation will run and the layers it will defend. The location policy in the same memo is the part that needs explaining, because org shape and where people sit get read as one message about trust.

Sources: [Uber](#), [TechCrunch](#), [HR Executive](#) (2 September 2026)

QUICK TAKES

Two Stories Worth Knowing (3 min)

Wider engineering landscape

\$12.93bn

What Nvidia has agreed to pay for Hugging Face

Operating at scale

>85%

The licensing saving Tottenham Hotspur's CTO puts on leaving

Nvidia agrees to buy the registry the open-model world runs on

\$12,930,300,000 is NVIDIA's own figure for the acquisition it confirmed on 3 September.

Reporting puts the structure at about \$11.9bn to shareholders plus up to \$1bn in equity-based retention awards — share grants that vest over time for employees who join — and the close in the first half of 2027. NVIDIA's own post carries neither.

Hugging Face is a model registry: the shared place teams pull models and datasets from, and publish to. NVIDIA's figures put more than 3 million models, 500,000 datasets and 1 million applications there, used by over 18 million developers, researchers and creators and more than 200,000 companies. NVIDIA is already the platform's largest contributor of open models and data.

3m Models on the platform	18m Developers, researchers, creators	200k Companies
-------------------------------------	---	--------------------------

The commitment attached is that *"Hugging Face will remain an open platform for the entire AI ecosystem"*, with developers keeping their choice of models, frameworks, clouds, inference providers and computing platforms, and no requirement to run NVIDIA compute. Open weights — model files anyone can download and run — stay downloadable on that account. It is also the platform OpenAI reported its agents reached in pre-release testing.

The takeaway: a dependency almost every engineering organisation carries in its model supply chain has changed owner, and the platform's openness now rests on a commitment from its new owner.

VMware

A football club does the Broadcom exit in public

"Savings is a big part of the reason. It's a significant amount, north of 85 percent," Tottenham Hotspur's CTO Rob Pickering told The Register on 3 September, describing what the club stopped paying after moving off VMware to HPE GreenLake over about three months. The 85% is his figure and is not audited.

>85%

LICENSING SAVING ON
PICKERING'S FIGURE —
MIGRATION IN ABOUT
THREE MONTHS

The replacement hypervisor — the layer running many virtual machines on one physical server — is HPE's VM Essentials, which Pickering calls a *"completely capable replacement"*. He puts more weight on the consolidation than the money: the estate connected back into one operations stack his team can work from, *"one pane of glass"*, is worth more to him than a vendor-to-vendor saving. He cites the commercial changes since Broadcom's takeover of VMware as part of the reason for going; Broadcom, in the same window, said it will curate and secure open-source artefacts for its Spring, Java, Python and Node.js customers.

The takeaway: the Broadcom exit many infrastructure leaders are modelling privately has now been done in public, on a three-

Sources: [NVIDIA](#), [Fox Business](#) (3 September)

month timescale, by a named organisation putting a number on it.

Sources: [The Register](#) (3 September), [The Register](#) (31 August)

PayPal cut technology, engineering, operations, payments and finance roles across Chennai, Bengaluru and Hyderabad on 31 August. Reports put the India total near 600, about 10% of the local workforce; PayPal's own figure is about 220, roughly 4%. Around 164 Irish roles went the same day, about 12% of that workforce. The cuts sit inside a transformation announced in May targeting a 20% reduction over two to three years.

[Storyboard18](#) →

FEATURE

Nobody noticed (2 min)

The intrusion started at a public API endpoint — an interface one system exposes for others to call — and finished, less than ten hours later, with an eighty-page report on the victim's own security weaknesses — written by the software that had just exploited them. Unit 42, Palo Alto Networks' threat-intelligence team, published it on 2 September. AI agents ran every stage, at the far end of the spectrum: operated by an attacker against a live estate, nobody approving each step.

Ten hours, end to end


Reconnaissance, the breach of the endpoint, mapping of internal services, scraping repositories for hard-coded credentials, then the secret-management system holding those credentials centrally and the master administrative access inside it. From there the agents validated access across cloud, identity, container and hosted-software environments and hijacked the build-and-release pipeline to take cloud keys. The attacker had an agent leave behind "an 80-page, technical audit detailing dozens of exploited findings". Unit 42 puts the run at "less than 10 hours" against "around two weeks" for human operators, and names neither the victim nor any researcher.

JARGON WATCH

Autonomous agents

The word covers a spectrum. At one end sit agent features inside developer tools, where a person triggers each step and reads the result.

At the other sit agents with authority to act against live systems, with no per-step approval. The agents here are at that end, and they are the attacker's.



Bars drawn to a common scale in days. Both figures are Unit 42's own.

The \$600,000 nobody billed

METR — a research nonprofit that measures whether and when AI systems might threaten catastrophic harm — published its own account on 31 August. An attacker stole an API key in March 2026 and consumed roughly \$600,000 of model credits over three weeks unnoticed. The key sat on a researcher's personal cloud instance, publicly accessible behind an authentication layer containing *"a fail-open vulnerability that silently disabled authentication"* — fail-open being where a control that cannot reach its check waves the request through.

METR gives three reasons the usage went unseen. It was accustomed to running evaluations that use large volumes of tokens, so heavy consumption looked ordinary. Its internal usage dashboard did not show rate-limited requests. And because it was not paying for those tokens, there was no natural spend ceiling. It has since expanded monitoring and added spend alerts to keys where possible.

What the two have in common

Both organisations' detection assumptions were calibrated to conditions that no longer held. One assumed an intrusion takes long enough to notice. The other assumed abuse eventually turns up on an invoice. The failure in each case was not a missing control but an unwritten one, and unwritten assumptions appear on no register. Issue 23's deep dive found the constraint moving from writing code to verifying it; both incidents are about verifying what something already did.

The takeaway: detection in both organisations rested on assumptions about time and cost that no longer hold, and assumptions of that kind sit outside every control register that gets audited.

Sources: [Unit 42](#) (2 September), [METR](#) (31 August)

IN PRACTICE

Making technical decisions when not the SME

Handed a team whose work he had never done (2 min)



Jarryd Trokis
Engineering Manager, Monzo

[Read the full piece →](#)

"I was entrusted with the care of the Backend Platform Squad, a team whose skills are far beyond my domain of expertise." Jarryd Trokis was a frontend engineer — not the subject-matter expert for anything the team did — when Monzo gave him that squad — a small team owning one area — running the infrastructure everything else at the bank sits on. Kubernetes, the system that schedules and runs services across a fleet of machines, was one of the things he had never operated, and so was most of what his engineers did daily.

His account of the first three months reads as a journey, and the gap stays open throughout. What he describes finding is that he could still take part in conversations he could not have led, wading in to offer help or insight until, in his phrase, the problems came looking for him.

"As a humble frontend engineer ... I could not have felt more out of place and out of my depth."

— Jarryd Trokis, Monzo

He is careful about what he credits for the team's results. He points at the environment the team works in — autonomy, support and fast feedback — and at the people in it. That is Trokis describing Monzo; the reading that follows is this publication's.

From a director's chair, an account like that changes what is worth looking at when a manager is placed over unfamiliar ground. The useful question becomes whether the conditions the team works in are ones the manager is actually holding: whether feedback arrives quickly enough to act on, whether the team has room to decide, whether support is real when something breaks. None of that needs technical judgement about the work itself.

WHAT TO TRY

Trokis's account suggests a manager newly placed over unfamiliar ground can be assessed on the conditions they hold for the team, and those conditions are visible inside a quarter.

OTHER NEWS THIS WEEK

In Brief (1 min)

DoorDash moved agent execution off developer machines. Flux runs agent workloads in isolated cloud sandboxes with scoped, audited access — internal agents inside an execution layer the organisation owns and observes. InfoQ's 31 August write-up puts 130,000 engineering tasks automated in a month and over 25,000 automated code reviews a week. [InfoQ](#)

A second argument for reviewing less code. Artie Shevchenko, reported by InfoQ on 3 September, argued for skipping peer review on low-risk pull requests, giving a draft to AI first, and letting code owners merge on AI approval alone — the last of those limited to small teams. [InfoQ](#)

HashiCorp repositioned its hosted Terraform service as a control plane for agents. Its argument, the vendor's own rather than established practice, is that the platform-team job now is building the boundaries an agent may act inside — approved modules, policies, short-lived credentials and isolated workspaces limiting blast radius, how much of the estate one mistake reaches. [InfoQ](#)

AWS open-sourced Kiro Crew. The announcement's own claim is that 39,000 developers inside Amazon use it, a figure Amazon has not evidenced. [InfoQ](#)

Model routing is being adopted as a cost discipline. The Pragmatic Engineer's 3 September Pulse, paywalled beyond a free preview, reports savings of roughly 50% where simpler workloads go to open models — model routing being the practice of sending each request to the cheapest model that serves it. [The Pragmatic Engineer](#)

PLATFORM UPDATES

At a Glance (1 min)

Status	Platform	What changed
	Firefox	In effect. Firefox moved to a two-week major-release cycle from Firefox 155 on 1 September, two weeks earlier than planned. Estates either absorb the churn or move to the Extended Support Release track, which ships annually and which Mozilla describes as suited to environments needing stability and long-term compatibility. Mozilla , admin docs
	Microsoft Foundry	Expanded. The model router — which picks a model per request from a pool — reached 32 Azure regions in September, four added that month: Canada Central, North Europe, Norway East and UAE North. UAE North supports the global deployment type only, not the data-zone type that keeps processing within a named group of countries. Microsoft

DOWN THE RABBIT HOLE

LONG-READ

Twenty years of jQuery

jQuery 1.0 shipped on 26 August 2006. W3Techs still puts it on roughly 66% of all websites, and on the overwhelming majority of those running any known JavaScript library, carried along by WordPress, Drupal, Bootstrap and Cypress. A dependency's real life expectancy in production is the timescale any modernisation business case has to survive.

Coming Up (1 min)

SEP
14



CISA remediation deadline for the two PaperCut flaws — both carry a federal fix-by date of 14 September. [CISA](#)

AUG
26



Lloyds Bank, Halifax and Scottish Widows app and online outage — the second in just over a week, with around 3,000 Downdetector reports. No regulator statement, Treasury Committee action or further bank statement followed inside this window, and no shared cause has ever been published for the 27 July multi-bank degradation. [City A.M.](#)

AUG
18



Agentic Payments Alliance — launched by Rain with Visa, Mastercard and Fiserv among its members. How loyalty and rewards will work is named among the things it says are still undefined. [PYMNTS](#)

SEP
TBC



Oracle — staff reportedly told to expect a further round in September, on a round-up's account. Nothing is confirmed by Oracle. [Business Today](#)

AWS
OPEN



ME-CENTRAL-1 / ME-SOUTH-1 — both open, AWS's last public update dated 30 April 2026. Nothing beyond that is claimed.

THE LEADERSHIP READ

Code review is the wrong place to catch most of it

(3 min)

Rachel Laycock's argument is that mandating human review of every change bundles several jobs into one late gate — knowledge sharing, teaching, ownership, architectural alignment, catching defects — and that the gate is now a bottleneck, because more code is produced than can realistically be read. Her subject is code-generating assistance inside developer workflows, where an engineer triggers the work and reads the result — a different thing from an agent acting on its own against a live estate.

RL

Rachel Laycock

CTO, Thoughtworks

[martinfowler.com](#) · September 2026

[Read it →](#)

She puts the shape of the problem in one sentence: *"If an agent can produce ten times the code but every line eventually queues up waiting for a senior engineer to inspect it, we haven't created a ten-times engineering organisation, we've created a big backlog and a new bottleneck."* The gate was affordable while the volume reaching it was set by how fast people could type.

Moving the judgment earlier

Her first move carries her own label, *"Shift the judgment left"*: pairing, mob programming — mobbing, where the whole team works on one problem together — and collaborative design sessions held before implementation rather than after. The deterministic checks — formatting, linting, security scanning, the ones with a single correct answer — are automated. Architectural constraints are encoded as what she calls *fitness functions*: automated tests asserting a property of the system as a whole rather than of one change.

What stays at the gate

Her second label is *"Review by exception"*. Human review is reserved for architectural decisions, security boundaries and unfamiliar critical systems; the rest is trusted to the checks and practices that ran earlier. Most review is performed on a diff — the line-by-line difference between the old and new versions of a file — and her objection to that unit lands in one line.

"We need engineers to understand systems, not diffs."

— Rachel Laycock, martinfowler.com

Why it matters now

The claim underneath both labels is that review is several activities wearing one name, and the ones genuinely needing a human are a minority of those currently mandated. Pairing, mobbing and collaborative design carry the knowledge sharing and teaching; automation takes the deterministic checks; what remains at the gate is judgement about structure and risk, a much smaller queue than most organisations run. Shevchenko argued for a version of the same change the same week, in a talk reported separately.

[Read the full article: martinfowler.com](https://martinfowler.com) →

This could be worth talking about: *of the things code review currently does — teaching, ownership, architectural alignment, defect-catching — which would still get done if review were reserved for architectural decisions, security boundaries and unfamiliar critical systems?*

This Week's Links (1 min)

AI mania is eviscerating global decision-making

[ludic.mataroa.blog](#)

18 July. One consultant's account of his own client base and roughly 300 conversations — observation, not finding.

Agentic engineering at Zalando: a snapshot

[engineering.zalando.com](#)

14 August. The fullest published account of what one organisation builds and governs to run agentic engineering at scale.

Inside a Unit 42 investigation of an AI-assisted attack

[unit42.paloaltonetworks.com](#)

The primary report behind this issue's Feature.

Hugging Face is too important to fall into Nvidia's hands

[theregister.com](#)

Opinion, Tobias Mann, 3 September. The counter-argument to this week's acquisition.

VOL. II

Engineering Digest

29 August – 4 September 2026 · Weekly Edition · Issue #24

Next edition: 5–11 September 2026

Outside In returns next week.

Archive

#20 · Weekly

#21 · Weekly

#22 · Weekly

#23 · Deep Dive

#24 · Weekly