Engineering *Digest*

February 18 – March 18, 2026 | Issue #0

MONTHLY RETROSPECTIVE

IN THIS ISSUE

When Drones Hit the Cloud: The AWS Middle East Crisis
The AI Delivery Bottleneck
The Leadership Read: Agentic Engineering in Practice
Platform Updates, Other News & On the Radar
The AI Coding Arms Race Enters a New Phase
Revolut Gets Its Full UK Banking Licence
Outside In: Monzo's Backup Bank

FOREWORD

On March 1, Iranian drones hit two AWS data centres in the UAE. Within hours, 109 cloud services went dark across the Middle East — not degraded, unavailable. Banks lost mobile platforms. Payment processors went offline. And the industry confronted something it had spent two decades abstracting away: the cloud has a physical layer, and that physical layer sits in the real world where geopolitics happens. For engineering leaders who built their resilience strategies around multi-AZ deployments within a single region, the assumptions just broke.

Meanwhile, two other forces reshaped the landscape. AI coding adoption hit saturation — 95% weekly usage — but CircleCI's analysis of 28 million workflows revealed the uncomfortable corollary: teams are writing 59% more code and shipping it 7% more slowly. The bottleneck has moved from code generation to everything that happens after. And in UK banking, Revolut secured full PRA authorisation, giving an engineering-led digital platform the same regulatory standing as Barclays. The competitive dynamics in financial services just shifted permanently.

This is the inaugural issue of Engineering Digest. No sponsored content, no hype cycles, no listicles. Just the things that matter, framed for the people who have to make decisions about them.

THE MONTH AT A GLANCE

Seven stories in sixty seconds

AWS **Iranian drone strikes hit AWS data centres in UAE and Bahrain** (March 1) — two of three AZs in ME-CENTRAL-1 went offline. 109+ services unavailable. 16+ days without full restoration. CSIS calls it the first major military strike on hyperscale cloud infrastructure.

AI **Microsoft Copilot Wave 3 integrates Claude** (March 9) — Anthropic's Cowork technology powers Agent 365. E7 Frontier Suite launching May 1 at $99/user. Claude now in mainline Copilot chat.

AI **Anthropic launched the $100M Claude Partner Network** (March 12) — certified training, Claude Certified Architect programme, co-marketing. Anthropic at $19B annualised revenue run rate.

Process **PMI and Agile Alliance released the Manifesto for Enterprise Agility** (March 3) — first joint manifesto. Extends agile from team delivery to enterprise governance, funding, and leadership.

.NET **March security patches across .NET 10/9/8.** Emergency .NET 10.0.5 out-of-band fix for macOS debugging (March 12). C# MCP SDK now available.

Node **Node.js 24.14.0 LTS released** (Feb 24). Node 20 EOL April 30. New release schedule announced: one major per year from Node 27.

Security **March Patch Tuesday: 82+ CVEs, 8 critical, 2 publicly disclosed zero-days.** Notable: Azure MCP Server SSRF (CVE-2026-26118, 8.8 CVSS) and Excel XSS that could enable Copilot Agent mode data exfiltration.

CLOUD & INFRASTRUCTURE

When Drones Hit the Cloud: The AWS Middle East Crisis

109+	3	2 of 3	16+
SERVICES AFFECTED	FACILITIES STRUCK	AZS OFFLINE	DAYS WITHOUT RESTORATION

On March 1, Iranian drone strikes hit AWS data centres in the UAE and Bahrain. Two of three availability zones in ME-CENTRAL-1 went offline. Not degraded — offline. 109+ AWS services became unavailable. The physical damage was severe: structural, power, fire damage, and water damage from suppression systems. AWS initially described the event as "objects" causing a "localized power issue." Thirty-six hours later, on March 2, they confirmed it was drone strikes.

The banking impact was immediate. ADCB mobile banking went offline. Emirates NBD was disrupted. Payment platforms Alaan and Hubpay went down. Careem — the region's dominant ride-hailing app — was disrupted. As of March 17, more than 16 days after the strikes, there is no full restoration date. AWS's official recommendation: migrate workloads to alternate regions.

The data residency paradox

Here's where it gets structurally interesting. UAE and Bahrain Central Bank mandates require financial data to stay in-region. The region's infrastructure has been physically destroyed. You can't comply with data residency requirements when the region doesn't exist in a functional sense. The Dubai Financial Services Authority issued emergency guidance allowing temporary cross-region processing — a regulatory improvisation that nobody had planned for. Every organisation that built its compliance posture around "data stays in ME-CENTRAL-1" is now operating in legal grey territory.

> *This is the first major military strike on hyperscale cloud infrastructure. $300 billion in Gulf AI and data centre investments are now at risk. The cloud's physical layer is no longer an abstraction — it's a target.*
>
> — CSIS analysis; The Information, Semafor (March 2026)

What this means

CSIS describes this as the first major military strike on hyperscale cloud infrastructure. $300 billion in Gulf AI and data centre investments are now at risk, according to reporting from The Information and Semafor. AWS acknowledged that the "broader operating environment in the Middle East remains unpredictable." Multi-AZ redundancy within a single region — the standard resilience playbook — assumed that availability zones could fail independently through technical faults, not simultaneously through military action. That assumption is broken. For any organisation running production workloads in a geopolitically contested region, the conversation has shifted from "how many AZs do we need?" to "how many *regions* do we need, and what happens to our compliance posture when we can't use the one the regulator requires?"

Sources: CNBC, CBS News, The Register, CSIS, AWS Health Dashboard

AI & TOOLING

The AI Coding Arms Race Enters a New Phase

95%	26.9%	1.6M	29%
WEEKLY AI USAGE	AI-AUTHORED CODE	CODEX WEEKLY USERS	TRUST AI-GENERATED CODE

Two surveys published this month give the clearest picture yet of where AI adoption actually stands — and they tell the same story from different angles. The Pragmatic Engineer surveyed 906 engineers and engineering leaders (median 11–15 years experience): 95% use AI tools at least weekly, 75% for half or more of their work, 56% for over 70%. Laura Tacho's DX data — 121,000 developers across 450+ companies — found 92.6% monthly usage, with 26.9% of production code now AI-authored, up from 22%. This is saturation, not adoption.

Claude Code's eight-month ascent

The Pragmatic Engineer survey's most striking result: Claude Code is now the #1 most-used AI coding tool, overtaking GitHub Copilot and Cursor in just eight months since its May 2025 launch. It commands 46% of "most loved" mentions — 57% including Claude models generally. At small companies, where engineers choose their own tools, Claude Code leads at 75%. At enterprises with 10,000+ employees, Copilot dominates at 56%, driven by Microsoft procurement bundling. The market is bifurcated between what engineers choose and what procurement buys. Staff+ engineers are the heaviest agent users at 63.5%, and agent users are twice as excited about AI as non-users.

The trust paradox

Only 29% of developers trust AI-generated code — down from 40% on Stack Overflow's earlier confidence metric. Yet adoption keeps accelerating. Engineers are using tools they don't trust for the majority of their work. The METR study adds a telling data point: their early 2025 research found AI actually caused a 20% productivity slowdown, but the latest update admits they can't measure reliably anymore because developers refuse to work without AI. When you can't run the control group, the experiment is over.

> *95% weekly adoption, 29% trust. Engineers are using AI tools for the majority of their work while trusting them less than ever. The gap between adoption and confidence is the defining tension of 2026.*
>
> — Data from The Pragmatic Engineer AI Tooling Survey and Stack Overflow, 2026

The platform contest

OpenAI isn't conceding the agentic market. Codex tripled to 1.6 million weekly active users with the GPT-5.3 Codex launch, with over a million desktop downloads and 5x token usage growth. Enterprise customers include Cisco, Nvidia, Ramp, and Rakuten. Meanwhile, Anthropic is on a $19 billion annualised revenue run rate — up $6 billion in February alone. Ramp data shows Anthropic now captures 60% of business AI chatbot invoices, up from 10% a year ago. Microsoft's Copilot Wave 3 (March 9) brings Claude into mainline Copilot chat, with Agent 365 and the E7 Frontier Suite at $99/user launching May 1. The contest is no longer about code completion. It's about who owns the agent layer between the engineer and the codebase.

The productivity plateau nobody wants to discuss

DX's data reveals something the adoption headlines obscure: measurable productivity gains have plateaued at roughly 10%. Sonar's survey reinforces the pattern — 72% of developers who've tried AI tools use them daily, and 42% of code is AI-generated or assisted. Near-universal adoption, declining trust, and plateauing measurable productivity. The gap between the market narrative and the observable data is widening. For engineering leaders, the question isn't "should we use AI tools" — that's settled — it's "are we measuring the right things to know whether they're actually helping?"

Sources: The Pragmatic Engineer, Fortune, Sonar, ShiftMag/DX, GeekWire

ENGINEERING LEADERSHIP

The AI Delivery Bottleneck: Teams Write More Code, Ship Less

59%	70.8%	72 min	<5%
THROUGHPUT INCREASE	SUCCESS RATE (5YR LOW)	RECOVERY TIME	TEAMS AT AI SPEED

CircleCI's 2026 State of Software Delivery report — 28 million CI/CD workflows analysed across thousands of teams — reveals the paradox at the heart of the AI coding boom. Average throughput jumped 59% year-over-year, the biggest increase since they started tracking. But look at the main branch, where code actually ships, and the picture inverts: delivery throughput for the median team is *down* 7%. Teams are producing more code than ever and deploying it more slowly.

Main branch success rates have fallen to 70.8%, the lowest in five-plus years, against a benchmark of 90%. Recovery times are up 13% to 72 minutes on average. Feature branch throughput for the median team is up 15%, which means more work-in-progress accumulating in branches that aren't making it to production. The pipeline from "code written" to "code shipped" has a bottleneck, and AI is making it wider.

The distribution is the story

The top 5% of teams saw throughput jump 97%, with main branch throughput up 26% and feature branch throughput up 85%. The median team got a 4% improvement. The bottom quartile got nothing. AI is amplifying existing delivery strengths, not distributing them. The teams that already had strong CI/CD pipelines, good testing practices, and reliable deployment processes captured nearly all the benefit. Everyone else just generated more code that piles up in branches and fails in pipelines. Mid-sized companies (21–50 employees) are stuck in what CircleCI calls the "messy middle" — lowest throughput, recovery times averaging three hours.

> *Writing code is no longer the constraint. Review, validation, and integration is where AI-generated code piles up. Fewer than 1 in 20 teams have figured out how to ship at AI speed. The other 19 are drowning in work-in-progress.*
>
> — CircleCI 2026 State of Software Delivery

What this means for engineering leaders

AI tooling has solved the code generation problem. It has not solved — and is actively worsening — the code delivery problem. At a 70% success rate, a team running 500 daily changes loses the equivalent of 12 full-time engineers to pipeline recovery alone. The organisations that will pull ahead are the ones investing as heavily in CI/CD reliability, automated testing, security scanning, and review processes as they are in AI coding tools. Right now, that's fewer than 1 in 20.

Sources: CircleCI, CircleCI 2026 Report

BANKING & FINTECH

Revolut Gets Its Full UK Banking Licence

13M	$75B	£3B	£120K
UK CUSTOMERS	VALUATION	UK INVESTMENT PLANNED	FSCS PROTECTION

After a four-year application process and 18 months in mobilisation, the PRA lifted restrictions on March 11 and Revolut Bank UK Ltd became a fully licensed bank. Eligible deposits are now FSCS-protected up to £120,000 — the same safety net customers get at Barclays, HSBC, or Lloyds. This is the moment the fintech narrative shifts from "digital challengers nibbling at margins" to "a direct competitor with regulatory parity."

The licence unlocks what Revolut has been waiting for: lending and credit products. These are the real revenue engines of banking, and they've been locked behind full authorisation. Current accounts are rolling out gradually to new customers; existing customers remain on the e-money entity for now, with migration over several months. Revolut has committed £3 billion to UK investment and plans to create 1,000 high-skilled jobs. They've also applied for a US OCC charter.

The competitive reality

Bain's latest analysis projects challenger banks will control 35% of the UK market by 2030, up from roughly 20% today. Chris Skinner's data adds nuance: over 62% of UK customers already use at least one challenger bank, but only 14% have fully abandoned traditional institutions. The battleground

isn't customer acquisition — Revolut already has 13 million UK customers. It's becoming the primary banking relationship. FSCS protection and credit products are the tools that make that shift possible.

The IPO question

Revolut is reportedly exploring a dual London/New York IPO at its $75 billion valuation, which would make it one of the largest fintech listings in history. The Bank of England governor reportedly blocked a chancellor-arranged meeting between Revolut and regulators over concerns about political interference in the licensing process — a signal of how seriously the establishment takes this, and how politically charged fintech regulation has become in the UK.

For anyone building payment infrastructure, loyalty programmes, or customer-facing banking technology: this is a competitor with 13 million customers that just gained the full regulatory platform to offer lending, credit, and protected banking services — powered entirely by an engineering-led operating model that traditional banks cannot replicate without rebuilding from the ground up.

Sources: Bloomberg, CNBC, PYMNTS, Tech.eu, City AM, FinTech Futures

OTHER NEWS THIS MONTH

In Brief

Open banking goes mainstream — Amazon and eBay launched Pay by Bank in the UK via TrueLayer. Account-to-account payments via Faster Payments, biometric authentication through banking apps, no card details needed. The moment open banking stops being a fintech buzzword and becomes a checkout option. *Open Banking UK, TrueLayer*

GitHub removes premium models from free Copilot Student plan (March 12) — GPT-5.4, Claude Opus/Sonnet removed, affecting ~2M students. Students keep Auto mode. Community response: 1,836 downvotes, 818 comments. The economics of bundled frontier model access are starting to bite. *GitHub Community*

Monzo deployed full-stack backup bank during January outage — "Monzo Stand-in" kept card payments, cash withdrawals, and transfers running for 2 hours that would have been offline. Second real-world activation. *Monzo Engineering Blog*

GoodFirms survey: 91% of software companies use AI to cut development costs — framing AI adoption as a cost play rather than a productivity play. *GoodFirms*

Chrome moving to 14-day release cycle — Selenium blog says mostly a non-event for automation users, but teams running browser-specific test suites should verify compatibility cadence. *selenium.dev*

PLATFORM UPDATES

At a Glance

PLATFORM	WHAT CHANGED
.NET	Security servicing for .NET 10/9/8. Emergency .NET 10.0.5 out-of-band fix for macOS debugging (March 12). C# MCP SDK now available.
Node.js	24.14.0 LTS released (Feb 24). Node 20 EOL April 30 — migration should be underway. New release schedule: one major per year from Node 27.
GitHub	Copilot coding agent updates: model picker, self-review, security scanning, custom agents, CLI handoff. Student plan restructured — premium models removed from free tier.
AWS	S3 turned 20. Route 53 Global Resolver GA. Stateful MCP in Bedrock AgentCore. ME region recovery ongoing — no restoration date.

THE LEADERSHIP READ

Agentic Engineering in Practice: What the Pragmatic Summit Told Us

Simon Willison — prolific open-source contributor, creator of Datasette, and one of the most respected voices in AI-assisted development — sat down for a fireside chat at Gergely Orosz's Pragmatic Summit in San Francisco. The recording was published on March 14. It's the most useful thing published this month for anyone managing an engineering team that uses AI tools.

The adoption curve is real — and your people are at different points on it

Willison described distinct phases engineers move through: starting with ChatGPT for occasional questions, progressing to AI agents that write bits of code, then reaching a tipping point where agents produce more code than the human. The newest phase — one Willison says he's now in — is where "you don't read the code." The agent writes it, runs it, tests it, and you evaluate the output rather than reviewing the implementation. Understanding where each person on your team sits on this curve matters, because they need fundamentally different support at each stage.

Mental exhaustion is the new bottleneck — not laziness

> *After two hours of that I'm done for the day. I'm mentally exhausted. People worry about skill atrophy and being lazy. I think this is the opposite of that. You have to operate firing on all cylinders if you're going to keep your trio or quadruple of agents busy solving all these different problems.*
>
> — Simon Willison, Pragmatic Summit, 2026

Running 3–4 AI agent sessions in parallel is technically possible. Willison does it routinely. But it's cognitively draining in a way traditional coding isn't — the context-switching, quality evaluation, and prompt refinement all happening simultaneously. Two hours of high-intensity agent orchestration can be more productive than a full day of traditional development. The engineers using AI most effectively may look like they're working less. Traditional productivity metrics — hours logged, lines of code, PRs merged — are becoming meaningless as signals of output quality.

Speed is up, but capacity has a ceiling

AI makes individual engineers dramatically faster. But one engineer still cannot scale to a thousand projects. Context-switching between agent sessions is the real bottleneck — not the agent's throughput, but the human's ability to maintain context. This connects directly to the CircleCI data elsewhere in this issue: throughput is up 59%, but fewer than 5% of teams are shipping at AI speed. The constraint has moved from code generation to everything that happens after.

Maintain a "retry list"

The most practical suggestion: keep track of tasks that AI tools failed at, and re-test them every six months as models improve. Willison describes regularly being surprised when something impossible three months ago now works perfectly. The capabilities of frontier models change faster than anyone's intuitions about them — and the teams that systematically re-test their assumptions will find capabilities that others miss.

What this means for engineering management

Productivity measurement needs rethinking — output is disconnecting from hours. Team practices need to account for the mental intensity of AI-assisted work and build in sustainable rhythms. Investment in tooling access and model selection is becoming as important as investment in headcount. The engineers who lean furthest into AI agents aren't the ones who need less management attention — they're navigating the steepest learning curve, making the most consequential decisions about when to trust and when to verify, and operating at the highest cognitive load.

Read the full post: simonwillison.net/2026/Mar/14/pragmatic-summit/ →

Monzo's Backup Bank: Engineering Resilience When the Cloud Goes Down

In January 2026, Monzo experienced an outage. What happened next was remarkable: instead of customers losing access for two hours, Monzo activated "Stand-in" — a fully independent backup bank that kept card payments, cash withdrawals, and transfers running throughout the incident. Customers barely noticed. In a month where drone strikes took an entire AWS region offline for weeks, Monzo's architecture is a case study in what serious resilience engineering actually looks like.

What Stand-in actually is

Monzo Stand-in isn't a degraded mode or a cached read replica. It's a fully independent system capable of processing core banking transactions — card payments, ATM withdrawals, incoming and outgoing transfers — without the primary platform being available. It runs on two independent infrastructure stacks, designed as the "backup of last resort." Monzo first described the architecture in a 2025 blog post titled "Tolerating Full Cloud Outages" and has been building toward it since.

Battle-tested, twice

The January 2026 incident was the second real-world activation. The first was a shorter, one-hour outage in 2024. Both times, Stand-in handled live customer transactions without data loss or service degradation. For a UK bank serving millions of customers, this is an extraordinary engineering achievement — and it directly addresses the operational resilience requirements that regulators have been tightening under DORA (the Digital Operational Resilience Act).

Why this matters beyond Monzo

Most organisations' disaster recovery plans are theoretical. They exist in runbooks and are tested in controlled environments with synthetic traffic. Monzo built a system that has now been validated under real production conditions, with real customer money, twice. The engineering investment is significant — maintaining two independent infrastructure stacks is expensive and operationally complex. But the alternative is downtime during which customers can't access their money. After the AWS Middle East crisis demonstrated that entire regions can go dark for weeks, the question for every engineering leader running financial services workloads isn't whether this level of resilience investment is justified. It's whether you can afford not to make it.

Sources: Monzo Engineering Blog, CrowdFund Insider, CX Today

ON THE RADAR

Coming Up

April 19	Claude Haiku 3 model retirement — migrate to Haiku 4.5
April 22	AWS Summit London — expect Bedrock and Trainium announcements
April 29	Zero Day Con, Dublin — Ireland's largest cybersecurity conference
April 30	Node.js 20 end of life
May 1	Microsoft Copilot Wave 3 / Agent 365 + E7 Frontier Suite launch
May 11–12	Dubai Fintech Summit

RECOMMENDED READING

This Month's Links

AI Tooling for Software Engineers in 2026	The Pragmatic Engineer
2026 State of Software Delivery	CircleCI
My fireside chat about agentic engineering	simonwillison.net
AWS drone strike analysis	CNBC
Revolut Launches UK Bank	revolut.com
Manifesto for Enterprise Agility	PMI
Tolerating Full Cloud Outages	Monzo Engineering Blog

Engineering Digest

February 18 – March 18, 2026 · Inaugural Edition · Issue #0

Next edition: March 25, 2026