Engineering *Digest*

February 18 – March 18, 2026 | Issue #0

MONTHLY RETROSPECTIVE

IN THIS ISSUE

The AI Coding Arms Race Enters a New Phase
The AI Delivery Bottleneck
The Leadership Read: Agentic Engineering in Practice
Platform Updates, Other News & On the Radar
Revolut Gets Its Full UK Banking Licence
When Drones Hit the Cloud: AWS and the Middle East Crisis
Outside In: Monzo's Backup Bank

FOREWORD

Something clarifying happened over the past month. AI adoption crossed from widespread to near-universal — 95% weekly usage in the Pragmatic Engineer survey, 92.6% monthly in DX's 121,000-developer dataset — but the infrastructure around it didn't keep pace. CircleCI's analysis of 28 million workflows shows the consequences: throughput up 59%, main branch delivery *down* 7%, success rates at a five-year low. Engineers are writing more code than ever. They're shipping it more slowly. That gap is the defining challenge for engineering leadership right now.

Meanwhile, the physical world reminded the cloud that it still exists. Drone strikes in the UAE and Bahrain knocked out AWS Middle East infrastructure for weeks — 109+ services affected, two-thirds of availability zones offline, 16+ days without full restoration. For engineering leaders who believed multi-AZ was sufficient resilience, the data residency trade-offs just became viscerally real. And in UK banking, Revolut secured full PRA authorisation, giving an engineering-led digital platform the same regulatory standing as HSBC.

This is the inaugural issue of Engineering Digest — a monthly retrospective for people who lead engineering teams and need signal, not noise. No sponsored content, no hype cycles, no listicles. Just the things that matter, framed through the lens of the people who have to make decisions about them.

THE MONTH AT A GLANCE

Seven stories in sixty seconds

AI **Microsoft Copilot Wave 3 integrates Claude** (March 9) — Anthropic's Cowork technology powers Agent 365. Claude in mainline Copilot chat. E7 Frontier Suite at $99/user launching May 1.

AI **Anthropic launched the $100M Claude Partner Network** — certified training, Claude Certified Architect programme, co-marketing. Plus: Cowork plugins update (Feb 24) with private marketplaces, 12 MCP connectors, per-user provisioning.

AWS **ME drone strikes hit two UAE data centres** (March 1) — Bahrain facility also damaged. 109+ services affected. Physical recovery ongoing 16+ days. AWS recommends migrating workloads to alternate regions.

Process **PMI and Agile Alliance released the Manifesto for Enterprise Agility** (March 3) — first joint manifesto. Extends agile from team delivery to enterprise governance, funding, and leadership.

.NET **March security patches across .NET 10/9/8.** Emergency .NET 10.0.5 out-of-band fix for macOS debugging (March 12). C# MCP SDK now available.

Node **Node.js 24.14.0 LTS released** (Feb 24). Node 20 EOL April 30. New release schedule announced: one major per year from Node 27.

Security **March Patch Tuesday: 83 CVEs, 8 critical, 2 zero-days.** Notable: Azure MCP Server SSRF (CVE-2026-26118, 8.8 CVSS) and Excel XSS that could cause Copilot Agent mode data exfiltration.

AI & TOOLING

The AI Coding Arms Race Enters a New Phase

95%	26.9%	1.6M	29%
WEEKLY AI USAGE	PRODUCTION CODE AI-AUTHORED	CODEX WEEKLY USERS	TRUST AI-GENERATED CODE

Two surveys published in the past month give us the clearest picture yet of AI adoption in engineering — and they tell the same story from different angles. The Pragmatic Engineer surveyed 906 engineers and engineering leaders (median 11–15 years experience): 95% use AI tools at least weekly, 75% for half or more of their work, 56% for 70% or more. Laura Tacho's DX study — 121,000 developers across 450+ companies — found 92.6% monthly usage, with 26.9% of production code now AI-authored, up from 22%. This is saturation, not adoption.

Claude Code's eight-month ascent

The Pragmatic Engineer survey's most striking result: Claude Code is now the #1 most-used AI coding tool, overtaking GitHub Copilot and Cursor in just eight months since its May 2025 launch. It commands 46% of "most loved" mentions — 57% when you include Claude models generally. At small companies, where engineers choose their own tools, Claude Code leads at 75%. At enterprises with 10,000+ employees, Copilot dominates at 56%, driven by Microsoft procurement bundling. The market is bifurcated between what engineers choose and what procurement buys.

The trust paradox

Here's the number that should give everyone pause: only 29% of developers trust AI-generated code, down from 40% on Stack Overflow's earlier confidence metric. Yet adoption keeps accelerating. Engineers are using tools they don't trust for the majority of their work. Sonar's survey reinforces the pattern — 72% of developers who've tried AI tools use them daily, and 42% of code is now AI-generated or assisted. The METR study adds another wrinkle: their early 2025 research found AI actually caused a 20% productivity slowdown, but their latest update admits they can't measure reliably anymore because developers refuse to work without AI. When you can't run the control group, the experiment is over.

> *Staff+ engineers are the heaviest agent users at 63.5%, and agent users are twice as excited about AI as non-users. The enthusiasm gap between power users and everyone else is widening — even as overall trust in AI-generated code declines.*
>
> — Data from The Pragmatic Engineer AI Tooling Survey 2026

OpenAI's Codex counter-move

OpenAI isn't conceding the agentic market. Codex tripled to 1.6 million weekly active users following the GPT-5.3 Codex launch in early February, with over a million desktop app downloads and 5x token usage growth. Enterprise customers include Cisco, Nvidia, Ramp, and Rakuten. OpenAI is positioning Codex as "the standard agent" — expanding beyond coding into broader enterprise domains. The contest is no longer about code completion. It's about who owns the agent layer between the engineer and the codebase.

The productivity plateau nobody wants to discuss

DX's data reveals something the adoption headlines obscure: measurable productivity gains have plateaued at roughly 10%. That's real, but it's not the transformation the market is pricing in. Near-universal adoption, declining trust, and plateauing measurable productivity — the gap between the hype narrative and the observable data is widening. For engineering leaders, the practical question isn't "should we use AI tools" (that's settled) but "are we measuring the right things to know whether they're actually helping?"

Sources: The Pragmatic Engineer, Fortune, Sonar, Stack Overflow

BANKING & FINTECH

Revolut Gets Its Full UK Banking Licence

13M	$75B	£3B	£120K
UK CUSTOMERS	VALUATION	UK INVESTMENT PLANNED	FSCS PROTECTION

After a four-year application process and 18 months in mobilisation, the PRA lifted restrictions on March 11 and Revolut Bank UK Ltd became a fully licensed bank. Eligible deposits are now FSCS-protected up to £120,000 — the same safety net customers get at Barclays, HSBC, or Lloyds. This is the moment the fintech narrative shifts from "digital challengers nibbling at margins" to "a direct competitor with regulatory parity."

The licence unlocks what Revolut has been waiting for: lending and credit products. These are the real revenue engines of banking, and they've been locked behind full authorisation. Current accounts are rolling out gradually to new customers; existing customers remain on the e-money entity for now, with migration over several months. Revolut has committed £3 billion to UK investment and plans to create 1,000 high-skilled jobs.

The competitive reality

Bain & Co's latest analysis projects challenger banks will control 35% of the UK market by 2030, up from roughly 20% today. Chris Skinner's data adds nuance: over 62% of UK customers already use at least one challenger bank, but only 14% have fully abandoned traditional institutions. The battleground isn't customer acquisition — Revolut already has 13 million UK customers. It's becoming the primary banking relationship. FSCS protection and credit products are the tools that make that possible.

The IPO question

Revolut is reportedly exploring a dual London/New York IPO at its $75 billion valuation, which would make it one of the largest fintech listings in history. The Bank of England governor reportedly blocked a chancellor-arranged meeting between Revolut and regulators over concerns about political interference in the licensing process — a signal of how seriously the establishment takes this, and how politically charged fintech regulation has become in the UK.

For anyone building payment infrastructure, loyalty programmes, or customer-facing banking technology: this is a competitor with 13 million customers that just gained the full regulatory platform to offer lending, credit, and protected banking services — powered entirely by an

engineering-led operating model that traditional banks cannot replicate without rebuilding from the ground up.

Sources: Bloomberg, CNBC, Reuters, PYMNTS, Tech.eu, City AM, Chris Skinner / The Finanser

ENGINEERING LEADERSHIP

The AI Delivery Bottleneck: Teams Write More Code, Ship Less

59%	70.8%	72 min	<5%
THROUGHPUT INCREASE	SUCCESS RATE (5YR LOW)	RECOVERY TIME	TEAMS AT AI SPEED

CircleCI's 2026 State of Software Delivery report — 28 million CI/CD workflows analysed across thousands of teams — reveals the paradox at the heart of the AI coding boom. Average throughput jumped 59% year-over-year, the biggest increase since they started tracking. But look at the main branch, where code actually ships, and the picture inverts: delivery throughput for the median team is *down* 7%. Teams are producing more code than ever and deploying it more slowly.

The numbers get worse the deeper you look. Main branch success rates have fallen to 70.8% — the lowest in five-plus years, against a benchmark of 90%. Recovery times are up 13% to 72 minutes on average. Feature branch throughput for the median team is up 15%, which means more work-in-progress accumulating in branches that aren't making it to production. The pipeline from "code written" to "code shipped" has a bottleneck, and AI is making it wider.

The distribution is the story

The top 5% of teams saw throughput jump 97%. The median team got 4%. The bottom quartile got nothing. AI is amplifying existing delivery strengths, not distributing them. The teams that already had strong CI/CD pipelines, good testing practices, and reliable deployment processes captured nearly all the benefit. Everyone else just generated more code that piles up in branches and fails in pipelines. Mid-sized companies (21–50 employees) are stuck in what CircleCI calls the "messy middle" — lowest throughput, recovery times averaging three hours.

> *Writing code is no longer the constraint. Review, validation, and integration is where AI-generated code piles up. Fewer than 1 in 20 teams have figured out how to ship at AI speed. The other 19 are drowning in work-in-progress.*
>
> — CircleCI 2026 State of Software Delivery

What this means for engineering leaders

AI tooling has solved the code generation problem. It has not solved — and is actively worsening — the code delivery problem. At a 70% success rate, a team running 500 daily changes loses the equivalent of 12 full-time engineers to recovery alone. The teams that will pull ahead are the ones investing as heavily in CI/CD reliability, automated testing, security scanning, and review processes as they are in AI coding tools. Right now, that's fewer than 1 in 20.

Sources: CircleCI, Thoughtworks

CLOUD & INFRASTRUCTURE

When Drones Hit the Cloud: The AWS Middle East Crisis

109+	3	2 of 3	16+
SERVICES AFFECTED	FACILITIES STRUCK	AZS OFFLINE	DAYS TO RECOVERY

On March 1, drone strikes hit AWS data centre facilities in the UAE and Bahrain as part of the broader US-Iran military conflict in the region. Two of three availability zones in the me-south-1 region went offline. Over 109 AWS services were affected — not degraded, not throttled, but unavailable. Physical damage to infrastructure required hardware replacement, not software failover. More than two weeks later, full restoration had not been confirmed. AWS's official guidance: migrate workloads to alternate regions.

The banking impact

Financial institutions across the Gulf were hit hardest. Banks operating under UAE and Bahrain Central Bank data residency requirements — which mandate that customer financial data remain within national borders — found themselves in an impossible position. Their data had to stay in-region by law, but the region's infrastructure was physically destroyed. Some institutions reported being unable to process transactions for days. The Dubai Financial Services Authority issued emergency guidance permitting temporary cross-region data processing, but the regulatory patchwork across different Gulf states created confusion about what was actually permitted and for how long.

The data residency paradox

This is the scenario that cloud resilience planning never properly accounted for. Multi-AZ architecture — the standard recommendation for high availability — assumes that availability zones fail independently due to software bugs, power outages, or network partitions. It does not assume

that a military strike takes out multiple facilities simultaneously within the same geographic region. For organisations bound by data sovereignty regulations, the standard playbook of "fail over to another region" is legally prohibited. The architecture that regulators required for data protection became the architecture that prevented disaster recovery.

> *The cloud's physical layer has always been someone else's data centre. When those data centres sit in a conflict zone, "high availability" means something different than the whitepaper suggested.*

The implications for engineering leadership

Every organisation running workloads in the Middle East — or any geopolitically sensitive region — needs to revisit three assumptions. First, multi-AZ is not sufficient when the threat model includes physical attacks on concentrated infrastructure. Second, data residency regulations and disaster recovery requirements can be fundamentally incompatible, and that tension needs to be resolved before the crisis, not during it. Third, cloud provider SLAs do not cover acts of war — the financial exposure for extended outages falls entirely on the customer. CSIS analysis noted this was the first major military strike on hyperscale cloud infrastructure, but it is unlikely to be the last. The question for engineering leaders is whether their business continuity plans account for a world where cloud regions can be physically destroyed.

Sources: Reuters, CNBC, BBC News, The National News, CSIS

OTHER NEWS THIS MONTH

In Brief

Open banking goes mainstream — Amazon and eBay launched Pay by Bank in the UK via TrueLayer, bypassing card networks entirely. Biometric auth via banking apps, instant settlement over Faster Payments. A new payment rail engineering teams will need to support. *Open Banking UK, TrueLayer*

GitHub removes premium models from free Copilot Student plan (March 12) — GPT-5.4, Claude Opus/Sonnet removed, affecting ~2M students. Students keep Auto mode. Community response: 1,836 downvotes, 818 comments. *GitHub Community*

Monzo deployed full-stack backup bank during January outage — "Monzo Stand-in" kept card payments, cash withdrawals, and transfers running for 2 hours that would have been offline. *Monzo Engineering Blog*

GoodFirms survey: 91% of software companies use AI to cut development costs — framing AI adoption as a cost play rather than a productivity play. *GoodFirms*

Chrome moving to 14-day release cycle — Selenium blog says mostly a non-event for automation users, but

teams running browser-specific test suites should verify compatibility cadence. *selenium.dev*

PLATFORM UPDATES

At a Glance

PLATFORM	WHAT CHANGED
.NET	Security servicing for .NET 10/9/8. Emergency .NET 10.0.5 out-of-band fix for macOS debugging (March 12). C# MCP SDK now available.
Node.js	24.14.0 LTS released (Feb 24). Node 20 EOL April 30 — migration should be underway. New release schedule: one major per year from Node 27.
GitHub	Copilot coding agent updates: model picker, self-review, security scanning, custom agents, CLI handoff. Student plan restructured — premium models removed from free tier.
AWS	S3 turned 20. Route 53 Global Resolver GA. Stateful MCP in Bedrock AgentCore. ME region recovery ongoing.

THE LEADERSHIP READ

Agentic Engineering in Practice: What the Pragmatic Summit Told Us

Simon Willison — prolific open-source contributor, creator of Datasette, and one of the most respected voices in AI-assisted development — sat down for a fireside chat at Gergely Orosz's Pragmatic Summit in San Francisco. The recording was published on March 14. It's the most useful thing published this month for anyone managing an engineering team that uses AI tools.

The adoption curve is real — and your people are at different points on it

Willison described distinct phases engineers move through: starting with ChatGPT for occasional questions, progressing to AI agents that write bits of code, then reaching a tipping point where agents produce more code than the human. The newest phase — one Willison says he's now in — is where "you don't read the code." The agent writes it, runs it, tests it, and you evaluate the output rather than reviewing the implementation. Understanding where each person on your team sits on this curve matters, because they need fundamentally different support at each stage.

Mental exhaustion is the new bottleneck — not laziness

> *After two hours of that I'm done for the day. I'm mentally exhausted. People worry about skill atrophy and being lazy. I think this is the opposite of that. You have to operate firing on all cylinders if you're going to keep your trio or quadruple of agents busy solving all these different problems.*
>
> — Simon Willison, Pragmatic Summit, 2026

Running 3–4 AI agent sessions in parallel is technically possible. Willison does it routinely. But it's cognitively draining in a way traditional coding isn't — the context-switching, quality evaluation, and prompt refinement all happening simultaneously. Two hours of high-intensity agent orchestration can be more productive than a full day of traditional development. The engineers using AI most effectively may look like they're working less. Traditional productivity metrics — hours logged, lines of code, PRs merged — are becoming meaningless as signals of output quality.

Speed is up, but capacity has a ceiling

AI makes individual engineers dramatically faster. But one engineer still cannot scale to a thousand projects. Context-switching between agent sessions is the real bottleneck — not the agent's throughput, but the human's ability to maintain context. This connects directly to the CircleCI data elsewhere in this issue: throughput is up 59%, but only 5% of teams are shipping at AI speed. The constraint has moved from code generation to everything that happens after code is generated.

Maintain a "retry list"

The most practical suggestion: keep track of tasks that AI tools failed at, and re-test them every six months as models improve. Willison describes regularly being surprised when something impossible three months ago now works perfectly. The capabilities of frontier models change faster than anyone's intuitions about them — and the teams that systematically re-test their assumptions will find capabilities that others miss.

What this means for engineering management

Productivity measurement needs rethinking — output is disconnecting from hours. Team practices need to account for the mental intensity of AI-assisted work and build in sustainable rhythms. Investment in tooling access and model selection is becoming as important as investment in headcount. The engineers who lean furthest into AI agents aren't the ones who need less management attention — they're navigating the steepest learning curve, making the most consequential decisions about when to trust and when to verify, and operating at the highest cognitive load.

Read the full post: simonwillison.net/2026/Mar/14/pragmatic-summit/ →

OUTSIDE IN

Monzo's Backup Bank: Engineering Resilience When the Cloud Goes Down

In January 2026, Monzo experienced an outage. What happened next was remarkable: instead of customers losing access for two hours, Monzo activated "Stand-in" — a fully independent backup bank that kept card payments, cash withdrawals, and transfers running throughout the incident. Customers barely noticed.

What Stand-in actually is

Monzo Stand-in isn't a degraded mode or a cached read replica. It's a fully independent system capable of processing core banking transactions — card payments, ATM withdrawals, incoming and outgoing transfers — without the primary platform being available. It runs on two independent infrastructure stacks, designed as the "backup of last resort." Monzo first described the architecture in a 2025 blog post titled "Tolerating Full Cloud Outages" and has been building toward it since.

Battle-tested, twice

The January 2026 incident was the second real-world activation. The first was a shorter, one-hour outage in 2024. Both times, Stand-in handled live customer transactions without data loss or service degradation. For a UK bank serving millions of customers, this is an extraordinary engineering achievement — and it directly addresses the operational resilience requirements that regulators have been tightening under frameworks like DORA (the Digital Operational Resilience Act).

Why this matters beyond Monzo

Most organisations' disaster recovery plans are theoretical. They exist in runbooks and are tested in controlled environments with synthetic traffic. Monzo's approach is different: they built a system that has now been validated under real production conditions, with real customer money, twice. The engineering investment is significant — maintaining two independent infrastructure stacks is expensive and operationally complex. But the alternative is downtime during which customers can't access their money, and in regulated banking, that's an existential risk.

For engineering leaders in financial services — or any domain where downtime has regulatory or reputational consequences — the question Monzo has answered is whether it's possible to build genuinely independent failover for a modern cloud-native banking platform. It is. The question for everyone else is whether the investment is worth it before the outage forces the decision.

Sources: Monzo Engineering Blog, CrowdFund Insider, CX Today

ON THE RADAR

Coming Up

April 19	Claude Haiku 3 model retirement — migrate to Haiku 4.5
April 22	AWS Summit London — expect Bedrock and Trainium announcements
April 29	Zero Day Con, Dublin — Ireland's largest cybersecurity conference
April 30	Node.js 20 end of life
May 1	Microsoft Copilot Wave 3 / Agent 365 + E7 Frontier Suite launch
May 11–12	Dubai Fintech Summit

RECOMMENDED READING

This Month's Links

AI Tooling for Software Engineers in 2026	The Pragmatic Engineer
2026 State of Software Delivery	CircleCI
My fireside chat about agentic engineering	simonwillison.net
Military Strikes on Cloud Infrastructure: Strategic Implications	CSIS
Revolut Launches UK Bank	revolut.com
Manifesto for Enterprise Agility	PMI
Tolerating Full Cloud Outages	Monzo Engineering Blog

Engineering Digest

February 18 – March 18, 2026 · Inaugural Edition · Issue #0

Next edition: March 25, 2026