Engineering *Digest*

February 17 – March 17, 2026 | Issue #0

INAUGURAL MONTHLY EDITION

IN THIS ISSUE

The AI Tooling Arms Race Enters a New Phase

AI Is Generating Code Faster Than Teams Can Ship It

The Leadership Read: Why Would Anyone Be an Engineering Manager in 2026?

Platform Updates, Other News & On the Radar

Revolut Gets Its Full UK Banking Licence

Open Banking's Landmark Year and the Push to Open Finance

Outside In: Can AI Agents Build Real Stripe Integrations?

FOREWORD

Three forces collided over the past month that together define the position engineering leadership finds itself in right now. AI tooling went from useful to ubiquitous — 95% weekly usage according to the Pragmatic Engineer's definitive survey, Claude Code crowned the most-loved tool barely eight months after launch. The delivery infrastructure buckled under the weight: CircleCI's analysis of 28 million workflows shows throughput up 59% but main branch delivery actually *down* 7%, with success rates at a five-year low. And the fintech landscape shifted again, with Revolut earning the same banking credentials as Barclays and an EY analysis valuing Open Banking at £43 billion for the UK economy.

For engineering leaders, the question is no longer whether to adopt AI but how to ensure the rest of the machine — testing, security, review, delivery pipelines — catches up with the new speed of code generation. The tooling is moving faster than the processes designed to govern it. That tension runs through every article in this inaugural edition.

Welcome to Issue 0. This is a digest built for people who run engineering teams and need signal, not noise. No listicles, no hype cycles, no sponsored content. Just the things that matter, framed through the lens of the people who have to make decisions about them.

THE MONTH AT A GLANCE

Eight stories in sixty seconds

AI **Pragmatic Engineer survey: 95% weekly AI usage, Claude Code #1 tool** — 906 respondents, 75% using AI for half or more of their work. Claude Code overtook Copilot and Cursor in just eight months. *Pragmatic Engineer*

AI **OpenAI Codex hits 1.6M weekly users** — tripled since GPT-5.3 Codex launch. Enterprise customers include Cisco, Nvidia, Ramp. Positioning as "the standard agent" beyond coding. *Fortune*

AI **OpenAI launches Codex Security** — scanned 1.2M commits, found 792 critical and 10,561 high-severity issues. Free research preview for Enterprise, Business, Edu customers. *The Hacker News*

AI **Cursor launches Automations** — always-on cloud agents triggered by events from Slack, GitHub, PagerDuty. Hundreds of automations per hour. JetBrains ACP integration live. *TechCrunch*

Banking **Revolut secures full UK banking licence** — FSCS protection for 13M customers. Credit products coming, dual London/NY IPO at $75B valuation on the table. *CNBC*

Banking **Open Banking valued at £43B potential annual benefit** — EY/OBL analysis. Already generated £8.3B cumulative value. 17.5M+ user connections, 53% YoY payment growth. *FinTech Global*

Process **CircleCI: AI boosted throughput 59% but main branch delivery slowing** — 28M workflows analysed. Success rates at 70.8%, a five-year low. Only 1 in 20 teams shipping at AI speed. *CircleCI*

AI **Snyk: 48% of AI-generated code has security vulnerabilities** — time-to-exploit shrinking, expected to accelerate 50% by 2027. The speed-safety trade-off is real. *Snyk*

AI & TOOLING

The AI Tooling Arms Race Enters a New Phase

95%	55%	4.7M	1.6M
WEEKLY AI USAGE	USE AGENTS REGULARLY	COPILOT PAID SUBS	CODEX WEEKLY USERS

The Pragmatic Engineer's 2026 AI tooling survey is the most comprehensive snapshot we have of how engineers actually use AI — 906 respondents, mostly engineers and engineering leaders with a median 11–15 years of experience. The headline numbers are staggering: 95% use AI tools at least

weekly. 75% use them for half or more of their work. 56% use them for 70% or more. This is no longer adoption; it's saturation.

Claude Code's eight-month dominance

The survey's most striking finding is Claude Code's trajectory. Launched in May 2025, it overtook both GitHub Copilot and Cursor to become the #1 most-used AI coding tool in under eight months. It commands 46% of "most loved" mentions — 57% when you include engineers citing Claude models generally. This happened in a market where GitHub Copilot has 4.7 million paid subscribers and is used by 90% of the Fortune 100.

The split between enterprise and individual choice tells the real story. At companies with 10,000+ employees, Copilot dominates at 56% usage — driven by Microsoft procurement bundling and enterprise licensing agreements. At small companies, where engineers pick their own tools, Claude Code leads at 75%. The market is bifurcated: what procurement buys vs. what engineers choose.

> *Staff+ engineers are the heaviest agent users at 63.5%, and agent users are twice as excited about AI as non-users. But only 29% of developers trust AI-generated code. The enthusiasm gap between power users and everyone else is widening.*
>
> — Data from The Pragmatic Engineer AI Tooling Survey, March 2026

OpenAI Codex: the enterprise counter-move

OpenAI isn't ceding the agentic coding market. Codex tripled to 1.6 million weekly active users following the GPT-5.3 Codex launch in early February, with over a million desktop app downloads and 5x token usage growth. Enterprise customers include Cisco, Nvidia, Ramp, and Rakuten. OpenAI is explicitly positioning Codex as "the standard agent" — expanding beyond coding into broader enterprise domains. Alongside this, Codex Security launched on March 7: an AI security agent that scanned 1.2 million commits and identified 792 critical and 10,561 high-severity findings. It analyses repositories, identifies vulnerabilities, validates them in a sandbox, and proposes fixes.

Cursor's event-driven bet

Cursor launched Automations on March 5 — a fundamentally different approach. Instead of agents you invoke, these are always-on cloud agents triggered by events: a Slack message, a GitHub push, a PagerDuty alert, a webhook. Hundreds of automations per hour, running security reviews on every commit, generating PR summaries, triaging incidents. The marketplace expanded to 30+ plugins (Atlassian, Datadog, GitLab, Glean, Hugging Face, PlanetScale), and Cursor is now available in JetBrains IDEs via the Cursor Agent Protocol. This isn't competing with Claude Code or Copilot directly — it's building the automation layer on top.

GitHub Copilot: the incumbent adapts

GitHub isn't standing still. Copilot's agentic capabilities are now GA in JetBrains with custom agents, sub-agents, and plan agents. Support for AGENTS.md and CLAUDE.md instruction files signals an acknowledgment that developers work with multiple AI tools. The coding agent gained self-review, security scanning, and cross-agent memory. But the economics of bundled frontier model access are starting to bite: GitHub removed premium model selection (Claude Opus, Sonnet, GPT-5.4) from the free student plan, affecting around 2 million students. The community response — 1,836 downvotes and 818 comments — was sharp.

Sources: The Pragmatic Engineer, Fortune, The Hacker News, TechCrunch, GitHub Blog

BANKING & FINTECH

Revolut Gets Its Full UK Banking Licence

13M	$75B	£3B	£120K
UK CUSTOMERS	VALUATION	UK INVESTMENT PLANNED	FSCS PROTECTION

After a four-year application process and 18 months in mobilisation, Revolut Bank UK Ltd received PRA approval on March 11 to operate as a fully licensed bank. This is the moment the fintech story changes from "disruptors nibbling at the edges" to "a direct competitor with regulatory parity." Eligible deposits are now FSCS-protected up to £120,000 — the same protection customers get at Barclays, HSBC, or Lloyds.

Current accounts are rolling out gradually to new customers. Existing customers remain on the e-money entity for now, with migration happening over several months. Credit products and lending are coming next — the capabilities that have been locked behind full banking authorisation and the real revenue engine for any bank.

The competitive reality

Revolut now stands shoulder-to-shoulder with high street banks on regulatory credibility, but it brings a fundamentally different operating model. 13 million UK customers, 70 million globally, built on a digital-first, engineering-led platform. They've committed £3 billion specifically to UK investment and plan to create 1,000 high-skilled jobs. With a US OCC charter application pending — part of a wave of 14 de novo applications in 2025, the most in four years — they're building a platform that directly competes for the same customers traditional banks serve, across jurisdictions.

The IPO question

Revolut is reportedly exploring a dual London/New York IPO at its $75 billion valuation, which would make it one of the largest fintech listings in history. The Bank of England governor reportedly blocked a chancellor-arranged meeting between Revolut and regulators over concerns about

political interference in the licensing process — a signal of how seriously the establishment is taking this, and how politically charged fintech regulation has become.

For anyone building payment infrastructure, loyalty programmes, or customer-facing banking technology in the UK: this is a competitor with 13 million customers that just gained the full regulatory platform to offer lending, credit, and FSCS-protected banking services — powered entirely by an engineering-led operating model that traditional banks cannot replicate without rebuilding from the ground up.

Sources: Bloomberg, CNBC, PYMNTS, Tech.eu, City AM

GENERAL SOFTWARE ENGINEERING

AI Is Generating Code Faster Than Teams Can Ship It

59%	70.8%	72 min	5%
THROUGHPUT INCREASE	SUCCESS RATE (5YR LOW)	RECOVERY TIME	TEAMS SHIPPING AT AI SPEED

CircleCI's 2026 State of Software Delivery report reveals a paradox that should trouble every engineering leader. The analysis of 28 million CI workflows shows the largest increase in development activity since 2019 — average throughput up 59% year-over-year. But look at the main branch and the picture inverts: delivery throughput for the median team is *down* 7%. Teams are writing more code than ever and shipping it more slowly.

Main branch success rates have fallen to 70.8%, the lowest in five-plus years. The benchmark is 90%. Recovery times are up 13% to 72 minutes. Put those numbers together at scale: a team with 500 daily changes at a 70% success rate loses the equivalent of 12 full-time engineers to recovery work alone. That's not a process problem; it's a structural failure.

The gap is not evenly distributed

The top 5% of teams saw throughput jump 97% — they figured out how to ship at AI speed. The median team got a 4% improvement. The bottom quartile got nothing. Feature branch throughput for the median team is up 15%, which means engineers are producing more work in progress, but the pipeline from "code written" to "code deployed" has a bottleneck. Mid-sized companies (21–50 employees) are stuck in what CircleCI calls the "messy middle": lowest throughput, recovery times averaging three hours.

The security dimension

Snyk's data adds a second layer to the problem: 48% of AI-generated code contains security vulnerabilities, and the time-to-exploit is shrinking — expected to accelerate by 50% by 2027. Sonar's State of Code survey puts the scale in perspective: 42% of code is now AI-generated or assisted, predicted to exceed half by 2027. If nearly half the code is AI-generated and nearly half of that is insecure, the maths on your security posture just changed fundamentally.

> *Only 1 in 20 teams have figured out how to ship at AI speed. The other 19 are generating more code, creating more work in progress, and watching their delivery infrastructure buckle under the load.*
>
> — CircleCI 2026 State of Software Delivery

The implication is clear: AI tooling has solved the code generation problem. It has not solved — and may be actively worsening — the code delivery problem. The teams that will win are the ones investing as heavily in CI/CD reliability, automated testing, and security scanning as they are in AI coding tools. Right now, that's about 1 in 20.

Sources: CircleCI, Snyk, Sonar

BANKING & FINTECH

Open Banking's Landmark Year and the Push to Open Finance

£43B	£8.3B	17.5M+	53%
POTENTIAL ANNUAL VALUE	CUMULATIVE BENEFIT	USER CONNECTIONS	YOY PAYMENT GROWTH

An EY analysis commissioned by Open Banking Limited puts the potential annual value of Open Banking at £43 billion for the UK economy at full maturity. That number is speculative at the upper end, but the trajectory underneath it is not: £8.3 billion in cumulative economic benefit already generated, with annual benefit expected to reach £7.4 billion within five years. 17.5 million user connections, 16 million active users, and 53% year-over-year growth in payment volumes.

The regulatory push

2026 is the year the regulatory infrastructure catches up with the technology. HM Treasury is expected to grant the FCA new powers to set open banking rules this year. The FCA is publishing its roadmap for Open Finance this month — extending data sharing beyond bank accounts to credit, mortgages, pensions, and insurance. Variable Recurring Payments are live with transactions flowing

in Q1 2026. The Future Entity — the successor to Open Banking Limited — is being established to provide governance for the next phase.

Beyond banking

The Open Finance agenda is where this gets architecturally interesting. When data sharing extends from current accounts to mortgages, pensions, and insurance, the API surface area expands enormously. Every financial institution becomes a potential data provider and consumer. The FCA's nine sector reports replacing traditional portfolio letters signal a more structured, outcome-oriented approach to supervision — and the digital pound blueprint, expected from HM Treasury and the Bank of England this year, could reshape the settlement layer entirely.

For engineering teams working in payments, lending, or financial data: the regulatory framework is about to demand more from your APIs, your security posture, and your data governance. The organisations that invested in clean, well-documented Open Banking interfaces in the first wave will have a structural advantage in the Open Finance expansion. Those that bolted on compliance at the last minute will face the same scramble again, at larger scale.

Sources: FinTech Global, Open Banking Limited, FCA

OTHER NEWS THIS MONTH

In Brief

AWS ME drone strike recovery continues (Day 16+) — physical repair of UAE and Bahrain facilities ongoing. No full restoration date confirmed. AWS still recommends migrating ME workloads to alternate regions. *The Register, CNBC*

Microsoft integrating Claude into Copilot Wave 3 — Anthropic's Cowork technology powering Agent 365, launching May. Claude Opus 4.6 live in Foundry. UK/EU tenants have Claude disabled by default. *eWeek, Microsoft Learn*

Anthropic launched $100M Claude Partner Network — enterprise training, Claude Certified Architect certification, co-marketing. Claude now available on all three major cloud providers. *Anthropic*

Node.js security releases scheduled March 24 — 2 high, 5 medium, 2 low severity issues across Node 25.x/24.x/22.x/20.x. *nodejs.org*

Cloudflare 2026 Threat Report — inaugural report found CVE-2026-22813 (9.4 CVSS) in markdown rendering, plus weaponisation of trusted cloud tooling as an emerging attack pattern. *Cloudflare Blog*

Monzo published representation learning for fraud detection — engineering blog details how learned embeddings improve fraud model accuracy beyond hand-crafted features. *Monzo Blog*

PLATFORM UPDATES

At a Glance

PLATFORM	WHAT CHANGED
.NET	Security servicing for .NET 10/9/8 (March 10). CVE-2026-26130 Security Feature Bypass across all three. Emergency .NET 10.0.5 out-of-band fix for macOS debugging (March 12).
Node.js	24.14.0 LTS released (March 11). Node 20 EOL April 30 — migration should be underway. New release schedule: one major per year from Node 27. Security releases March 24.
GitHub	Copilot agentic capabilities GA in JetBrains. AGENTS.md/CLAUDE.md instruction file support. Coding agent: self-review, security scanning, cross-agent memory. 30+ marketplace plugins.
Cursor	Automations launched (March 5). JetBrains ACP integration. 30+ marketplace plugins. Security automation templates for always-on vulnerability scanning.
Patch Tue	82 CVEs patched by Microsoft (8 critical). Notable: CVE-2026-26144 Excel information disclosure flaw could exfiltrate data via Copilot.

THE LEADERSHIP READ

Why Would Anyone Be an Engineering Manager in 2026?

Will Larson — CTO at Imprint and one of the most widely read voices on engineering management — published a piece on LeadDev in January that has been quietly circulating among engineering leaders ever since. Its core argument is uncomfortable: the management model that defined the last decade is dead, and nobody has agreed on what replaces it.

The hypergrowth model is over

The engineering management playbook of the 2010s was defined by growth. Growing headcount meant expanding scope. Expanding scope meant clear career ladders. There was a legible path: IC to team lead to manager to director, with each step bringing more people, more budget, more influence. That model assumed a world of perpetual hiring, rising valuations, and organisational expansion. That world no longer exists for most engineering organisations.

Now the reality is flat or shrinking teams, AI changing what engineering work looks like, and intense pressure to "stay technical." Managers are being told, in effect, that they spent the last decade doing the wrong things — that the time invested in hiring processes, team rituals, and stakeholder management was less valuable than maintaining deep technical fluency. The industry wants managers to be simultaneously more technical *and* better people managers, without acknowledging that these demands pull in opposite directions.

> *The industry spent a decade telling engineering managers to focus on people, process, and organisational design. Now it's telling them that was wrong. The dissonance isn't just demoralising — it's structurally unsound.*
>
> — Adapted from Will Larson, LeadDev, January 2026

The AI dimension makes it worse

Connect Larson's argument to the data in this issue. CircleCI shows delivery infrastructure buckling — main branch success at a five-year low, recovery times climbing, only 5% of teams shipping at AI speed. Traditional metrics are breaking down. Lines of code, PRs merged, velocity points — none of these measure what matters when an engineer with three AI agents running in parallel produces more meaningful output in two hours than a traditional workflow generates in a day. Managers need new metrics, new ways of measuring output, new models for team structure — and they're being asked to figure this out while simultaneously being told they should be writing more code themselves.

Why it still matters

Larson doesn't argue that engineering management is pointless. He argues that the *old* model is, and that people who want to manage engineering teams in 2026 need to engage with an entirely new set of problems: how to measure productivity when AI decouples effort from output, how to structure teams when headcount is flat, how to maintain technical credibility while doing the organisational work that still needs doing. The managers who thrive will be the ones who build a new playbook rather than mourning the old one.

Read the full article: LeadDev — Why would anyone be an engineering manager in 2026? → · *Also: lethain.com — Good engineering management is a fad*

OUTSIDE IN

Can AI Agents Build Real Stripe Integrations? Stripe Built a Benchmark to Find Out

Most AI coding benchmarks test isolated problems: write a function, fix a bug, complete a snippet. Stripe's engineering team did something different. They built production-realistic evaluation environments — 11 diverse setups, deliberately biased toward difficulty — and set frontier AI agents loose on real Stripe integration tasks. The results, published on March 2, are one of the first rigorous assessments of what AI agents can actually do when the task looks like real engineering work.

The results surprised the researchers

Claude Opus 4.5 scored 92% on average across four full-stack API integration tasks. GPT-5.2 scored 73% across two "gym" problem sets. All models worked productively for long durations — the best runs averaged 63 turns. The agents navigated UIs, debugged live issues, handled underdocumented behaviour, and self-verified through test purchases. They used a goose-based harness with an MCP server providing terminal access, browser control, and Stripe-specific search.

Where they struggled

Ambiguity was the biggest failure mode. When requirements were unclear, agents made assumptions rather than asking for clarification — and those assumptions were often wrong. Browser interactions remained unreliable. Most notably, some agents accepted error states as successful completions: the code ran, a response came back, and the agent moved on without checking whether the response actually indicated success. This is exactly the pattern that creates subtle production bugs humans catch through experience and contextual awareness.

Why this matters

Stripe open-sourced the benchmark as part of their AI toolkit. The significance isn't the scores — it's the methodology. Production-realistic evaluation environments, with real APIs, real UIs, and real ambiguity, produce fundamentally different results than clean-room coding challenges. A 92% score on a full-stack Stripe integration is a different claim than a 92% score on LeetCode. For engineering teams evaluating where to deploy AI agents in their own workflows, this is the standard of evidence to look for: not "can it write code?" but "can it ship working integrations against real APIs with real edge cases?"

Read the full post: stripe.com/blog/can-ai-agents-build-real-stripe-integrations →

ON THE RADAR

Coming Up

March 24	Node.js security releases — 2 high, 5 medium, 2 low severity across 25.x/24.x/22.x/20.x
April 19	Claude Haiku 3 model retirement — migrate to Haiku 4.5

April 22	AWS Summit London — expect Bedrock and Trainium announcements
April 29	Zero Day Con, Dublin — Ireland's largest cybersecurity conference
April 30	Node.js 20 end of life
May	Microsoft Copilot Wave 3 / Agent 365 launch (includes Claude integration)

RECOMMENDED READING

This Month's Links

AI Tooling for Software Engineers in 2026	The Pragmatic Engineer
5 Key Takeaways from the 2026 State of Software Delivery	CircleCI
Can AI Agents Build Real Stripe Integrations?	Stripe Engineering
OpenAI Codex Security Scanned 1.2 Million Commits	The Hacker News
Cursor Is Rolling Out a New Kind of Agentic Coding Tool	TechCrunch
The Engineering Leadership Blind Spot of 2026	Waydev
Why Would Anyone Be an Engineering Manager in 2026?	LeadDev

Engineering Digest

February 17 – March 17, 2026 · Inaugural Edition · Issue #0

Next edition: March 20, 2026