Engineering *Digest*

February 18 – March 18, 2026 | Issue #0

MONTHLY RETROSPECTIVE

IN THIS ISSUE

When Drones Hit the Cloud: The AWS Middle East Strikes
The AI Delivery Bottleneck
The Leadership Read: Agentic Engineering
The AI Coding Landscape Shifts
Revolut Gets Its Full UK Banking Licence
Other News • Platform Updates • On the Radar

FOREWORD

Engineers adopted AI coding tools faster than anyone predicted — 95% use them weekly now. But the delivery data tells a different story: main-branch success rates just hit a five-year low. The month's most important question isn't whether your teams are using AI. It's whether all that generated code is actually reaching production.

⚠ ONGOING P0 — AWS MIDDLE EAST RECOVERY

Day 18 · ~40% services restored. Physical repair of UAE and Bahrain facilities ongoing. No full restoration date confirmed. AWS continues to recommend migrating ME workloads to alternate regions. *See lead article and Deep Dive DD-1 for full analysis.*

THE MONTH AT A GLANCE

Eight stories in sixty seconds

⚠ ACTION REQUIRED

AWS **Iranian drone strikes hit 3 AWS facilities in UAE and Bahrain** (Mar 1) — 109+ services offline, 2/3 UAE AZs down. 18+ days without full restoration. AWS recommends migrating ME workloads. *See lead article.*

Security **Chrome zero-days patched** (Mar 13) — exploited in the wild. Update immediately. *The Hacker News*

Node **Node.js 20 EOL April 30** — 24.14.0 LTS released (Mar 11). Migration should be underway. *nodejs.org*

WORTH KNOWING

AI **Claude Code is the #1 AI coding tool** — Pragmatic Engineer survey: 95% weekly AI usage, Claude Code overtook Copilot and Cursor in eight months. *See Quick Take 1.*

Banking **Revolut secured its full UK banking licence** (Mar 11) — FSCS-protected deposits, 13M customers. Credit products and dual London/NY IPO ahead. *See Article 4.*

DevEx **CircleCI: throughput up 59%, but delivery is slowing** — main branch success at 70.8%, a five-year low. The gap between generating code and shipping it is widening. *See Quick Take 2.*

WATCHING

AI **Microsoft Copilot Wave 3 launches May with Claude integration** — Agent 365 at $15/user, E7 Frontier Suite at $99/user. *Windows Central*

Regulati **EU AI Act high-risk obligations apply August 2, 2026** — organisations deploying AI in hiring, credit scoring, or critical infrastructure must comply. Planning window is narrowing.

PRIORITY 0 • CLOUD & INFRASTRUCTURE

When Drones Hit the Cloud: The AWS Middle East Strikes

109+	3	2 of 3	18+
SERVICES AFFECTED	FACILITIES HIT	UAE AZS OFFLINE	DAYS RECOVERY

On March 1, Iranian drone strikes hit two AWS data centres in the UAE and damaged a third in Bahrain. Two of three availability zones went offline — not degraded, offline. 109+ services became unavailable. The physical damage was severe enough that AWS

Context: What is multi-AZ?

AWS regions contain multiple Availability Zones — physically separate data centres within the same geographic area. Multi-AZ

initially described "objects" causing a "localized power issue," then confirmed the strikes 36 hours later.

Banking impact was immediate. ADCB mobile banking went down. Emirates NBD was disrupted. Payment platforms Alaan and Hubpay stopped processing. The effects cascaded beyond the Middle East: US-EAST-1 experienced degraded performance, Snowflake reported issues, and Claude AI saw availability problems. As of mid-March, there is no full restoration date. AWS's recommendation: migrate workloads to alternate regions.

The data residency paradox

UAE and Bahrain Central Bank mandates require financial data to stay in-region. The region's infrastructure was physically destroyed. Failing over resolves the outage but creates a legal problem — you can't comply with data residency requirements when the region doesn't exist in a functional sense. The Dubai Financial Services Authority issued emergency guidance for temporary cross-region processing, a regulatory improvisation nobody had planned for.

Multi-AZ redundancy — the standard resilience playbook — assumed availability zones fail independently through technical faults, not simultaneously through military action. That assumption is broken. The conversation has shifted from "how many AZs?" to "how many *regions*?" — and what happens to your compliance posture when you can't use the one the regulator requires.

Sources: CNBC, The Register, Reuters, Data Center Knowledge, Marketplace

deployment spreads workloads across 2–3 AZs so that if one fails, the others keep running. The design assumption: AZs fail independently through hardware faults, not simultaneously through kinetic military action.

Cascading impact

The strikes didn't stay contained. US-EAST-1 experienced degraded performance. Snowflake reported service issues. Claude AI saw availability problems. $300 billion in Gulf AI and data centre investments are now at elevated risk.

This is the first major military strike on hyperscale cloud infrastructure. Data centres are now military targets.

— CSIS; The Information, Semafor (March 2026)

QUICK TAKES

Two stories, one thread

AI & TOOLING

DEVELOPER EXPERIENCE & PRODUCTIVITY

The AI Coding Landscape Shifts

The Pragmatic Engineer's March survey of 906 engineers confirmed what most delivery managers already suspect: AI coding tools are now baseline infrastructure. 95% use them weekly, 75% for half or more of their work, and 55% use agents regularly. Claude Code is the most-used tool — overtaking Copilot and Cursor in eight months since launch, with 46% "most loved" mentions.

The cost and procurement picture is shifting fast. Anthropic's COBOL modernisation capability sent IBM shares down 13%. Microsoft is integrating Claude into Copilot Wave 3, launching Agent 365 in May at $15/user and the E7 Frontier Suite at $99/user. Claude is now available on all three major clouds. The $100M Claude Partner Network launched March 12 with enterprise certification and training.

For delivery managers, the question is no longer "should teams use AI tools?" — it's which agent platform your procurement locks you into, and whether the tools your engineers actually prefer are the ones your enterprise agreement covers.

Sources: Pragmatic Engineer, CNBC, Windows Central

The AI Delivery Bottleneck

70.8% MAIN BRANCH SUCCESS RATE

-7% MAIN BRANCH THROUGHPUT

+59% AVERAGE THROUGHPUT

72 min RECOVERY TIME (+13%)

CircleCI's 2026 report — 28 million workflows analysed — reveals the other side of the AI adoption story in Quick Take 1. Average throughput jumped 59%, but on the main branch where code actually ships, throughput for the median team *fell* 7%. Success rate hit 70.8% — a five-year low against a 90% benchmark.

Read these two quick takes together: if 95% of engineers are generating code with AI and main-branch delivery is simultaneously falling, the bottleneck has moved. Writing code is no longer the constraint. Review, validation, and integration is where everything piles up. The top 5% of teams saw 97% throughput gains; the median got 4%. AI is amplifying existing delivery strengths, not distributing them.

Sources: CircleCI, Full Report

BANKING & FINTECH

Revolut Gets Its Full UK Banking Licence

13M	$75B	£3B	£120K

UK CUSTOMERS	VALUATION	UK INVESTMENT PLANNED	FSCS PROTECTION

After a four-year application process and 18 months in mobilisation, the PRA lifted restrictions on March 11 and Revolut Bank UK Ltd became a fully licensed bank. Eligible deposits are now FSCS-protected up to £120,000 — the same safety net customers get at Barclays or HSBC. This is the moment the fintech narrative shifts from "challenger nibbling at margins" to "direct competitor with regulatory parity."

The licence unlocks lending and credit products — the real revenue engines of banking. Revolut has committed £3 billion to UK investment and 1,000 high-skilled jobs. They're exploring a dual London/New York IPO at their $75 billion valuation. The Bank of England governor reportedly blocked a chancellor-arranged meeting over political interference concerns — a signal of how seriously the establishment takes this.

For anyone building payment infrastructure or customer-facing banking technology in the UK: this is a competitor with 13 million customers that just gained the full platform to offer lending, credit, and protected deposits — powered by an engineering-led operating model traditional banks cannot replicate without rebuilding from scratch.

Sources: Reuters, CNBC, Bloomberg

OTHER NEWS THIS MONTH

In Brief

PMI and Agile Alliance released the Manifesto for Enterprise Agility (Mar 3) — first joint manifesto, extending agile from team delivery to enterprise governance, funding, and leadership. *PMI, Agile Alliance*

Anthropic discovered 24,000+ fraudulent accounts from Chinese AI labs — DeepSeek, Moonshot, and MiniMax systematically scraping Claude capabilities. *Mean CEO Blog*

MCP security concerns mounting — RSA Conference 2026 submissions heavily focused on Model Context Protocol risks. Agentic AI attack surface is becoming a first-class security concern. *CIO*

UK open banking heading toward open finance — FCA expects VRP live transactions in Q1, statutory instrument for new regulatory framework expected 2026. *Open Banking UK*

March Patch Tuesday: 79 CVEs, 2 zero-days — SQL Server elevation of privilege and .NET denial of service both exploited. Patch now. *BleepingComputer*

PLATFORM UPDATES

At a Glance

PLATFORM	WHAT CHANGED
.NET	March security patches across .NET 10/9/8. Emergency .NET 10.0.5 out-of-band fix for macOS debugging (Mar 12).
Node.js	24.14.0 LTS (Mar 11). Node 20 EOL April 30 — migration should be underway. New release schedule: one major per year from Node 27.
GitHub	Copilot agentic capabilities GA in JetBrains. VS 2026 18.3.0 with Copilot memories. AGENTS.md instruction file support.
AWS	S3 turned 20. Route 53 Global Resolver GA. Stateful MCP in Bedrock AgentCore. ME region recovery ongoing.

THE LEADERSHIP READ

Agentic Engineering in Practice: What the Pragmatic Summit Told Us

Simon Willison — creator of Datasette and one of the most respected voices in AI-assisted development — sat down for a fireside chat at Gergely Orosz's Pragmatic Summit in San Francisco. Published March 14. It's the most useful thing released this month for anyone managing an engineering team that uses AI tools.

The adoption curve is real — and your people are at different points on it

Willison described distinct phases: ChatGPT for occasional questions, then AI agents writing bits of code, then agents producing more code than the human. The newest phase — where Willison operates now — is where "you don't read the code." The agent writes it, runs it, tests it, and you evaluate the output. Understanding where each person on your team sits on this curve matters, because they need fundamentally different support at each stage.

TDD becomes essential, not optional

"Every single coding session I start with an agent, I start by saying here's how to run the test." When AI writes the code, tests become the specification language — the mechanism through which you verify the agent did what you intended. Teams without strong test coverage are flying blind with agents.

Mental exhaustion is the new bottleneck

> *After two hours of that I'm done for the day. I'm mentally exhausted. People worry about skill atrophy and being lazy. I think this is the opposite of that. You have to operate firing on all cylinders if you're going to keep your trio or quadruple of agents busy solving all these different problems.*
>
> — Simon Willison, Pragmatic Summit, 2026

Running 3–4 agent sessions in parallel is technically possible but cognitively draining. Two hours of high-intensity agent orchestration can outproduce a full day of traditional development. The engineers using AI most effectively may look like they're working less. Traditional productivity metrics are becoming meaningless.

Maintain a "retry list"

Keep track of tasks that AI failed at and re-test every six months as models improve. Willison describes regularly being surprised when something impossible three months ago now works perfectly. The teams that systematically re-test their assumptions will find capabilities others miss.

What this means for management

Productivity measurement needs rethinking. Team practices must account for cognitive intensity. Tooling investment is becoming as important as headcount. The engineers leaning furthest into agents aren't the ones who need less attention — they're navigating the steepest learning curve and operating at the highest cognitive load.

Read the full post: simonwillison.net/2026/Mar/14/pragmatic-summit/ →

ON THE RADAR

What's Coming

DEADLINES & REQUIRED ACTIONS

April 19	Claude Haiku 3 model retirement — migrate to Haiku 4.5
April 30	Node.js 20 end of life — migration to v24 recommended
August 2	EU AI Act — high-risk AI obligations take effect

EVENTS & CONFERENCES

April 22	AWS Summit London — expect Bedrock, Trainium, and ME recovery announcements
April 29	Zero Day Con, Dublin — Ireland's largest cybersecurity conference
May 1	Microsoft Copilot Wave 3 / Agent 365 launch

RECOMMENDED READING

This Month's Links

AI Tooling for Software Engineers in 2026	The Pragmatic Engineer
5 Key Takeaways from the 2026 State of Software Delivery	CircleCI
My fireside chat about agentic engineering	simonwillison.net
The Future of Software Engineering with AI	The Pragmatic Engineer
Why Model Context Protocol is suddenly on every executive agenda	CIO
Manifesto for Enterprise Agility	PMI / Agile Alliance

Engineering Digest

February 18 – March 18, 2026 · Monthly Retrospective · Issue #0

Mar 11–18 · Weekly · Mar 18 · Deep Dive

Next edition: March 25, 2026