Engineering *Digest*

March 13–20, 2026 | Weekly Edition | Issue #1

THIS WEEK'S BRIEFING

IN THIS ISSUE

NVIDIA GTC 2026: The Trillion-Dollar AI Infrastructure Bet
AWS Middle East Recovery: Day 18 Update
Outside In: How Monzo Ships Hundreds of Changes a Day
The Agent Workflow Wars: Skills, Manus, and What Ships Next
The Leadership Read: How AI Is Changing Staff+ Engineering
Platform Updates, Other News & On the Radar

FOREWORD

Jensen Huang stood on stage this week and outlined a trillion dollars in AI infrastructure orders. That number is staggering, but the more interesting signal from GTC came from the software side: NVIDIA is building an open-source agent platform, OpenAI is shipping reusable enterprise workflows, and Meta's Manus just landed on the desktop. The silicon is settled. The fight for the agent layer on top of it is just starting, and that is the layer your teams will actually have to work with.

ONGOING STORY TRACKER

AWS ME Recovery — Day 18

AWS Middle East (UAE & Bahrain): Physical repair of data centre facilities ongoing. No full restoration date confirmed. AWS continues to recommend migrating ME workloads to alternate regions. *AWS Health Dashboard*

THE WEEK AT A GLANCE

Seven stories in sixty seconds

⚠ ACTION REQUIRED

Node.js

Node.js security releases scheduled for March 24 — affecting v25, v24, v22, and v20. Patch promptly. Node 20 EOL April 30. *nodejs.org*

.NET **March Patch Tuesday: 84 CVEs, 8 critical** — includes .NET DoS (CVE-2026-26127), ASP.NET Core DoS, Office RCE. Patch .NET 10.0.4/9.0.14/8.0.25. *Zero Day Initiative*

WORTH KNOWING

AI **NVIDIA GTC 2026: $1 trillion in Blackwell + Vera Rubin orders** — Vera Rubin shipping H2 2026 with 10x cheaper inference. NemoClaw open-source agent platform announced. DGX Spark desktop AI.

AI **OpenAI launches ChatGPT Skills** (March 13) — reusable enterprise workflows for Business/Enterprise. Follows Agent Skills open standard. GPT-5.4 mini rolling out.

AI **Meta's Manus launches desktop app** (March 18) — "My Computer" feature gives the AI agent direct access to local files and applications. *CNBC*

WATCHING

AWS **AWS Middle East recovery Day 18** — physical repair of UAE and Bahrain facilities ongoing. No full restoration date. Migration guidance unchanged.

GitHub **Copilot coding agent now starts work 50% faster** (March 19) — shorter startup times for issue-to-PR workflows and iteration cycles. *GitHub Blog*

AI TOOLING & INFRASTRUCTURE

NVIDIA GTC 2026: The Trillion-Dollar AI Infrastructure Bet

NVIDIA's annual GTC developer conference (March 16–19) delivered a week of announcements that collectively signal where enterprise AI infrastructure is heading. The headline: Jensen Huang disclosed that Blackwell and Vera Rubin orders already total over $1 trillion through 2027. For engineering leaders, the numbers matter less than what they imply — AI compute is being procured at a scale that guarantees the tooling, platforms, and agent ecosystems built on top of it

NemoClaw: What to Know

Open source — Apache 2.0 expected

Hardware-agnostic — runs on any infrastructure

Enterprise-focused — security and privacy controls built in

Pitched to: Salesforce, Cisco, Google, Adobe, CrowdStrike

Complements: AI-Q agent blueprint also

will dominate enterprise software for the rest of the decade.

$1T+	10x	4x
ORDERS THROUGH 2027	CHEAPER INFERENCE (RUBIN VS BLACKWELL)	FEWER GPUS NEEDED (RUBIN VS BLACKWELL)

Vera Rubin: the next generation ships this year

The Vera Rubin platform is in full production and will ship in H2 2026. NVIDIA claims 4x fewer GPUs required to train the same models compared to Blackwell, and a 10x reduction in inference token costs. That last figure is the one that matters for teams evaluating AI spend: if inference costs drop by an order of magnitude, the economics of running AI agents in production shift from "experimental budget" to "cost of doing business." The Feynman architecture, previewed for 2028, promises another generational leap.

NemoClaw: NVIDIA enters the agent platform market

The most strategically significant announcement wasn't about silicon — it was software. NemoClaw is NVIDIA's new open-source enterprise AI agent platform, designed to let organisations deploy autonomous AI agents with built-in security and privacy controls. The platform is hardware-agnostic, meaning it runs on any infrastructure, not just NVIDIA's. NVIDIA has been pitching it to Salesforce, Cisco, Google, Adobe, and CrowdStrike.

This is NVIDIA doing what it did with CUDA two decades ago: giving away the software layer to lock in the ecosystem. If NemoClaw becomes the default agent deployment platform, NVIDIA captures the orchestration layer regardless of

announced at GTC

Also Announced at GTC

DGX Spark — desktop AI supercomputer for developers

DLSS 5 — next-gen rendering for gaming

Feynman — 2028 architecture preview, the generation after Rubin

Physical AI — robotics and simulation announcements

whose hardware runs it. For engineering teams, it means another open-source agent framework to evaluate — but this one comes backed by the company that manufactures the chips underneath it.

Sources: CNBC, NVIDIA Investor Relations, NVIDIA Blog, Wired, Forbes (March 9–19, 2026)

QUICK TAKES

Two stories worth five minutes

AI TOOLING

The Agent Workflow Wars Are Crystallising

Three announcements in five days painted a clear picture of where enterprise AI agents are heading. OpenAI launched **ChatGPT Skills** on March 13 — reusable, shareable workflows for Business and Enterprise plans that let teams define a process once and have ChatGPT follow it consistently. Skills follow the Agent Skills open standard, making them portable across products including Codex and the API.

On March 18, Meta's **Manus shipped a desktop app** with "My Computer" — giving the agent direct access to local files and applications. The same day, OpenAI rolled out **GPT-5.4 mini** as a rate-limit fallback, continuing the pattern of making frontier-class reasoning cheaper and more accessible.

The convergence is clear: agents are moving from chat interfaces to persistent workflow tools that operate on your files, follow your processes, and integrate with your existing stack. The teams that define their workflows as Skills now will have a head start when

DEVELOPER EXPERIENCE

The Largest AI Adoption Study Yet Says the Gap Is Widening

Jellyfish published its AI Engineering Trends benchmark this week — the most comprehensive quantitative study of AI's impact on software engineering to date, drawn from **700+ companies, 200,000 engineers, and 20 million pull requests**. The headline: over the past three months, companies in the top quartile of AI adoption have seen **2x the PR throughput** of low adopters. 64% of surveyed companies now generate a majority of their code with AI assistance.

CONNECT THE DOTS

Jellyfish's partnership with Harvard Economics analysed 100,000 engineers across 500 companies and found AI is making coding faster without measurable quality degradation. But the Harvard researchers also found something striking: productivity improvements aren't yet translating into business outcomes. Engineers are coding faster, but organisations aren't necessarily delivering more value to customers. The bottleneck has moved downstream.

these tools become the default mode of work.

Sources: OpenAI, CNBC, OpenAI Release Notes

Autonomous agent activity — PRs generated entirely by AI agents — is still low overall, but growing exponentially. The data also shows that aggressive AI adopters are pulling away from the pack, creating a measurable maturity gap between organisations that have invested in AI integration and those still experimenting. TaskRabbit, using the Augment Code integration tracked by Jellyfish, reported a 50% decrease in issue cycle time and 2x deployment rates.

Sources: Jellyfish, Jellyfish AI Engineering Trends, Jellyfish / Harvard Economics

OTHER NEWS THIS WEEK

In Brief

Bain & Co: challenger banks to own a third of UK banking market by 2030 — traditional banks have slipped from 95% to 80% of addressable revenue. Revolut's new UK banking licence accelerates the shift. *Chris Skinner / The Finanser*

Standard Chartered hires JP Morgan's Naveen Mallela as global head of payments — Kinexys co-head leaves after 11 years. Starts May 4, unifying collections, clearing, and payments. *FinTech Futures*

GitHub Copilot model picker incident (March 16) — Claude Sonnet 4.6 and Opus 4.6 disappeared from Copilot Pro. Resolved March 17 after 1,836 downvotes. *GitHub Community*

Node.js evolving release schedule — one major version per year from Node 27 (April 2027). Alpha phase introduced. Node 26 follows the existing model. *nodejs.org*

PLATFORM UPDATES

At a Glance

PLATFORM	WHAT CHANGED
.NET	Security servicing 10.0.4 / 9.0.14 / 8.0.25 (March 10). DoS vulnerability (CVE-2026-26127), ASP.NET Core DoS (CVE-2026-26130), privilege escalation (CVE-2026-26131). Emergency 10.0.5 out-of-band fix for macOS debugging.

PLATFORM	WHAT CHANGED
GitHub	Copilot coding agent 50% faster startup (March 19). Model picker incident resolved (March 17). Enterprise telemetry tracked separately.
Node.js	Security releases scheduled March 24 for v25/v24/v22/v20. Node 20 EOL April 30 — migrate to v24. New release schedule announced: one major per year from Node 27.
OpenAI	ChatGPT Skills beta for Business/Enterprise (March 13). GPT-5.4 mini rollout (March 18). Model picker simplified (March 17). GPT-5.1 models retired (March 11).

We are not freeing up more time in our calendar. We are taking on more work, and therefore there's higher pressure to go and ship these changes.

— Suhail Patel, Principal Engineer at Monzo, QCon London 2026

THE LEADERSHIP READ

How AI Is Changing the Staff+ Engineer Role

Maxime Najim — Distinguished Engineer at Target with two decades building systems at Netflix, Apple, and Amazon — published a detailed account this week of how AI coding agents are reshaping the work of senior technical leaders. The piece, published on LeadDev, is the most practical treatment of what the agentic shift actually looks like at staff+ level: less time writing implementations, more time defining what "correct" means, and a fundamentally different relationship between seniority and execution.

Execution is no longer the bottleneck — verification is

Najim describes being pulled into an unfamiliar domain and using an AI agent to read internal Confluence docs, trace data flows across 20+ repositories, and identify where documentation had drifted from reality. He used the output to draft a steel-thread proposal, iterated with principal engineers, and had alignment in 2–3 days. "This would normally take a month, maybe more." The new SDLC, he argues, is agentic and parallel: planning compresses from weeks to hours, agents explore multiple paths simultaneously, and human verification becomes the primary bottleneck.

From influence to activation

The most striking shift Najim describes is what he calls moving from "influence to activation." Staff+ engineers have always been expected to improve systems across organisational boundaries, but their main tool was persuasion — writing proposals, making the case, waiting for it to land on a

roadmap. Now, you point a coding agent at the problem, generate a working PR, and walk over to the owning team with a finished implementation instead of a backlog ticket. "The conversation shifts from 'can you prioritise this sometime next quarter?' to 'here's the solution, does this look right to you?'"

> *Your value stops being about how fast you can implement something and starts being about whether you can look at what got implemented and say 'yeah, I trust this.'*
>
> — Maxime Najim, Distinguished Engineer at Target, LeadDev (March 16, 2026)

The spec becomes the whole game

When the volume of AI-generated code outpaces what any person can read line by line, Najim argues, trying to review it all just makes you the bottleneck. The differentiator becomes spec quality — not a traditional requirements document, but a constraint system: what must not change, what stays true when things break, what's off-limits for compliance reasons, and what "done" means in production (monitoring, alerting, rollback plans, SLOs). The staff+ engineers who get comfortable in that space first, he concludes, will have an outsized impact on what their teams believe they can take on.

Read the full article: How AI is changing my work as a staff+ engineer — LeadDev →

OUTSIDE IN

Monzo: Move Fast, Don't Break Trust

At QCon London this week, Suhail Patel — principal engineer leading Monzo's platform group — described how the bank ships hundreds of changes to production every day across 3,000+ microservices. The talk, "Move Fast, Don't Break Trust," is one of the most detailed public descriptions of how a UK-regulated bank has built a developer platform that treats speed and compliance as the same problem, not competing priorities.

Every service follows the same structure via a bootstrap generator. The bank's "Backend 101" guide — mandatory for all engineers — has been encoded as LLM skills so Claude and Cursor generate code that conforms to Monzo's conventions by default, drawing on 3,000 repositories as training examples. LLMs now generate 10 to 20 candidate implementations within hours; non-engineering staff handle small bugs directly with AI rather than raising tickets.

But Patel pushed back on AI code review as the answer: "You don't need a fancy AI code review. Instead, spend the time writing really high-quality CI checks." The emphasis is on deterministic validation — incidents become automated checks, quality gates are mandatory, and bypassing them requires multi-party authorisation. Monzo's "tenancies" system provides lightweight isolated

staging environments that provision in minutes, with hundreds running simultaneously, so developers test against realistic infrastructure without blocking each other.

Patel's three principles: invest in composable, best-in-class tools to encode organisational processes; invest in rapid iteration capabilities; standardise on a small set of technology choices and keep improving those abstractions. Observability runs on Prometheus, Grafana, and OpenTelemetry, with alerts tracing back to specific changes and responsible teams.

Read the full report: QCon London 2026 — Shipping Constantly with Humans and Beyond — InfoQ →

ON THE RADAR

Coming Up

DEADLINES & REQUIRED ACTIONS

March 24	Node.js security releases — v25, v24, v22, v20. Patch promptly.
April 19	Claude Haiku 3 model retirement — migrate to Haiku 4.5
April 30	Node.js 20 end of life. AWS Lambda Phase 1: security patches stop for nodejs20.x

EVENTS & CONFERENCES

April 22	AWS Summit London — expect Bedrock and Trainium announcements
May	Microsoft Copilot Wave 3 / Agent 365 launch (includes Claude integration)

RECOMMENDED READING

This Week's Links

The Future of Software Engineering with AI: Six Predictions	The Pragmatic Engineer
AI Tooling for Software Engineers in 2026	The Pragmatic Engineer
NVIDIA GTC 2026: Keynote highlights	CNBC
Move Fast, Don't Break Trust: Shipping Constantly at Monzo	InfoQ
Evolving the Node.js Release Schedule	nodejs.org

Engineering Digest

March 13–20, 2026 · Weekly Edition · Issue #1

Previous issues: Mar 11–18 · Weekly · Deep Dive DD-1 · Feb 18 – Mar 18 · Monthly

Next edition: March 20–27, 2026