Engineering *Digest*

March 20–27, 2026 | Weekly Edition | Issue #2.5

THIS WEEK'S BRIEFING

1st

Starling becomes the first UK bank to deploy agentic AI in a customer-facing context — acting on accounts, not just answering questions

IN THIS ISSUE

Starling's Agentic AI Banking Assistant
GitHub Copilot Data Policy vs Cursor Self-Hosted Agents
Nebius us-central1 Post-Mortem
XBOW Embeds AI Pen-Testing in Microsoft Security Copilot
The Preceptorship Model for Junior Devs
Platform Updates, Other News & On the Radar

FOREWORD

Agentic AI stopped being a research concept this week and started touching real bank accounts. Starling's assistant doesn't just interpret — it acts, moving money, setting goals, and executing on behalf of customers. That shift from advisory to autonomous changes the engineering surface area dramatically: trust models, rollback strategies, and regulatory boundaries all need rethinking when the AI has write access to your customers' money.

PRIORITY 0 — ONGOING

AWS Middle East Region Disruption — Day 24+

Second drone-related disruption hit Bahrain on March 24 — the second time in a month. AWS continues advising migration to alternate regions with no recovery timeline. Reuters

THE WEEK AT A GLANCE

Five stories in sixty seconds

ACTION REQUIRED

Security **Node.js security releases patch 7 vulnerabilities across v20–v25.** Two high-severity flaws: remote process crash via TLS SNICallback and HTTP __proto__ header crash. Patch to v20.20.2, v22.22.2, v24.14.1, or v25.8.2 immediately. *nodejs.org*

WORTH KNOWING

Fintech **Revolut posts record £1.7B profit, rolling out UK bank accounts.** 68.3M customers globally after securing full UK banking licence. Lending products next — the competitive pressure on UK neobanks just intensified. *Reuters*

AI **OpenAI winds down Sora to redirect compute.** The AI video generator is being retired and the Disney collaboration cancelled. Shift from consumer to business focus — compute is being reallocated to the next major model. *NBC News*

WATCHING

DevEx **Cursor ships self-hosted cloud agents.** Code and tool execution stays in your own network. A direct counterpoint to GitHub's centralised data approach — worth watching for teams evaluating code assistant governance. *cursor.com*

Reg **FCA targeted support gateway opens — rules take effect April 6.** Banks and pension providers can now apply to offer tailored recommendations. Engineering teams at authorised firms should be preparing integration work. *FCA*

BANKING & FINTECH

Starling Launches the UK's First Agentic AI Banking Assistant

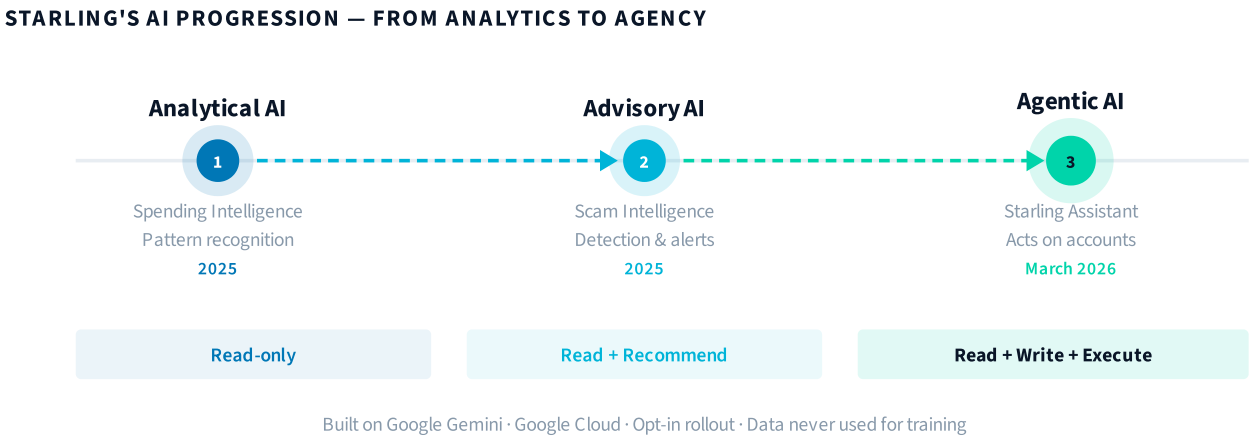

Data: Starling Bank

What Starling built

Starling's "Starling Assistant," launched March 20 for personal account holders on an opt-in basis, is a conversational banking interface that responds to voice and natural language prompts. It doesn't just answer questions — it acts. The assistant can set up personalised saving goals, organise bill payments, share spending insights, and provide tailored banking guidance, all executing directly on the customer's account. Built on Google Gemini running within Starling's Google Cloud environment, all customer data stays within that perimeter and is never used for model training.

Eight years of infrastructure, three phases of AI

The launch is the culmination of eight years of AI development at Starling. The progression is worth studying: Starling moved from analytical AI (Spending Intelligence — pattern recognition in spending data, launched 2025) to advisory AI (Scam Intelligence — detecting and alerting on marketplace scams) to agentic AI (taking actions on customer accounts). Each layer built on the infrastructure and trust model of the one before it. The Starling Assistant now surfaces both earlier tools alongside new agentic capabilities through a single conversational interface.

The competitive signal

This matters because it sets a benchmark. Starling is the first UK bank to give an AI write access to customer accounts in a production setting. The opt-in model and the progressive capability build suggest a deliberate approach to regulatory and customer trust. But the competitive clock has started. Revolut posted £1.7B in profit this week and is expanding its UK banking operations with 68.3M customers globally. The FCA's targeted support gateway opens April 6. For engineering teams at UK banks, the question is no longer whether to build agentic capabilities, but how quickly your architecture, trust model, and compliance framework can support them.

What engineering leaders should watch

The hardest problem here is not the AI model — it is the control plane around it. When an assistant has the ability to move money, set up recurring transfers, and modify account configurations, the

blast radius of a model error or prompt injection is fundamentally different from a chatbot that just retrieves information. Starling's decision to build incrementally — proving out analytical and advisory layers before granting agentic access — is the engineering pattern worth emulating. The teams that skip the trust-building phases and jump straight to agentic will discover why Starling didn't.

KEY TAKEAWAYS

→ Starling is the first UK bank to deploy agentic AI that acts on customer accounts — not just answers questions

→ The progressive build from analytical → advisory → agentic AI is the architectural pattern — each layer proved the trust model before the next was built

→ The competitive benchmark is now set: Revolut scaling, FCA rules landing April 6, and agentic AI in production at a UK neobank

→ The control plane challenge — rollback, blast-radius containment, prompt injection defence — is harder when the AI has write access to accounts

Sources: Starling Bank, Reuters (March 20–27, 2026)

QUICK TAKES

Two Stories Worth Knowing

SECURITY

CVSS 9.8

FIRST AI AGENT TO RECEIVE CVE ATTRIBUTION

XBOW Embeds AI Pen-Testing in Microsoft Security Copilot

XBOW, the autonomous offensive security platform that made history by receiving the first-ever AI CVE attribution (CVE-2026-21536, a critical unauthenticated RCE), has announced a deeper integration with Microsoft. The collaboration embeds continuous AI-driven penetration testing directly into

AI TOOLING

Apr 24

GITHUB TRAINING DATA OPT-IN TAKES EFFECT

Two AI Tooling Strategies Diverge

GitHub announced that from April 24, interaction data from Copilot Free, Pro, and Pro+ users will be used to train AI models. Business and Enterprise plans are unaffected. The nuance matters: private repository content at rest is excluded, but interaction data generated while working in private repos — inputs, outputs, code snippets, and associated

Security Copilot and Sentinel through three components: a Pentest Manager Agent, a Sentinel Connector, and a Pentest Analysis Agent.

Available as a public preview at RSAC 2026, this shifts the model from periodic manual pen-testing engagements to continuous, automated offensive assessment. The practical question for security teams is not whether AI pen-testing works — XBOW's CVE record proves it does — but how vulnerability management workflows handle the order-of-magnitude increase in discovered issues when scanning never stops.

Source: XBOW

context — is covered. Opt-out is available in settings.

The same week, Cursor shipped self-hosted cloud agents that keep code and tool execution entirely within your own network. These two announcements illustrate a stark divergence: GitHub is centralising data to improve its models; Cursor is decentralising execution to win trust. For engineering leaders evaluating code assistant governance, the question is no longer just capability — it is where your interaction data lives and who trains on it.

Sources: GitHub Blog, Cursor

OTHER NEWS THIS WEEK

In Brief

Claude Code vs Copilot vs Cursor: a no-fluff comparison for 2026. Useful reference for teams currently evaluating AI coding assistants — covers strengths, trade-offs, and where each tool fits. *cosmicjs.com*

Buy-now-pay-later regulation hits the UK perimeter. Deferred payment credit enters the FCA's regulatory scope from July 15, 2026 — engineering teams at BNPL providers should be planning compliance infrastructure now. *Osborne Clarke*

Nebius publishes detailed us-central1 post-mortem. A scheduled power maintenance event turned into a 9.5-hour multi-phase outage when provider work exceeded the announced scope — see Outside In for the full analysis. *nebius.com*

PLATFORM UPDATES

At a Glance

STATUS	PLATFORM	WHAT CHANGED
●	**Node.js**	Security releases for v20.x–v25.x: 2 high, 5 medium severity. Remote crash via TLS SNICallback is the priority. Patch now →
●	**AWS ME**	Second drone disruption to Bahrain region on March 24. No recovery timeline. Migration to alternate regions advised. Reuters →
●	**GitHub**	Copilot interaction data used for AI training from April 24 (Free, Pro, Pro+ plans). Business/Enterprise excluded. Policy →
●	**Cursor**	Self-hosted cloud agents shipped — code execution stays within your own network. Details →
●	**Starling**	Agentic AI assistant launched for personal accounts. Built on Google Gemini, opt-in, data stays within Starling's cloud perimeter. Details →

THE LEADERSHIP READ

The Preceptorship Model: Growing Engineers When AI Handles the Easy Work

Microsoft Azure CTO Mark Russinovich and VP of Developer Community Scott Hanselman have published a paper proposing a nursing-style "preceptorship" model for developing junior software engineers. Their argument: as AI automates the tasks that used to build foundational skills — writing boilerplate, debugging simple issues, reading documentation — organisations face a pipeline crisis that takes years to materialise and years to repair.

How it works

Preceptorship pairs early-career developers with senior engineers at 3:1 to 5:1 ratios. The juniors work on real product teams — prompting, debugging, and reviewing AI-generated code — while seniors externalise their judgment, turning daily work into structured learning. AI tools are deliberately configured for Socratic coaching rather than direct code generation: asking "why did you choose this approach?" instead of just producing the answer. Each cohort runs for at least one year.

> *"The path forward is a culture of preceptorship at scale. We must enable senior mentorship with AI systems that capture reasoning, surface misconceptions, and turn daily work into teachable moments."*

— Mark Russinovich and Scott Hanselman, Microsoft

The cautionary data point

The paper references Microsoft's own EDS Systems Engineering Development program, which was paused with a three-month recovery estimate. Actual recovery took over 18 months. Russinovich and Hanselman's point is stark: the talent pipeline is not a headcount lever — it is a supply chain, and supply chains are rebuilt, not restarted. The short-term incentive to stop hiring juniors because AI fills the gap will produce a future without the engineers capable of governing the AI systems being deployed today.

Read the full article: leaddev.com →

OUTSIDE IN

Nebius us-central1 Post-Mortem: When Scheduled Maintenance Becomes an Outage

~9.5h	4+	1h 46m	1h 22m
TOTAL INCIDENT DURATION	UNPLANNED POWER EVENTS	VM CONNECTIVITY LOSS	S3 ENDPOINT DOWN

What happened

On March 10, Nebius — a cloud GPU provider relevant to teams running AI workloads — experienced a major outage in its us-central1 region during what was supposed to be a controlled power infrastructure maintenance window. The provider-led electrical work did not remain within the announced maintenance profile. Instead of a single controlled event, it produced an initial multi-rack power outage at 15:18 UTC followed by at least four additional unplanned power interruptions over the next several hours, each disrupting recovery work already in progress.

Why the duration compounded

The site was operating in a temporary power configuration during maintenance. When an electrical anomaly caused incorrect switching and rack ATS unit failures, UPS capacity was insufficient to prevent the cascade. But the real duration driver was architectural: customer-facing networking, VM connectivity, Managed Kubernetes, and storage all depended on shared regional networking and a Kubernetes-based control plane. Power returned to different parts of the site at different times, meaning some hosts recovered before the control services needed to restart their workloads were healthy. The result: many workloads required manual reconciliation rather than automatic recovery.

The honest admission

Nebius acknowledged that during the incident, they did not follow the most effective restoration algorithm — too much recovery was handled as parallel service-by-service restoration instead of a stricter dependency-ordered regional recovery sequence. Their action items include hard stop-and-review criteria for ongoing electrical work and regular recovery drills with dependency-ordered restoration sequences. For any team relying on GPU cloud providers for AI workloads, this post-mortem is worth reading in full.

Read the full post-mortem: nebius.com →

ON THE RADAR

Coming Up

APR 06

FCA targeted support rules take effect. Authorised banks and pension providers can begin offering tailored recommendations. Engineering teams should have integration work ready. FCA →

APR 24

GitHub Copilot training data opt-in takes effect. Personal-plan interaction data will be used for model training unless users opt out. Review your organisation's Copilot tier and data governance policies. GitHub →

JUN 03

Money20/20 Europe 2026. Amsterdam, June 3–5. Agentic finance and AI governance expected to dominate the agenda after Starling's launch this week.

Q2 TBC

FCA consultation on retained EU payments law. Will include the approach to open banking and stablecoin payments. Related statutory instrument expected 2027/2028.

RECOMMENDED READING

This Week's Links

Starling launches UK's first Agentic AI financial assistant	starlingbank.com
A model for growing the next generation of developers	leaddev.com
Nebius us-central1 incident post-mortem	nebius.com
Claude Code vs GitHub Copilot vs Cursor: 2026 comparison	cosmicjs.com

Engineering Digest

March 20–27, 2026 · Weekly Edition · Issue #2.5

Next edition: March 27 – April 3, 2026

ARCHIVE
Mar 18 · Weekly #0
Mar 18 · Deep Dive
Mar 18 · Weekly #0.5
Mar 20 · Weekly #1
Mar 27 · Weekly #2.5