# Engineering *Digest*

May 8–15, 2026 | Weekly Edition | Issue #10 | ~18 min read

THIS WEEK'S BRIEFING

**67.48%**

Open-source talent is in motion — and AI impact views are the split that defines retention risk in 2026.

IN THIS ISSUE


- FCA safeguarding regime: reconciliation pipelines as audit evidence
- DORA ROI of AI-assisted development — the J-curve warning
- Wise: observability re-architecture through ownership culture
- Mini Shai-Hulud: when valid SLSA attestations aren't enough
- Agentic coding and the IC/EM boundary collapse
- Platform Updates, On the Radar & Down the Rabbit Hole


## FOREWORD

Three stories this week share a structural pattern: something that was a process concern has become a production obligation. The FCA's safeguarding regime turns reconciliation into audit evidence. The Mini Shai-Hulud supply-chain worm shows that valid provenance certificates mean nothing when the pipeline generating them is compromised. The DORA team gives engineering leaders a financial model rigorous enough to take to a CFO — along with a warning about what happens when the underlying platform is weak.

PRIORITY 0 — ONGOING

## AWS ME-CENTRAL-1 Infrastructure Recovery — Day 75+

RECOVERING

| | |
|---|---|
| **May 15** | No new ME-CENTRAL-1-specific public statement found this week. Multi-month recovery timeline remains the operating assumption per AWS guidance. |
| **Mar 1** | Original infrastructure event. AWS confirmed extended recovery timeline. |

Note: the Northern Virginia cooling/power incident (May 8, us-east-1) is a separate event — resolved within 24 hours and covered in Platform Updates below. The ME-CENTRAL-1 situation remains distinct and ongoing.

**67.48%**

NUMBER OF THE WEEK

Two-thirds of open-source professionals are actively looking for a new role or open to the right opportunity, according to the 2026 Open Source Professionals Job Survey (OSPJS — the Linux Professional Institute and Open Source JobHub's third annual survey of FOSS — Free and Open Source Software — professionals). The survey also finds that 83.58% of respondents rate employer-provided training as important when choosing a role. For engineering leaders managing open-source-heavy teams, this is a talent-risk signal with a specific implication: teams that invest visibly in structured learning and thoughtful AI tooling adoption are operating in a market where two-thirds of the relevant talent pool is already in motion. The variable that separates firms in this environment is likely not compensation alone — it is whether the organisation's approach to AI-augmented work looks like a growth path or an extraction. Full report →

THE WEEK AT A GLANCE

# Four more stories worth knowing

Action required

Security

**vm2 sandbox-escape wave: 13+ CVEs in early May.** AI agent frameworks and SaaS automation services running untrusted JavaScript are the highest-risk targets. Check whether your agentic workflows use vm2 or any derivative sandbox library.

*The Hacker News →*

Worth knowing

AI Tooling

**Cursor 3 ships parallel agents, a PR review workflow, and a public SDK.** Coding agent tooling reaches CI/CD integration maturity: multiple agents can now run concurrently against the same codebase, with the SDK enabling custom integrations. *ReleaseBot →*

Regulation

**HM Treasury and FCA payments overhaul takes shape.** The PSR folds into the FCA and stablecoins, along with agentic AI payments, are explicitly addressed in the new unified framework — the first UK regulatory instrument to name AI-initiated payments as a distinct category. *AutoRek →*

Watching

Open Source

**OSI 2026 State of Open Source: geopolitical pressure is now a board-level risk topic.** The report frames open-source dependency as a strategic concern — compliance complexity, security exposure, and operational maintenance burden all rising. Context for the supply-chain story in Quick Takes. *OSI →*

LEAD

# The FCA's new safeguarding regime turns reconciliation pipelines into audit evidence

| 7 May | D+1 | Monthly | Annual |
|---|---|---|---|
| Effective date — CASS 15 live | Reconciliation frequency — daily | Dedicated safeguarding return | CASS 15 audit — first within 6 months of period end |

## What went live on 7 May

JARGON WATCH

### PS25/12 & CASS 15

As of 7 May 2026, the FCA's (Financial Conduct Authority, the UK's primary financial services regulator) Safeguarding Supplementary Regime is in force under PS25/12 for all authorised payment institutions — firms authorised by the FCA to provide payment services, distinct from a bank — and e-money institutions (EMIs — regulated firms licensed to issue electronic money, such as digital wallets and prepaid cards). The regime replaces the principles-based safeguarding framework of CASS 10A with the prescriptive obligations of the new CASS 15 chapter.

> Policy Statement 25/12 is the FCA's published rulebook instrument that brought the Safeguarding Supplementary Regime into force on 7 May 2026. The "25" is the year of issue; "12" the sequence number. It amends the FCA Handbook to insert CASS 15 — a new chapter of the Client Assets sourcebook — alongside the pre-existing CASS 10A (the FCA chapter on payment and e-money safeguarding, now supplemented by CASS 15). The key change for engineering teams: what were principles-based expectations became enforceable rules with daily, monthly, and annual data-delivery obligations.

The headline change is not the existence of safeguarding — the requirement to ringfence customer funds in dedicated accounts, segregated from the firm's own money, so they can be returned if the firm fails — but the standard of evidence that safeguarding now demands: production-grade data infrastructure across four distinct obligations.

## The four engineering obligations

First, D+1 reconciliation (a reconciliation run completed by the close of the business day following the reference date) is now a daily production requirement: internal reconciliation against the D+1 segregation requirement, and external reconciliation against third-party bank data, both must run every reconciliation day. Shortfalls must be remedied from the firm's own funds the same day.

Second, a dedicated monthly safeguarding return replaces the previous monthly questions — covering the D+1 segregation requirement, reconciliation completion status, and data on relevant assets and custodians. Third, most firms must arrange an annual CASS 15 audit under SUP 3A (the FCA Handbook chapter governing auditor appointment), with the first report due within six months of the period end. Fourth, firms must maintain a resolution pack — account maps, safeguarding records, and reconciliation logs — that is current and immediately usable by an insolvency practitioner, not a static document filed once.

> *The FCA expects an increase in adverse audit opinions in the short term, as greater scrutiny surfaces previously hidden deficiencies in safeguarding practices.*
>
> — Osborne Clarke analysis of PS25/12 (May 2026)

## From process to production

The practical pressure is one of reliability, not existence. Most payment institutions and EMIs already run reconciliation processes — but those were built for internal reporting, tolerating batch delays and manual corrections invisible to an auditor. CASS 15 changes the test: the question is no longer whether reconciliation ran, but whether the output is accurate enough for FCA audit. Reconciliation must be treated as a first-class production system; the resolution pack as a data product rather than a document — its accuracy only as good as the pipelines that feed it.

## What comes next

The FCA will review the regime's effectiveness after the first full audit period — firms that discover pipeline deficiencies after the first CASS 15 opinion is issued face remediation inside an active supervisory relationship, which is the harder path.

For teams building in the Middle East, the ADGM (Abu Dhabi Global Market — the international financial centre and free zone in Abu Dhabi) published AML/CFT (Anti-Money Laundering / Counter Financing of Terrorism) amendments also effective from May 1, creating a parallel regulatory two-front week for any engineering leadership spanning both UK and ME payment operations.

**KEY TAKEAWAYS**

→ Daily reconciliation pipelines and resolution packs are now audit evidence — a team whose pipeline was good enough for internal reporting may not be good enough for CASS 15 scrutiny.

→ The first CASS 15 audit opinions are ~12 months away — the implementation window is narrow, and adverse opinions are the expected starting point, not an aberration.

→ Resolution packs are data products, not documents — their accuracy depends on the quality of the pipelines that feed them.

*Sources: Osborne Clarke, FCA PS25/12, Bratby Law (May 2026)*

OUTSIDE IN

## Wise — Leading an observability re-architecture through ownership culture

When Marta Lima joined Wise's Observability domain, the team was mid-migration: rebuilding the entire observability stack — migrating to the LGTM stack (an open-source observability toolkit composed of Loki for log aggregation, Grafana for dashboarding, Tempo for distributed tracing, and Mimir for long-term metrics storage; maintained by Grafana Labs and widely used as an alternative to proprietary APM platforms) — concurrently with a Kubernetes platform migration (Kubernetes — open-source container orchestration system, the standard for running containerised applications across clusters), without interrupting the engineering organisation. The stack is built in-house — a deliberate choice reflecting Wise's requirement to handle high data volumes under strict regulatory constraints across multiple countries — and the migration took over a year. What made it tractable, Lima argues, was not technical planning alone but Wise's domain ownership culture: teams are fully accountable for their domains, with genuine authority to move fast when the data supports it. The platform team's role is to provide the underlying infrastructure that lets product engineers focus on product; the observability team's role is to give every engineer real-time visibility into their service's behaviour — which is the ability to understand the internal state of a system from its external outputs (logs, metrics, and traces), distinct from monitoring, which watches for predefined conditions. Neither depends on a central coordination layer to make decisions.

The transferable pattern for engineering leaders managing platform modernisation without dedicated SRE headcount is the accountability structure that came first: when a team owns the outcome end-to-end, concurrent migrations become tractable because coordination overhead is absorbed within the domain rather than escalated.

**The takeaway:** Wise's LGTM observability migration succeeded not because it was well-planned but because domain ownership meant the team absorbed the coordination complexity that would otherwise have required cross-org negotiation.

Marta Lima, Engineering Lead · Wise Blog (May 2026) →

QUICK TAKES

# Two Stories Worth Knowing

**Security & Compliance**

JARGON WATCH

### Mini Shai-Hulud

The name given by security researchers to a self-propagating supply-chain worm documented in 2026 by TeamPCP, a threat actor. Named after the giant sandworms in Frank Herbert's *Dune* for its ability to bore through build pipelines and surface at unexpected points. This was its third documented wave in 2026; prior targets included Aqua Security Trivy and Bitwarden CLI.

**170+**

packages reached; 12M+ weekly downloads for @tanstack/react-router alone

## Valid provenance wasn't enough: Mini Shai-Hulud turns CI/CD trust into the attack surface

On 11 May 2026, TeamPCP (the threat-actor designation used by Wiz and security researchers for the group behind the Mini Shai-Hulud worm campaigns) exploited a `pull_request_target` misconfiguration — a GitHub Actions trigger that runs with write permissions even from an external fork's pull request — in TanStack's monorepo. The attacker used this to poison the pnpm build cache (pnpm is a fast Node.js package manager; cache poisoning here injected malicious content that subsequent builds fetched as legitimate), then extracted OIDC

**Developer Experience**

**39%**

First-year ROI on AI-assisted development — 500-person org model

## DORA gives AI-assisted development a CFO model — and warns about the J-curve

"The greatest returns on AI investment come not from the tools themselves but from a strategic focus on the underlying organizational system: the quality of the internal platform, the clarity of workflows, and the alignment of teams. Without this foundation, AI creates localized pockets of productivity that are often lost in downstream chaos." That is the central claim of the DORA ROI of AI-Assisted Software Development report — where DORA refers to the DevOps Research and Assessment team at Google Cloud, not the EU Digital Operational Resilience Act (a different DORA in financial-services regulation context).

The illustrative model for a 500-person engineering organisation ($176k fully loaded salary) shows ~$11.6m first-year return against ~$8.4m investment: 39%

tokens (OpenID Connect — short-lived identity tokens used to authenticate GitHub Actions workflows against external services like the npm registry) directly from runner process memory. Those tokens were used to publish malicious packages to the npm (Node Package Manager registry — the primary JavaScript/TypeScript distribution channel) and PyPI ecosystems without ever stealing credentials.

The published packages carried valid SLSA Build Level 3 attestations and valid Sigstore signatures — SLSA (Supply-chain Levels for Software Artifacts, a Google-originated framework defining build-pipeline integrity levels) and Sigstore (an open-source toolkit for signing software supply chains using short-lived cryptographic certificates) both confirmed the pipeline ran correctly. They could not confirm the pipeline itself was the attack surface.

84 malicious artifacts across 42 @tanstack packages were published on May 11. The worm self-propagated to Mistral AI, UiPath, and over 160 other packages within hours. Environments that ran @tanstack packages on May 11 are treated as fully compromised in the incident response pattern: the remediation cycle involves rotating all accessible secrets (GitHub tokens, npm tokens, AWS credentials, Kubernetes service accounts, CI/CD secrets), removing persistence hooks in `.claude/` and `.vscode/` directories, and blocking C2 domains git-tanstack.com and *.getsession.org. An SBOM (Software Bill of

ROI, ~8-month payback. The headline figure is less useful than the J-curve warning that accompanies it: a temporary productivity dip before long-term gains, caused by the learning curve, a "verification tax" from reviewing AI-generated code, and the need to adapt downstream processes like testing and change approval to handle increased code volumes. Leaders who mistake the dip for failure risk pulling funding during the period that generates the eventual return. An interactive calculator is available at dora.dev/ai/roi/report.

---

**The takeaway:** Engineering leaders now have a Google-branded financial model to take to a CFO conversation about AI tooling — and a specific term ("the J-curve") for the temporary dip that finance teams will otherwise treat as evidence the investment isn't working.

*Sources: InfoQ, DORA report, Google Cloud (May 2026)*

Materials — a structured list of all components in a software build) with valid attestations offers no protection when the build system generating the attestations is the attack surface. Structural mitigations include auditing `pull_request_target` workflows, pinning GitHub Action versions to commit SHAs rather than tags, and applying GitHub's OIDC token scoping mitigations.

---

**The takeaway:** SLSA attestations and Sigstore signatures confirm that the build pipeline ran correctly — they cannot confirm the pipeline was not itself compromised; a provenance chain that starts at a compromised build environment is valid all the way down.

*Sources: Wiz, TanStack postmortem, Snyk, StepSecurity (May 2026)*

> *Fidelity Investments is trimming roughly 800 jobs — about 1% of its 80,000-person global workforce — while simultaneously opening 2,000 positions, 400 of them in tech or product roles. The restructuring explicitly targets senior leadership layers while adding early-career, hands-on engineering capacity. This is a financial-services org-shape signal worth watching: Fidelity's answer to the question of how AI-augmented delivery changes levelling appears to be fewer senior-strategic layers and more execution headcount. Whether that ratio holds as agent-first workflows mature is the open question for engineering leaders thinking about their own org design over the next 18 months.*
>
> Boston.com (May 11, 2026) →

PLATFORM UPDATES

# At a Glance

| Status | Platform | What changed |
| --- | --- | --- |
| ● | **GHES 3.21 RC** | **GitHub Enterprise Server** (GHES — the self-hosted version of GitHub, as distinct from github.com) 3.21 release candidate is available. Self-hosted teams should begin upgrade planning. Key items: org custom properties reach GA, hierarchy view for Projects ships, secret scanning permission model changes, and the REST API advances to version 2026-03-10 — a breaking change affecting GitHub App integrations. App owners should audit the migration path before enterprise rollout. *The GitHub Blog →* |
| ● | **.NET May 2026** | Microsoft's May 2026 servicing update covers **.NET 8 LTS** (LTS — Long-Term Support, the channel receiving security and stability patches for 3+ years), .NET 9, and .NET 10, plus .NET Framework 3.5 through 4.8.1. Two EoP (Elevation of Privilege — a vulnerability class allowing an attacker to gain higher permissions than granted) CVEs (Common Vulnerabilities and Exposures — standardised identifiers for disclosed security vulnerabilities) affect the runtime across all three current versions. Teams on any of these channels should ship this patch cycle. *.NET Blog →* |
| ● | **Patch Tuesday May 2026** | 130 CVEs, 30 Critical. The two actionable items for most teams: **Windows DNS Client RCE** (RCE — Remote Code Execution, a vulnerability class allowing an attacker to run arbitrary code from a remote location; CVSS — Common Vulnerability Scoring System, 0–10 severity scale — 9.8, unauthenticated, zero interaction required) and two M365 Copilot information disclosure CVEs (CVE-2026-26129, CVE-2026-26164). An Azure DevOps CVSS 10.0 item is cloud-remediated by Microsoft; no customer action needed. *CrowdStrike →* |
| ● | **AWS us-east-1** | A thermal event caused a 24-hour EBS (Elastic Block Store — AWS's persistent block storage service for EC2 instances) and EC2 (Elastic Compute Cloud — AWS's virtual server service) availability zone failure in us-east-1 (Northern Virginia) on May 8. Cooling stabilised by 1:50 |

| Status | Platform | What changed |
| --- | --- | --- |
| | | PM; most volumes recovered the same day. Engineering angle: multi-region resilience posture and dependency concentration in a single AZ (Availability Zone — an isolated data centre location within an AWS region). **Note:** this is a separate incident from the ongoing AWS ME-CENTRAL-1 P0 story above. *Mashable →* |
| ● | **Claude Security** | Anthropic's Claude Security is now in public beta for all Enterprise customers — native AI vulnerability scanning and patch suggestions accessible from the Claude.ai sidebar, with no custom API integration required. Positioned as an AI-native alternative to traditional SAST (Static Application Security Testing — automated analysis of source code for vulnerabilities before the code runs). *Infosecurity Magazine →* |

IN PRACTICE

*Calibration & performance conversations*

## Engineering performance calibration as decision-making infrastructure

*Pensero editorial team · Pensero, April 2026 · Read the full piece*

Calibration sessions often function as opinion-reconciliation meetings — managers arrive with different standards, negotiate toward rough consensus, and leave with decisions that are difficult to revisit later. The specific framework this piece argues for reframes calibration as decision-making infrastructure rather than a calendar ritual: a structured process with a defined data layer, a two-tier conversation model, and explicit handling of variables that current levelling rubrics do not address.

The data layer rests on complexity-weighted delivery metrics (metrics that weight engineering output by the estimated complexity of the work, rather than counting story points or commits — which prevents large but simple work from distorting comparisons). The specific metrics proposed: delivery per headcount, defect rate, AI adoption share, roadmap alignment, cycle time, capitalizable delivery (work that qualifies as capital expenditure under accounting rules, relevant

for teams whose CFO tracks software development cost categories), and knowledge gaps. The key structural argument is that VP and CTO-level calibration — team-to-team comparison for resource allocation — requires different data than EM-level calibration, which is about individual promotion evidence. Conflating the two produces sessions where neither conversation happens well.

The 2026-specific variable the piece introduces is agent-authored delivery (code or output produced by an AI coding agent rather than directly by a human engineer). Most levelling rubrics were designed around code-output signals — PRs merged, velocity, lines of code — that do not yet distinguish human from agent contribution. An engineer whose throughput has increased significantly this quarter may be demonstrating either mastery of AI-augmented workflows or a dependency on agent output that lacks the signal quality the framework was designed to measure. Calibration sessions in Q2 2026 will encounter this split for the first time, often without explicit data to resolve it.

> *Calibration works when the data reflects the work, not just the volume of output. The agent-authored delivery variable is the first place where this distinction becomes visible in the numbers.*

The concrete mechanism the piece proposes is treating agent-authored delivery share as a new column in the calibration data view — surfacing which engineers are extending AI-augmented capability in ways that compound their judgement, and which are generating throughput that looks like productivity but reflects authorship volume rather than decision quality. Without this column, calibration sessions reward the latter as readily as the former.

---

**WHAT TO TRY**

One anchor for the next calibration session: before the meeting, ask whether the data on the table distinguishes between team-to-team comparison and individual promotion evidence — and, if anyone on the team has meaningfully shifted toward agent-augmented delivery this cycle, whether that shift is visible in the metrics or invisible beneath a throughput number.

THE LEADERSHIP READ

## Agentic coding is making engineering-management skills part of the IC job

Sage Lazzaro's piece in LeadDev is not primarily about AI. It is about what happens to role boundaries, levelling frameworks, and career paths when the competencies that engineering

managers have spent years developing — breaking down problems, delegating, setting context, managing toward outcomes — become the core skill of an IC (individual contributor — an engineer without direct reports, whose primary work is technical execution) working with AI agents.

### The boundary collapse

The specific argument is that agentic coding (the practice of using AI coding agents — tools like GitHub Copilot, Cursor, or Claude Code — to plan, implement, test, and in some cases merge code with reduced or restructured human oversight) is collapsing the IC/EM (engineering manager — an engineer with direct reports, responsible for people management, team direction, and delivery coordination) boundary from both directions simultaneously. ICs managing five AI agents in parallel need to prioritise tasks, set context, review outputs, and manage toward a delivery outcome — structurally the same work as managing a team of junior engineers. Meanwhile, experienced EMs are being actively recruited as senior ICs at agent-first companies, because their skill set maps directly to orchestrating agent workflows.

Scott Breitenother, co-founder and CEO of Kilo — an agentic coding startup that launched in early 2025 — says: "Instead of having teams of humans working together, each individual contributor ends up being like a manager. But instead of a manager of humans, they're a manager of agents." Every engineer Kilo has hired was previously either an engineering manager or a senior peer mentor. Breitenother's hypothesis is that the adjustment from managing humans to managing agents is smaller than the adjustment from managing no-one to managing agents: the delegation and context-setting muscle is already there.

### Where levelling frameworks are misaligned

The second concrete mechanism is the mismatch between how current levelling frameworks measure performance and what the work now requires. Levelling rubrics built on code-output signals — PRs merged, velocity, lines of code delivered — reward throughput. An IC managing multiple agents can generate high throughput while the work that actually distinguishes senior contribution (context quality, delegation decisions, output review, delivery management) is invisible in the metrics. Vinay Perneti, VP of Engineering at Augment Code, puts it plainly: "Instead, I think you really want to embrace outcomes."

The data the article cites is pointed: staff and principal engineers now regularly out-earn engineering managers by 15 to 25% at top companies, according to IT hiring firm HR Oasis. Google has cut 35% of managers overseeing fewer than three people. Amazon required every

team to increase its ratio of ICs to managers by 15% by Q1 2025. These are separate decisions taken by separate organisations, but the direction is consistent: the management layer is contracting while the IC layer is expected to absorb more of what was previously management work.

### The career-path implication

For engineering leaders, the useful tension here is not that "AI changes everything." It is that careers and team structures are shifting before the levelling frameworks that govern them have caught up. An engineer who is demonstrably excellent at managing agent workflows may be measured — and paid, and promoted — against a rubric designed for a world where that skill did not exist. That gap will surface in calibration sessions, in retention conversations, and in hiring decisions. Recognising it explicitly is more useful than waiting for the frameworks to self-update.

*Read the full article: leaddev.com →*

---

**This could be worth talking about:** *If your levelling rubrics were redesigned today, would they reward the quality of delegation and outcome-management decisions — or would they still reward throughput, with agent-assisted throughput looking indistinguishable from human-authored throughput?*

DOWN THE RABBIT HOLE

Long-read

**If Outside In gave you the culture-and-architecture frame, here's the stack behind it.** ByteByteGo's breakdown of the Wise infrastructure gives the granular numbers: the Mimir cluster (part of the LGTM stack) processing 6 million metric samples per second, Temporal — the workflow engine for automated database failovers, distinct from the JavaScript Temporal API in Node.js 26 that led Issue 9 — and the Mission Zero cost programme that ties it together. The technical depth behind Wise's ownership model is in this piece.

ByteByteGo: The Tech Stack Powering Wise →

ON THE RADAR

# Coming Up

MAY 2026

**First CASS 15 audit cycle** — FCA commentary expected after the first full audit period completes under the Supplementary Safeguarding Regime. Watch for adverse-opinion rates and any supervisory guidance update. No new FCA publication yet; monitor PS25/12 publications page for updates.

OCT LTS DATE

**Node.js 26 LTS path** — Not re-covering as a story (Issue 9 lead). For platform teams with native add-ons: the NODE_MODULE_VERSION 147 bump in Node.js 26 requires a rebuild cycle before the October LTS date. Check your native dependency inventory now. Node.js 26.0.0 →

MAY 1

**ADGM AML reforms** — Abu Dhabi Global Market (ADGM — the international financial centre in Abu Dhabi with its own regulatory framework operated by the FSRA — Financial Services Regulatory Authority) AML/CFT amendments effective May 1. ME fintech teams should assess impact on KYC (Know Your Customer — regulatory obligations requiring firms to verify customer identity and monitor transactions) workflow design and reporting pipelines. BDO UAE / Law Corporation →

MAY 2026

**GHES 3.21 upgrade runway** — GitHub App owners integrating against the REST API should check the 2026-03-10 migration path (the GitHub REST API version identifier for the March 10, 2026 release — a breaking-change version; the old 2022-11-28 version will be deprecated) before enterprise rollout of 3.21. GitHub Changelog →

2026 CONTEXT

**Open-source governance as board risk** — The OSI 2026 State of Open Source report provides structural context behind the Mini Shai-Hulud supply-chain story: geopolitical pressure, compliance complexity, and operational maintenance burden are now board-level risk topics for organisations with significant open-source dependency. Context only; outside the window. OSI →

SOURCES

# Primary References

**Lead — FCA Safeguarding Regime:** *Osborne Clarke, FCA PS25/12, Bratby Law*

**Quick Take 1 — Mini Shai-Hulud:** *Wiz, TanStack postmortem, Snyk, StepSecurity*

**Quick Take 2 — DORA ROI:** *InfoQ, DORA report, Google Cloud*

**Editorial Moment — LPI Survey:** *LPI announcement, Full OSPJS report*

**Mid-issue Accent — Fidelity:** *Boston.com*

**Platform Updates:** *GitHub Blog, .NET Blog, CrowdStrike, Mashable, Infosecurity Magazine, Claude support docs*

**In Practice:** *Pensero*

**Leadership Read:** *LeadDev — Sage Lazzaro*

**Outside In:** *Wise Blog — Marta Lima*

**Down the Rabbit Hole:** *ByteByteGo*

**On the Radar:** *Node.js 26.0.0, ADGM AML, OSI 2026*

**P0 Status context:** *Mashable (NVA incident, not ME-CENTRAL-1)*

# Engineering Digest

May 8–15, 2026 · Weekly Edition · Issue #10

Next edition: May 15–22, 2026

Archive

Mar 2026 · Weekly | Mar 2026 · Weekly | Apr 2026 · Weekly | May 2026 · Weekly (Issue 9) | May 8–15 · Weekly (Issue 10)