

Engineering *Digest*

Engineering decisions, every Friday.

22–29 May 2026

THIS WEEK'S BRIEFING

~7 hrs

A single data-centre hall lost cooling — and a multi-billion-dollar exchange went dark for most of a working day.

IN THIS ISSUE

GitHub Copilot moves to usage-based AI Credits
Coinbase, "no pure managers" & the layoff wave
Outside In: incident.io on customer-first reliability

The AWS US-EAST-1 thermal-event post-mortems
The Leadership Read: a guide to calibration
In Practice, Other News & On the Radar

FOREWORD

A NOTE FROM THE EDITOR

A small change of dress this week. From this issue, the digest opens at Volume II — a sharper cover, a standing tagline, and a few markers that should make each edition easier to tell apart at a glance. The brief behind it is unchanged: the engineering decisions worth a Friday morning, nothing more.

The thread this week is cost made visible. GitHub Copilot turns a fixed per-seat line-item into a metered one on 1 June. A failed cooling loop in a single AWS hall put a number on single-region concentration. And a layoff wave still running through late May is, at root, a price

being put on the management layer itself. Three different ledgers, one question underneath them: what is actually being paid for, and by whom.

\$47bn

NUMBER OF THE WEEK

The six largest US banks posted roughly \$47bn in combined quarterly profit while cutting around 15,000 jobs — the shape of a decision to run AI as cost takeout rather than a revenue engine, and the strategic split the banks in this sector's orbit are now choosing between. [Implicator.ai](#) →

THE WEEK AT A GLANCE

Seven Stories in Sixty Seconds (2 min)

Action required

Security

Microsoft Defender zero-days under active exploitation. Two flaws — RedSun (CVE-2026-41091) and UnDefend (CVE-2026-45498) — are confirmed by CISA as exploited in the wild, with patches rolling out. For Windows-dependent orgs the leadership task is to confirm patching is owned and tracked, not to remediate it.

[BleepingComputer](#)

Worth knowing

Tooling

GitHub Copilot usage-based billing goes live 1 June. Flat Premium Request Units give way to metered AI Credits — the issue's Lead unpacks what changes for the budget. [github.blog](#)

Research

DORA publishes its "ROI of AI" report. DORA — the long-running delivery-research survey, not the EU regulation — finds returns come from the organisational system, not the tools. Covered as the Issue 10 Quick Take; no fresh treatment here. [InfoQ](#)

M&A

Anthropic acquires Stainless for \$300M+. The hosted SDK generator used by OpenAI, Google and Cloudflare will wind down, concentrating the developer-tooling layer under one vendor. [TechCrunch](#)

Payments

Visa and Mastercard scale agentic commerce. The Trusted Agent Protocol and Agent Pay move toward mainstream for the 2026 holiday season — a signpost for any loyalty or commerce-media roadmap planning agent checkout. [American Banker](#)

Watching

Fintech

The Middle East fintech frontier keeps building. The UAE's FIT programme is roughly 85% complete, Saudi open finance is expanding, and new regional cloud regions are landing — concrete integration drivers for ME bank partners. [Mobile Ecosystem Forum](#)

Tooling

"The harness is the product." The four big labs now price the agent runtime incompatibly — see the number band later in this issue. [The New Stack](#)

WAYS OF WORKING

GitHub Copilot Turns a Fixed Cost Into a Metered One (3 min)

From 1 June, GitHub Copilot stops being a flat per-seat purchase and becomes a usage-metered one — and the budget conversation changes with it. Seat prices hold and code completions stay unlimited, but the agentic work that teams have leaned into hardest is now billed by the token. A line-item that was easy to forecast becomes one that moves with how teams actually work.

What actually changes

GitHub is replacing Premium Request Units with "GitHub AI Credits." Each plan includes a monthly allotment of credits matched to its seat price — Copilot Business at \$19 with \$19 of credits, Enterprise at \$39 with \$39 — and paid plans can purchase more. Crucially, the company confirms that "code completions and Next Edit suggestions remain included in all plans and do not consume AI Credits." The metering applies to the agentic features: the work where, as GitHub puts it, "agentic usage is becoming

JARGON WATCH

Premium Request Units (PRU)

The retiring model: a fixed monthly allowance of "premium requests" bundled into each seat. Predictable, but blunt — heavy and light users cost the same.

GitHub AI Credits

The replacement: a monthly credit allotment consumed by token usage (input, output and cached tokens) at per-model rates, with paid plans able to buy more.

the default, and it brings significantly higher compute and inference demands."

A precision note on language. The agentic features being metered here are agent-mode workflows inside the IDE — a human still triggers a task and reviews the result — not autonomous production agents with authority to act on systems unattended. The bill scales with how much of that supervised, multi-step work a team does.

"A line-item teams could forecast to the seat becomes one that moves with how the team works that month."

— Engineering Digest, on the Copilot cutover

Why it lands as a budget decision

The shift converts a fixed cost into a variable one, and that is the whole story for engineering leadership. Teams that have pushed agent-mode usage furthest — exactly the teams a year of adoption pressure encouraged — are the ones most exposed to rising bills. We covered the DORA ROI model in Issue 10; the Copilot change is the moment that model meets the actual invoice, where the question of where AI returns come from stops being theoretical and starts appearing on a statement.

	UNTIL 31 MAY	FROM 1 JUNE
Pricing unit	Flat seat + Premium Request Units	Flat seat + metered AI Credits (by token)
Business seat	\$19 / user / month	\$19 / user / month + \$19 credits included
Enterprise seat	\$39 / user / month	\$39 / user / month + \$39 credits included
Code completions	Unlimited	Unlimited — no credit cost
Agentic features	Counted as premium requests	Metered by token consumption

Data: [The GitHub Blog](#)

The conversation it forces

The leadership conversation is concrete and imminent: how to meter and cap agent usage across teams, and how to renegotiate the AI line in the budget before consumption, not after. The trade-off engineering leaders are now navigating is between letting agent usage run free

and keeping spend legible — a finance discipline that, until 1 June, the flat model quietly handled on their behalf. GitHub also paused new Pro, Pro+ and Student sign-ups in late April, a smaller signal pointing the same way: the economics of supervised agent work are being repriced in real time.

Source: [The GitHub Blog](#) (announced 12 May; effective 1 June 2026)

QUICK TAKES

Two Stories Worth Knowing

Operating at Scale · DR

~7 hrs

Coinbase offline in the US-EAST-1 thermal event

When a single hall loses its cooling

On 7–8 May, a single data-centre hall in AWS's US-EAST-1 region — its busiest, in Northern Virginia — lost cooling. The temperature spike inside one availability zone (one of the physically separate sites AWS runs within a region) triggered a power loss that physically damaged EC2 and EBS hardware on the affected racks. EC2 and EBS are AWS's core compute and block-storage building blocks. The DR post-mortems that landed this week, dissecting that earlier event, are the news here — and the reading they offer is uncomfortable.

AWS's own status note described "an increase in temperatures within a single data center" that "caused impairments for instances in the Availability Zone." The fallout reached well beyond AWS: Coinbase was knocked offline for the better part of a working day — reported at around seven hours — alongside disruption at FanDuel and CME Group. Recovery was slow precisely because the failure was physical; damaged hardware does not reboot.

Management & Teams

~700

Coinbase roles cut — about 14% of staff

CONCEPT OF THE WEEK

Player-coach

A manager who keeps doing hands-on delivery work alongside leading the team — coaching while still carrying their own output. It removes the "pure manager" whose role is only coordination, and raises the bar on what a management title is expected to produce.

The post-AI org shape, named out loud

Coinbase announced on 5 May that it would cut about 14% of staff — roughly 700 roles — and rebuild with "no pure managers," replacing them with "player-coaches" who lead while still contributing, capping the org at five layers below the CEO and floating "AI-native pods," potentially one-person teams directing agents. CEO Brian Armstrong's framing was blunt: "Layers slow things down and create coordination tax." The

The diagram below traces how one cooling fault cascaded into a customer-facing outage. The post-mortem coverage focuses on the cross-region disaster-recovery lesson: a workload concentrated in one region inherits that region's single points of failure, however many availability zones sit underneath it.




Diagram: Engineering Digest, from *IT Pro* reporting

The takeaway: a workload sitting in one region carries that region's single points of failure no matter how many availability zones it spans — and physical damage turns a short outage into a long one, which is what makes cross-region DR a live design question again rather than a box already ticked.

Source: *IT Pro* (original event 7–8 May; DR analysis published this week)

wave kept running through the window — Webflow and Groupon cut staff on 27 May — which is what keeps the 5 May announcement live.

The signal is the cross-company pattern, not any one firm: management layers thinned and org charts flattened as the default response to AI. The contrast below frames the strategic split underneath it.

INCUMBENTS	FINTECH-NATIVE
AI as cost takeout — ~15,000 jobs cut, \$47bn combined profit (six largest US banks, Q1)	AI-native operating model — superior revenue per employee, API-first delivery
Efficiency engine: do the same work with fewer people	Revenue engine: new product surface per head of staff

The takeaway: the player-coach model and layer-capping reshape levelling, span-of-control and the IC-versus-management track directly — a "pure manager" title is being recast as a cost to justify rather than a default rung on the ladder.

Sources: *Fortune*, *InformationWeek*, *Implicator.ai*

Anthropic's price for its Managed Agents runtime — the "harness" that holds agent state, retries and tool hygiene. OpenAI charges no first-party runtime fee; Google and Microsoft bundle priced primitives inside broader platforms. Whichever harness a team standardises on now sets its multi-year AI cost structure and switching cost.

IN PRACTICE

Manager-of-managers craft (2 min)

A classic worth revisiting



Jay Vilalta
Director of Engineering · personal blog
September 2025

[Read the piece →](#)


Managing managers: a director's short guide

Jay Vilalta · personal blog, September 2025 · [Read the full piece](#)

"I'm not here to catch anyone doing something wrong. I want to understand how we can better support you and your manager."

That line is the hinge of Vilalta's account of managing managers, written by a director who has run dozens of them. He holds regular skip-levels — a one-to-one between a leader and someone two levels down — with every individual contributor, and frames them explicitly as support rather than surveillance. The distinction is the whole point: a skip-level read as a check-up on the manager poisons the well; read as "how do we support you and your manager," it surfaces problems early without undermining the chain.

His second mechanism is heavy structure for new managers. The jump from individual contributor to people-leader overwhelms most people, so Vilalta raises first-time managers to weekly one-to-ones and leans on concrete scaffolding — role-play, delegation coaching, pairing with an experienced mentor — rather than assuming the title confers the skill. The process flow below traces the cadence he describes.



The third move is a catalogue of recurring organisational dysfunctions: unclear or shifting priorities, weak tooling, poor cross-team communication, misaligned incentives, thin planning. Vilalta's discipline is to treat these as structural rather than individual — they rarely stem from one person failing, and the manager-of-managers' job is to fix the system that produces them, not to find someone to blame. That reframe is what separates a director who scales their managers from one who simply audits them.

WHAT TO TRY

The skip-level framing is the transferable piece here — whether a skip-level reads as support or as surveillance is set entirely by how the first one is opened, and Vilalta's single sentence is a clean model for setting that tone.

OTHER NEWS THIS WEEK

In Brief

Railway publishes its GCP account-suspension post-mortem. The full report on the 19 May outage shows a control-plane dependency on Google Cloud cascading beyond GCP, despite redundant data-plane environments. blog.railway.com

Stripe Sessions 2026 fallout continues. The Machine Payments Protocol and related agent-commerce primitives keep drawing analysis as the rails for how money moves between agents take shape. stripe.com

Brex and Ramp finance agents cross 1M conversations a month. Production agents with real action authority — blocking out-of-policy spend at swipe, auto-coding invoices — mark the credible end of the agent spectrum, alongside Capital One's \$5.15bn Brex deal. fintechbrainfood.com

EMs ask what's expected of them now. A practitioner essay argues the engineering-manager role has shifted again toward teaching and AI-stakeholder management — useful context for the management-shape pattern in this week's Quick Take. emdiary.substack.com

PLATFORM UPDATES

At a Glance

Status	Platform	What Changed
●	GitHub Copilot	Billing model changes 1 June: flat Premium Request Units replaced by metered AI Credits. Near-term budget decision for any team running agent-mode workflows.

OUTSIDE IN

incident.io: Measuring Reliability From the Customer's Outcome



Mike

incident.io engineering

28 May 2026

[Read the post →](#)

The unit is "good minutes"

incident.io measures reliability in good minutes over total minutes. Each month is divided into one-minute intervals, and a minute counts as "good" when the error rate stays below 10%. Transient, self-healing blips therefore do not register as outages — the metric tracks what customers actually experienced, not whether a component briefly wobbled. An SLO, or service-level objective, is simply the reliability target a team commits to; incident.io's is 99.99% a month for both alert ingestion (taking in the signals that trigger an alert) and notification delivery, with a five-minute latency threshold on getting a notification out.

99.99%

monthly SLO target

<5 min

notification latency

2×

active-active telecom providers

"Instead of thinking about reliability as an exercise in figuring out what we can control, and ignoring anything beyond that, we think about what we'll be really proud to offer to customers."

No third-party degradation gets to be an excuse

The sharper half of the philosophy is what incident.io refuses to write off. It runs dual active-active telecom providers — a redundancy added after a single October 2025 incident, the only AWS-driven outage in a twelve-month review otherwise near-perfect — so a degradation at one carrier is not accepted as a reason a customer missed a page. It also subtracts customer-configured delays from its latency maths, refusing to take credit for time it asked the customer to wait. The result is a reliability number measured from the outside in, against the experience promised, rather than from the inside out against the components a team happens to own.

The takeaway: defining reliability from the customer's experienced outcome — and declining to treat a vendor's degradation as an excuse — sets a higher and more honest bar than component uptime, and it is a framing an SLO or on-call review can adopt without rebuilding anything.

[Read the full post: *incident.io/blog/customers-over-control* →](#)

ON THE RADAR

Coming Up

JUN
01



GitHub Copilot AI Credits cutover — flat Premium Request Units retire and token-metered billing begins. The budget decision in this week's Lead lands on this date.

2026
HOLIDAY



Visa & Mastercard agentic commerce goes mainstream — the card networks target the 2026 holiday season for millions of consumers shopping via AI agents under the Trusted Agent Protocol and Agent Pay.

2026
Q4



Microsoft Saudi Azure region general availability — new in-country cloud capacity supporting data-residency compliance for ME bank and fintech partners.

THE LEADERSHIP READ

How Calibration Slides Toward Stack-Ranking

(3 min)



Alex Karp

LeadDev

2022 · a classic worth revisiting

[Read the article →](#)

Calibration is meant to make performance reviews fairer — managers compare ratings across a group so that one team's "exceeds expectations" means the same as another's. In a guide for LeadDev, Alex Karp argues that the process has a built-in drift, and he names it as a three-stage slippery slope. Each stage looks reasonable on its own. Taken together they carry an organisation from a fair comparison to a zero-sum contest, often without anyone deciding to make that move.

The three stages

Karp's progression runs in order, and the danger is that each step makes the next feel natural.

- 1. Dissociation and fungibility.** To cut the cost of calibration, companies abstract each person into a rating and a short summary. Senior leaders, working only from those abstractions, start to treat people as interchangeable — fungible — which quietly hollows out the point of the exercise.
- 2. Performance curves.** Ratings get fitted to a normal distribution on the assumption most people sit in the middle. This ignores how level and tenure shape a real distribution, and the harm arrives when managers are pressured to bend ratings to match the curve rather than the work.
- 3. Stack-ranking.** Performance becomes relative instead of absolute. A perfectly adequate engineer is rated poorly because peers scored higher — someone must be at the bottom —

and the process turns into a zero-sum competition that can end in termination for meeting the bar.

Despite good intentions, calibration tends to deteriorate through a slippery slope — from people abstracted into ratings, to curves, to stack-ranking, where someone must always be at the bottom.

— Alex Karp, LeadDev (paraphrased)

Seven ways to resist the drift

Against that slide Karp offers seven counter-recommendations. Two are structural and do most of the work. First, run calibrations at a local organisational level rather than rolling everyone into one giant pool, which is where fungibility creeps in. Second, eliminate performance curves entirely — the false objectivity of a forced distribution is the mechanism that produces stack-ranking, so removing it removes the slope's middle step.

The other five are about discipline and trust: include everyone in the discussion rather than only the edge cases; focus on how managers will actually communicate a review, not just the score; grant managers autonomy with light verification rather than overriding them from above; write explicit, detailed definitions of each rating so "meets expectations" is not left to interpretation; and analyse outcomes across demographic groups to catch bias the process might be encoding. None of these is exotic. The argument is that without them, the gravity of cost-cutting and tidy distributions pulls every calibration cycle the same way.

For a manager-of-managers running a cycle this year, the value is the early-warning vocabulary. Once "we just need to fit the curve" enters the room, the slope has already started, and Karp's framework gives a name to what is happening while there is still time to step off it.

[Read the full article: leaddev.com →](#)

This could be worth talking about: *In the last calibration cycle, was any rating adjusted to fit a distribution rather than to fit the work — and if so, which of Karp's two structural fixes would have caught it?*

RECOMMENDED READING

This Week's Links

The harness is the product: how four labs price the agent runtime

thenewstack.io

The banking AI bifurcation: cost takeout vs the AI-native stack

implicator.ai

The Middle East is quietly becoming a top fintech frontier

mobileecosystemforum.com

VOL. II

Engineering Digest

22–29 May 2026 · Weekly Edition · Issue #12

Next edition: 5 June 2026

Archive

1 May · Weekly

8 May · Weekly

15 May · Weekly

22 May · Weekly

29 May · Weekly