

Engineering *Digest*

Engineering decisions, every Friday.

Week of 5–12 June 2026

THIS WEEK'S BRIEFING

50:1

The manager-to-IC ratio becoming the archetype as AI absorbs the coordination layer — and the engineering-manager role splits in two.

IN THIS ISSUE

Visa and OpenAI put permission-bound card payments inside ChatGPT

UK banks and fintechs launch a shared open-banking payments scheme

Outside In: Honeycomb stopped tracking action items

The engineering-manager role is splitting in two

Leadership Read: why post-mortems fail

Platform Updates, Other News & On the Radar

FOREWORD

Control is the thread running through this week. Visa and OpenAI are defining the rules a payment agent operates inside; the engineering-manager role is fracturing around what AI now coordinates and what still needs a human; and a consortium of UK banks and fintechs has sat down to write a shared rulebook for recurring account-to-account payments. Three different rooms, one shared question — who sets the boundaries, and how tightly.

12.1

NUMBER OF THE WEEK

The average manager now carries 12.1 direct reports, up from 8.2 in 2013, while global employee engagement sits near 21% — the flattening wave that hollowed out team structures is now reaching the managers themselves, the same pressure behind this week's role-split story.

THE WEEK AT A GLANCE

Seven Stories in Sixty Seconds

Action required

Security

Chromium V8 zero-day exploited in the wild. CVE-2026-11645 was added to CISA's KEV — the Known Exploited Vulnerabilities list — on 9 June, with a federal mitigate-by date of 23 June; teams running Chrome or Chromium-based stacks should confirm the patch is rolling. [CISA KEV](#)

Security

Cisco SD-WAN flaw exploited as a zero-day. CVE-2026-20245, a critical Cisco Catalyst SD-WAN Manager vulnerability (software-defined WAN edge networking), is under active exploitation with no patch yet available; edge appliances and Manager instances are the priority for delegation and monitoring. [Cisco PSIRT](#)

Worth knowing

Payments

Mastercard launches "Agent Pay for Machines." A multi-rail settlement layer spanning cards, accounts and stablecoins for agent-to-agent transactions, with 30+ launch partners including Adyen, Checkout.com, Stripe, Coinbase and Cloudflare — the competing half of the agent-commerce trust-layer contest. [Mastercard](#)

DevEx

DORA's "adoption trap" comes into focus. New benchmark commentary contrasts ~100% AI-tool coverage with 30–40% weekly-active use — coverage is not adoption, and the ROI gap follows from the difference. [DORA](#)

Watching

Org

Salesforce describes a "truly agentic" engineering operating model. Peer-org context from the company's own newsroom — marketing-adjacent, light on operational detail, worth tracking rather than acting on. [Salesforce](#)

WIDER ENGINEERING LANDSCAPE

Visa and OpenAI Put Permission-Bound Card Payments Inside ChatGPT (4 min)

Visa and OpenAI announced on 10 June a partnership to let AI agents make tokenized, permission-bound card payments directly inside ChatGPT. Two of the largest names in payments and AI are jointly setting the default trust layer for agent-initiated commerce. This is agentic commerce at the production end of the spectrum: agents with the authority to initiate and move money, operating inside boundaries the cardholder defines in advance.

How the controls work

The mechanism is tokenization. Rather than handing an agent a card number, Visa issues a network token — a stand-in credential scoped to rules — and pairs it with real-time authorization and fraud monitoring. The cardholder sets the boundaries: spending limits, merchant categories, and required approvals. An agent can transact freely within those rails and is stopped at the edge of them.

The scoping model is best understood as three nested controls, illustrated below. (The grocery cap is the digest's illustration of how per-use-case scoping works, not a Visa-stated example.)

JARGON WATCH

Agentic commerce

Software agents that buy on a person's behalf. Here that means production agents with action authority — they initiate and move money, not just draft a suggestion a human then executes.

Tokenization / network token


A token is a stand-in card number scoped to rules, used in place of the real card number (PAN) so the underlying credential is never exposed to the agent or merchant.

Merchant-category controls

Limits that restrict where a credential can be spent — e.g. groceries but not gambling — set by category rather than per-merchant.

Codex

OpenAI's coding-agent product; named here as one surface where Visa-backed developer experiences will sit.



"Transactions will operate within clearly defined user permissions, policies and controls, such as spending limits, merchant categories or required approvals."

— Visa

The trust-layer contest

Visa is not alone. The same week, Mastercard launched "Agent Pay for Machines," a multi-rail settlement layer spanning cards, accounts and stablecoins with more than 30 partners. We covered the agent-commerce protocols reaching general availability in Issue 13; this is the consumer-distribution layer landing on top of them. The two firms are now competing on where agent payments reach people and how the permissions are expressed.

MECHANISM	APPROACH
Visa + OpenAI	Tokenized card credentials, distributed inside ChatGPT
Mastercard	Multi-rail: cards, accounts and stablecoins, across 30+ partners
Distribution	Single high-traffic surface (ChatGPT) vs. broad partner ecosystem
Settlement	Card network vs. cards / accounts / stablecoins
Control model	Spending limits, merchant categories, approvals vs. programmatic permissioning

What changes in the loyalty and commerce-media patch

For a fintech in the loyalty and commerce-media space, the standard is now being set by others. Bank customers will start asking how agent-initiated payments and per-use-case token scoping affect their rails, what data flows when an agent transacts, and where rewards and offers sit when the buyer is software. The decision in front of in-scope organisations is what to build against — and that decision is now anchored to whichever trust layer the market consolidates around.

Sources: [Visa Investor Relations](#), [SiliconANGLE](#), [Mastercard](#) (10 June 2026)

Two Stories Worth Knowing

Management & Teams

50:1

the manager-of-many archetype

CONCEPT OF THE WEEK

Levelling / career framework

The structured ladder an org uses to define what each engineering level means — scope, autonomy, impact — and how someone moves up. It is the tool a manager uses to make promotion and pay decisions defensible rather than personal.

The Engineering-Manager Role Is Splitting in Two

Meta's new applied-AI engineering team, LeadDev reports, runs with a 50-to-1 employee-to-manager ratio.

Amazon's Andy Jassy has mandated a 15% increase in the ratio of individual contributors (IC) to managers, and Gartner predicts that through 2026, 20% of organisations will use AI to flatten structures and eliminate more than half of their middle-management positions.

LeadDev's argument is that these numbers point to a role that is fracturing into two distinct jobs. One path stays close to the code — the tech-lead-manager (TLM), a manager who still owns technical direction for a single team. The other expands outward across many teams as a non-technical manager-of-many, where the job is coordination, prioritisation and people. What is being hollowed out is

Banking, Fintech & Loyalty

14+

banks and fintechs, one shared rulebook

JARGON WATCH

A2A (account-to-account)

Payments that move directly between bank accounts, bypassing the card networks.

VRP (variable recurring payments)

Pre-authorised payments that can vary in amount and timing — like a direct debit, but under open-banking permissions the payer controls.

Open-banking rulebook

A shared set of commercial and operational standards every participant follows, so the scheme works the same way across firms.

UK Banks and Fintechs Launch a Shared A2A Scheme

At Money20/20 in early June, a consortium of more than a dozen UK banks and fintechs — Barclays, HSBC, Lloyds, NatWest, Santander, Nationwide, Monzo, Revolut, Starling, GoCardless, Plaid, TrueLayer and Yapily among them — announced a shared open-banking scheme for recurring account-to-account (A2A) payments and variable recurring payments (VRP).

The scheme provides a common rulebook so recurring A2A payments behave consistently across participants, rather than each firm

the middle: the pure coordination layer — status, tracking, summaries — that agent-mode tooling inside developer platforms now absorbs. That work is being automated by agents running multi-step tasks under human triggers, not by autonomous production agents, and the span of control a manager can hold widens as a result.

The takeaway: the "engineering manager" is no longer one job, and orgs that level, hire and pay for it as one will mis-shape both tracks — the technical lead and the manager-of-many now need different ladders.

Sources: [LeadDev](#), [Fortune](#) (9 June 2026)

building its own bilateral arrangements. It is an industry-scheme launch — banks and fintechs collaborating commercially — not a regulator mandate or an FCA compliance action, which is what makes it notable: the ecosystem is organising itself around recurring payments that compete directly with cards.

The takeaway: recurring A2A is moving from pilot to shared infrastructure in the UK, and any firm whose rails or loyalty propositions assume card payments now has a credible non-card alternative forming around it.

Sources: [FinTech Futures](#), [The Paypers](#) (2 June 2026)

+98% / +91%

PRs merged rose 98% with AI assistance — but review time rose 91% alongside it. The throughput shifts; the bottleneck moves to human approval.
(Faros / LinearB benchmarks.)

IN PRACTICE

Career frameworks & ladder design (3 min)

A classic worth revisiting

Treating a Career Framework Redesign Like a Product

Megan Allison & Kelly Jacobs · Twilio, February 2024 · [Read the full piece](#)




Megan Allison & Kelly Jacobs

Redesigning Twilio's career development framework

Twilio · February 2024

[twilio.com](https://www.twilio.com) →

When Twilio set out to rework its career development framework (CDF) — the ladder defining what each level means and how people progress — Allison and Jacobs did not treat it as a policy edit. They ran it as a product, moving through the same sequence of stages a design team would use to ship anything else. The framework had users, and the redesign was built and tested against them rather than drafted in isolation and handed down.



The research stage included studying how other organisations build ladders — Figma, Intercom and the UK Government Digital Service (GDS) among them. The pattern they found across all of them was that the levels wove together increasing scope, autonomy, leadership and strategic impact, rather than stacking up craft skills alone. That diagnosis named the problem with the old CDF directly: it over-emphasised craft at the lower levels and read as a long list of boxes to tick.

"The CDF was too long and felt like an unobtainable checklist."

The redesign acted on that feedback in concrete terms. The team collapsed the behavioural categories from nineteen down to five, cutting the redundancy users had flagged, and built promotion "cheat sheets" that spelled out which skills mattered at each step and what good looked like in practice — replacing ambiguous language and opaque weighting with worked examples. The rebalance they landed on was deceptively simple — "craft, then strategy" — sequencing the two so that lower levels build craft and higher levels add scope and strategic impact, rather than demanding both at once.

The transferable mechanism is the method itself, more than any specific ladder shape. Treating levelling as a product means researching the field, prototyping a model, and testing it with the people who will be levelled against it, before it becomes policy.

WHAT TO TRY

The substitution test holds — a large engineering org ran its career-framework redesign as a product, and the method survives the swap. The next ladder revision is a candidate for the same treatment: researched and tested with its users, rather than drafted as a document and announced.

OTHER NEWS THIS WEEK

In Brief

Azure OpenAI publishes its retry-amplification post-mortem. The formal write-up of the 29 May 2026 incident details how automatic client retries amplified load — a retry storm that widened the blast radius (the set of systems an incident reaches) before traffic was shed. [Azure status history](#)

AI throughput benchmarks sharpen the review-bottleneck picture. Faros and LinearB data this week shows merged-PR volume rising far faster than review capacity, putting human approval — not generation — on the critical path. [LinearB](#)

Mastercard's agent-commerce ecosystem widens. The 30+ launch partners for Agent Pay for Machines — including Adyen, Checkout.com, Stripe, Coinbase and Cloudflare — signal that agent settlement is consolidating around a small number of multi-rail platforms. [Blockonomi](#)

OUTSIDE IN

Honeycomb Stopped Tracking Action Items — and Its Reviews Got Better

Lex Neva

LN

Incident Retrospective Ground Rules
Honeycomb · August 2024

honeycomb.io →

"If you have improvement ideas or action items, I would encourage you to note them down for later rather than raising them during this meeting." That is one of the ground rules Lex Neva uses to open an incident retrospective at Honeycomb — and it inverts how most teams run them.

The ritual that crowded out learning

The conventional retrospective drives toward a list of action items. Neva's argument is that doing so during the review pulls attention away from the harder, more valuable work: actually understanding what happened. When a team rushes to fixes, it stops asking why the incident unfolded the way it did. So Honeycomb's ground rules deliberately defer the action-item hunt.

"Once we have a better understanding of this incident, action items usually surface organically." The team found that the items they would have rushed to capture were mostly things they would have done anyway — and the ones that mattered emerged once the understanding was in place.

Removing a ritual, not adding one

The point worth sitting with is that the fix wasn't a better template — it was removing the ritual that crowded out learning. Honeycomb improved its incident reviews by taking something away, not by adding process. A post-mortem culture is shaped as much by what a team refuses to do in the room as by the artefact it produces afterward.

The takeaway: a peer org improved its incident reviews by deferring the action-item hunt until the team understood the incident — a reminder that the artefacts a process produces and the learning it generates are separate things.

[Read the full post: honeycomb.io →](#)

ON THE RADAR

Coming Up

JUN
23



Chromium V8 patch deadline — CVE-2026-11645 federal mitigate-by date under CISA KEV. Confirm Chrome/Chromium rollout is complete.

JUN
10



Agent-commerce ecosystem momentum — Visa–OpenAI and Mastercard's Agent Pay for Machines both landed this week; expect bank-customer questions to follow.

JUN
ONGOING



Cisco SD-WAN active exploitation — edge appliances remain the priority; track PSIRT advisories.

JUN
TBC



Azure dedicated-routing migration — completion targeted this month, following the 29 May retry-amplification incident.

THE LEADERSHIP READ

Why Post-Mortems Fail (5 min)

The post-mortem has quietly become one of engineering's most performed and least useful rituals. incident.io's argument is direct: post-mortems underperform because the document drifted from being written for people to being filed for process — completed to prove an incident was handled, then rarely read again.



incident.io engineering

The post-mortem problem

incident.io · 2026

[incident.io →](#)

The distinction that matters

The piece's central distinction is between two things a post-mortem can be. As communication, it is a narrative written by people, for people — an honest account of what happened, why it unfolded the way it did, and who owns what comes next. As a compliance artifact, it is a form filled in to satisfy a process, complete the moment the boxes are ticked. A blameless post-

mortem — one written to understand rather than to assign fault — only delivers value if it is also read.

"The actual problem is that somewhere along the way, the post-mortem stopped being a piece of communication and became a compliance artifact."

— incident.io

What good looks like

incident.io draws the contrast through concrete examples. The 2024 CrowdStrike incident — which the post describes as "bad content, a missed validation step, deployments that should have been staged but weren't" — is offered as a clean case where the interconnected chain of failures matters far more than naming any single root cause. The piece reaches for the Swiss cheese model to explain why: human systems are layers of cheese, each with random holes, and an incident happens only when a threat passes all the way through. A post-mortem that hunts for one root cause therefore misrepresents how incidents actually happen, because it stops at the first hole rather than mapping the alignment of several.

The fixes the piece argues for follow from the communication framing rather than from process. The first is timing: the best post-mortems, the author contends, are written while the incident "still stings a bit," because that is when people remember what they were actually thinking, not just what they did. The second is form — "tell a story, not a log" — walking a reader chronologically through what the team knew and when, so the document explains why rather than merely recording sequence. The third is specificity of ownership: "improve monitoring" is not an action, whereas a named owner adding a named alert by a named date is. Each fix is a way of writing for a future reader, not a checkbox added to a template.

Underneath the mechanics sits a distinction the post is careful about: the line between context and blame. "'Alex deployed the change' is context. 'Alex should have known better' is blame," the author writes — one describes what happened, the other assigns judgment. A blameless review, on this reading, is not one that omits names; it is one that records what people did and knew without grading their character, which is precisely what makes the account safe enough to be honest, and honest enough to be worth reading.

It is worth noting the throughline with this week's Outside In: Honeycomb improved its reviews by changing how it treats the artefact, not by adding tooling — a separate piece arguing, from the other end, the same underlying claim incident.io makes here. The value of a post-mortem is whether someone learns from it later, not whether the template was completed.

[Read the full article: incident.io →](#)

This could be worth talking about: *Was the team's last post-mortem written to be read by someone who wasn't in the incident — or to be filed?*

RECOMMENDED READING

This Week's Links

The engineering manager role is splitting in two [leaddev.com](#)

The post-mortem problem [incident.io](#)

Incident retrospective ground rules [honeycomb.io](#)

Redesigning our career development framework as a product [twilio.com](#)

2026 engineering benchmarks: AI throughput and review load [linearb.io](#)

PAPER

DOWN THE RABBIT HOLE

Span of Control: What's the Optimal Team Size for Managers?

A tangent off this week's role-split thread: Gallup finds the optimal span of control isn't a fixed number. The strongest predictor of engagement is whether a manager holds one meaningful conversation per report each week — and that holds regardless of how large the team gets.

VOL. II

Engineering Digest

Week of 5–12 June 2026 · Weekly Edition · Issue #14

Next edition: [Week of 12–19 June 2026](#)

Archive

15 May · Weekly

22 May · Weekly

29 May · Weekly

5 Jun · Weekly

12 Jun · Weekly