

Engineering *Digest*

Engineering decisions, every Friday.

26 June – 3 July 2026

THIS WEEK'S BRIEFING

17% → 43%

The share of CEOs planning to cut entry-level roles jumped in a single year as AI restructuring accelerated — while the senior bench those roles feed goes unbuilt.

IN THIS ISSUE

The AI-washing layoffs and the pipeline that isn't being built

The flat-rate era of AI coding is over

Open USD: the payments rail rewrites its own economics

In Practice: Atlassian rebuilds its operating model

The Leadership Read: Will Larson's revised rules

Outside In, Other News & On the Radar

FOREWORD

A technology becomes real to a business the moment it starts rewriting the org chart, the budget line, and the rails underneath both — and this week it did all three at once.

Employers cut roles and pinned the reason on automation. Coding-tool vendors quietly retired flat-rate pricing. A 140-firm consortium moved to change who earns the interest on the money the industry settles. Three different layers of the same shift, arriving in the same seven days.

\$3.4bn
NUMBER OF THE WEEK

The Middle East consumer-loyalty market is forecast to reach roughly \$3.4bn in 2026, growing about 15% a year toward \$5.6bn by 2030 — a reminder that the regional commerce-media patch loyalty and fintech firms operate in is expanding fast, led by digital wallets and payment-linked rewards in Saudi Arabia and the UAE. [Middle East Consumer Loyalty Business Report 2026](#) →

THE WEEK AT A GLANCE

The week in brief

Action required

Security

Oracle PeopleSoft zero-day still being exploited. CVE-2026-35273, a pre-authentication remote-code-execution flaw in PeopleTools rated CVSS 9.8, has been used by the ShinyHunters extortion crew to breach more than 100 organisations — 300-plus instances, higher education hardest hit. Teams running PeopleSoft should confirm Oracle's out-of-band patch is applied and check for compromise. [The Hacker News](#), [Help Net Security](#)

Security

Chrome's emergency V8 fix covers a zero-day still being exploited in the wild. CVE-2026-11645, an out-of-bounds memory bug in Chrome's V8 JavaScript engine (CVSS 8.8), is the fifth Chrome zero-day of 2026 and remains on CISA's Known Exploited Vulnerabilities list. Everyone runs a browser, so this stays a fleet-wide update to confirm. [The Hacker News](#)

Worth knowing

Model release

Anthropic ships Claude Sonnet 5. Launched 1 July and billed as its most agentic Sonnet model — able to plan and run tools autonomously at a level that recently needed larger models — at an introductory \$2 / \$10 per million input/output tokens through August. The relevance for leaders is cost: capable agentic work is getting cheaper per token even as usage climbs. [Anthropic](#)

Peer data

Spotify reports 99% weekly AI-tool adoption. Its chief architect's claim that "coding is no longer the constraint" is circulating as a headcount-math data point — if the bottleneck moves from writing code to reviewing it, team-shape assumptions move with it. [Spotify Engineering](#)

Watching**Security**

Arista EOS flaw will not be patched. CVE-2026-7473, an exploited tunnel-decapsulation bug now on CISA's KEV list, has no software fix planned — Arista's guidance is ACL-based mitigation. A useful reminder that "apply the patch" is not always an option and mitigation posture has to be designed. [SecurityWeek](#)

Commerce media

"Banks as media" reaches the mainstream. Card-linked offers are pushing banks and fintechs into retail-media territory, with financial-media-network ad spend forecast to grow at triple digits from a small base. The loyalty and commerce-media vertical is crossing from experiment to strategy. [The Financial Brand](#)

MANAGEMENT & TEAMS

The AI-washing layoffs and the pipeline that isn't being built

(3 min)

CONCEPT OF THE WEEK

The junior / apprenticeship pipeline

Engineering seniority is grown, not bought. Graduates learn by doing the

JARGON WATCH

AI-washing

Dressing an ordinary cost-cut in AI clothing — announcing job reductions as "AI-driven efficiency" when the real driver

small, well-scoped work, then progress into judgement roles over three to five years. Freezing junior hiring today does not just save this year's salary line — it removes the bottom of a conveyor belt whose top is where next decade's senior engineers come from.

is margin pressure. Also called AI redundancy washing. It inflates how much work AI has actually automated.

System verifier

The emerging senior role that reviews, corrects and takes ownership of what AI systems produce. It is a job that assumes years of accumulated judgement — the kind that used to be built by doing the junior work now being automated.

Seventeen per cent to forty-three per cent in a year

The share of CEOs planning to reduce junior and entry-level roles rose from 17% to 43% in a single year, according to Oliver Wyman — a near-tripling that lands as two large employers restructured in the same week. On 29 June, British American Tobacco announced it would remove or outsource around 9,000 roles, close to a fifth of its workforce, in an explicitly AI-driven cost programme. Days earlier, WARN filings showed Cisco cutting 471 Bay Area jobs, effective 13 July, as part of a restructuring the company frames around its pivot to AI.

What "AI-cited" actually counts

The headline numbers deserve a careful read. Challenger, Gray & Christmas has recorded AI cited in more than 100,000 US job-cut announcements so far in 2026, the leading stated reason for four consecutive months. But "AI-cited" is a mixed bag: it lumps genuine automation — work a system now does end to end — together with cost-cutting that borrows AI as a more palatable public explanation than "demand softened." Some of these roles were automated. Some were going regardless, and AI is the cover story. The two are not the same, and the survey framing rarely separates them.

"Stop hiring juniors today, and the senior verifiers of 2029 never get trained."

— Engineering Digest


The seniorised trap

PwC's 2026 Global AI Jobs Barometer surfaces the mechanism underneath the layoffs. Entry-level work has not vanished so much as hardened: roles that survive are being "seniorised,"

rewritten to demand skills a graduate cannot yet have. AI-exposed junior roles are now seven times more likely to require traditionally senior capabilities than the least-exposed ones. Since 2019 these seniorised entry roles have grown 35%, while other entry-level roles have fallen 10%. The ladder's bottom rung is being sawn off and reattached higher up.

The bottom rung is being cut — and rebuilt higher up

Entry-level hiring: two data sources, one direction of travel



Data: [Oliver Wyman](#) (CEOs planning to cut entry-level roles) and [PwC 2026 Global AI Jobs Barometer](#) via [Fortune](#) (seniorisation).

Why leadership owns this one

The decision here sits squarely with engineering leadership. The trade-off engineering leaders are now navigating is between a real short-term saving and a structural gap that only shows up years later. If entire cohorts stop entering the profession while AI absorbs the tasks juniors used to cut their teeth on, the "system verifier" seniors those AI systems will require in 2029 have nowhere to come from. The firms treating this as purely a headcount decision are the ones most exposed to it.

KEY TAKEAWAYS

- The share of CEOs planning to cut entry-level roles rose from 17% to 43% in a year (Oliver Wyman); Cisco (471 Bay Area roles) and BAT (~9,000) both restructured this week citing AI.
- "AI-cited" cuts conflate genuine automation with cost-cutting dressed as automation — the two should not be read as one number.
- Surviving entry roles are being "seniorised" (+35% since 2019 vs -10% for other entry-level work), which quietly closes the on-ramp that grows senior engineers.

QUICK TAKES

Two Stories Worth Knowing

Ways of working

\$120

Cursor's new per-seat price for running agents all day

The flat-rate era of AI coding is over

A quiet reality check landed at the start of July: the fixed-price seat for AI coding is being retired across the board, and heavy agentic use — the multi-step, tool-running work where a coding agent edits files and executes tasks on the developer's behalf — is where the meter now spins fastest. GitHub Copilot moved all plans to usage-based billing on 1 June, metering consumption through "AI Credits" on top of the base seat. Cursor added a \$120-per-seat tier aimed at developers running agents all day, with overages billed at four cents a request. Claude Code's Agent SDK now draws on separate metered credit pools for automation-scale work.

Metered, or usage-based, pricing means the invoice tracks consumption — tokens, requests, or agent time — rather than a flat monthly licence. That turns a predictable per-

Wider landscape

140+

Firms launching Open USD

Open USD: the payments rail rewrites its own economics

"It's a stablecoin built for the internet economy, designed by the businesses growing it," said Zach Abrams, interim CEO of the Open Standard, as more than 140 firms — Visa, Stripe, Mastercard, BlackRock, DBS, Standard Chartered among them — unveiled Open USD (OUSD) on 30 June. A *stablecoin* is a token pegged to a currency, here the dollar; the twist is the economics. Where issuers like Circle keep the interest earned on reserves, OUSD returns most of that reserve revenue to participants, minus a small fee. Circle, whose USDC model keeps that float, fell 13% to \$66 on the news; USDT holds roughly 62% of the market and USDC about 25%.

head cost into a variable one that rises with how hard teams lean on agents.

TOOL	FLAT ENTRY → METERED CEILING
Claude Code	\$20 Pro (included) → \$200 Max; Agent SDK now on separate metered credits
Cursor	\$20 Pro → \$120/seat "agents all day"; overages \$0.04/request
GitHub Copilot	\$10 Pro → usage-based AI Credits from 1 Jun 2026

The decision this forces is a finance one. Someone now owns the questions of metering, per-team caps, and who carries the overage when a developer runs an agent overnight against a large repository.

The takeaway: AI coding has moved from a fixed per-seat licence to a variable cost centre — and the leaders who get ahead of it are opening the budget conversation about caps and ownership before the first surprise invoice, not after.

Sources: [Developers Digest](#), [GetDX](#)

The takeaway: the settlement rail that banks and their fintech partners run on is changing who earns the float — and a revenue-sharing stablecoin backed by the networks themselves reshapes the merchant and loyalty economics that sit on top of it.

Sources: [Fortune](#), [CoinDesk](#)

IN PRACTICE

Ways of working in the agent era — operating-model change after adoption (2 min)

Avani Prabhakar

Rebuilding the operating model, not just the toolkit


Avani Prabhakar · Inside Atlassian, June 2026 · [Read the full piece](#)

Atlassian's newly created Chief People and AI Enablement Officer makes a claim most adoption stories skip: getting AI into an organisation is a change to how the company is run, and the work does not end when the tools are deployed. Her framing is human first, AI second, and the barrier she names is mindset, trust and confidence rather than access to tools. What makes the piece a useful In Practice read is that Atlassian built explicit operating-model machinery to move people through that change, rather than hoping usage would spread on its own.

Two mechanisms do the load-bearing work. *AI Working Agreements* are team-level documents — as Prabhakar puts it, "shared guidelines and principles that define how a team should use AI safely and where they should not use AI." *AI Retros* are regular sessions where teams assess what AI genuinely helped with and what it did not, then share those findings more widely. Neither is a tool; both are rituals that give a team permission to experiment inside agreed boundaries and a structured way to learn out loud.

We've evolved our measurement framework to focus on "super users" — those who are using AI to actually drive positive change.

The measurement shift is the transferable part. Atlassian tracks *super users*: those using advanced AI features at least 40 times a week inside their real workflows. Around 30% of staff clear that bar today, against near-universal daily use of its assistant. That funnel — everyone using it lightly, a third using it deeply — is a truer adoption signal than a raw "99% have tried it" figure.



WHAT TO TRY

A grounding question for any team past the "everyone has access" milestone: what share are genuinely super users versus occasional dabblers — and whether the operating model has rituals like working agreements and retros to move people from one to the other, or just a licence and a hope.

In Brief

Visa settles the first live agentic-commerce payments in Europe. On 2 July, Visa and banks across Europe reached the next phase of agentic commerce — a live, SCA-compliant transaction completed with Worldline and ING, and merchants including lastminute.com, Frasers and BrickDepot, routed through Visa's Trusted Agent Protocol and 30-plus European issuers. [Visa newsroom](#), [Worldline / ING / Visa](#)

Mastercard extends Agent Pay to machine-to-machine payments. Building on last month's coverage, Mastercard's Agent Pay for Machines lets AI agents transact autonomously across its network — including microtransactions worth fractions of a cent — with settlement across cards, accounts and stablecoins. Context for how the agentic-commerce rails are filling in. [Fortune](#)

OUTSIDE IN

Run out to meet it: craft when the ground keeps shifting (2 min)

KB

Kent Beck

Author of *Tidy First?*, in conversation with Charity Majors
newsletter.kentbeck.com, 2026

[Read the post →](#)

The durable stance underneath

Kent Beck's "Run Out to Meet It" is, on its surface, a conversation about a period of upheaval in software work. Read for its durable lesson, it is about a stance any engineering leader recognises: what to do when a discontinuity arrives and years of hard-won expertise suddenly count for less than they did. The title is the argument. When change is coming regardless, the move is to run out and meet it. Holding position is usually the riskier call.

The piece's sharpest framing is that everybody's ignorance has been reset to 100. The engineers who built careers on beautiful, readable code, it argues, are struggling more than anyone — while newer developers may hold a hidden advantage, because nobody has a head start on ground this new. When the terrain changes underneath it, deep expertise can weigh a person down.

Why it travels beyond one moment

Strip out the specific occasion and the pattern is one leaders face repeatedly — a platform shift, a market break, a reorganisation that resets who knows what. The instinct of the most experienced people is to protect the expertise that made them senior, which is precisely the reflex that leaves them flat-footed. The teams that adapt are the ones that treat a reset as a shared starting line, where the person most willing to look ignorant and relearn moves fastest. For an engineering leader, the transferable discipline is protecting the space for senior people to be beginners again without losing status for it.

The takeaway: durable craft is the willingness to meet a reset head-on — and to build a team culture where relearning carries no penalty, because the next discontinuity is always on its way.

Read the full post: newsletter.kentbeck.com →

ON THE RADAR

Coming Up

JUL

13



Cisco Bay Area cuts take effect — the 471 WARN-noticed roles land on 13 July, the concrete edge of the AI-restructuring wave this week's lead traces.

NOW
READ



DORA's "ROI of AI-Assisted Development" framing — the argument that AI's return shows up as capacity, not headcount reduction, is worth having to hand as the layoff-versus-productivity debate sharpens.

WATCH
RISK



Shared-transit concentration risk — the recent Cloudflare fibre-cut episode is a prompt to check how much of a service map runs through a single upstream transit layer that nobody owns.

WATCH
CRAFT



Calibration in the AI era — a practitioner cluster is questioning whether velocity-based calibration still holds once agent-assisted output inflates raw throughput. One to track before the next review cycle.

Will Larson's revised rules of engineering leadership (3 min)



Will Larson

lethain.com · author of *An Elegant Puzzle* and *The Engineering Executive's Primer*

15 June 2026

[Read the article →](#)

Will Larson has spent years publishing durable rules for running engineering organisations, and his latest post revises them for how the work has changed. The striking thing about the revision is how much of the craft it leaves untouched. The tools move faster, and the leadership fundamentals — durable teams, deep context, and good decisions made quickly — grow more valuable as they do.

Five revised rules

Larson names five. Migrations can now be done by an individual rather than a team. First-pass code is nearly free, but the cost of *working* code depends on the development harness and is not free. The base case of process should be optimised for agents. Durable, high-ownership teams with domain context matter even more. And quick, good, durable decision-making is a prerequisite to meaningfully benefit from any of it.

Two of the distinctions repay attention. The first is the gap between first-pass and working code: producing a draft is cheap, but the expensive, load-bearing part is the harness — the tests, review, and verification that turn a draft into something safe to ship. The second is his warning that high-judgement individuals still get hemmed in by a lack of domain context. Speed of output does not substitute for knowing the system; a team that has lived with its own problem space makes better fast decisions than a talented outsider can.

Your team and company can only benefit from this increased pace of execution if you can make durable decisions quickly, and those decisions are good.

— Will Larson, lethain.com

The management craft underneath

Read as leadership craft, the through-line is old-fashioned in the best sense. Larson's argument lands on decision quality and team durability — the ability to decide well and fast, and to keep hold of the domain knowledge that makes those decisions good. Faster execution amplifies whatever decision-making sits behind it; a team that decides badly now simply decides badly at higher speed. The rules that changed are about the tooling. The rule that carries the most weight is an old one: build teams that own their context and can commit to a call without thrashing.

[Read the full article: lethain.com](#) →

This could be worth talking about: *If faster execution only pays off when decisions are quick, good, and durable, which of those three is the team currently weakest on — and is the constraint the speed of deciding, the quality of the call, or how often decisions get reopened a month later?*

RECOMMENDED READING

This Week's Links

Entry-level work didn't disappear — it 'seniorized'	fortune.com
Software engineers top Cisco's Bay Area WARN notices	techtimes.com
Stripe, Visa and 140+ firms launch Open USD	fortune.com
How Atlassian is driving effective AI transformation	atlassian.com
Run out to meet it	newsletter.kentbeck.com
Revised rules of engineering leadership	lethain.com

VOL. II

Engineering Digest

26 June – 3 July 2026 · Weekly Edition · Issue #17

Next edition: 10 July 2026

