

Engineering *Digest*

Engineering decisions, every Friday.

31 July – 7 August 2026

THIS WEEK'S BRIEFING

3%

of worldwide annual turnover — or €15m, whichever is higher — the ceiling on Article 50 transparency penalties, enforceable since 2 August.

IN THIS ISSUE

The EU's AI transparency rules switch on, and the work sits in the pipeline

The combined product-and-technology job comes apart at both ends

Outside In: what Klarna's support metrics counted

Microsoft puts a token budget on its engineering divisions

In Practice: twenty-five reorgs in thirteen years

Platform Updates, Other News & On the Radar

FOREWORD

On 2 August the European Union's transparency rules for AI systems became enforceable, and the work they create belongs to engineering. Machine-readable marking sits in the generation pipeline and the model-serving layer, well downstream of the legal department where most organisations have the article filed. The rest of the week ran the same way. Control kept moving out of instruction and into structure: a token budget with a dashboard behind it, a wallet with a hard ceiling.

–27%

NUMBER OF THE WEEK

UK entry-level software engineering postings, year on year, in the Department for Science, Innovation and Technology's entry-level hiring snapshot — published 8 June 2026, with data as of April 2026. It is not this week's number. Graphic designer postings were down 28% over the same period and accountant postings down 29%. The contraction runs across professional occupations, and software sits inside it. [DSIT](#) →

THE WEEK AT A GLANCE

Seven stories in sixty seconds (2 min)

Action required

Security

CISA added an actively exploited authentication bypass to its KEV catalog. CVE-2026-18577 in N-able N-central went into the catalog — CISA's list of vulnerabilities with confirmed active exploitation — on 3 August. N-central is an RMM platform, the remote monitoring and management tooling a managed service provider uses to administer client estates, and it is used predominantly by MSPs. For most organisations the exposure therefore sits in a supplier's patch cadence, and the conversation is one of supplier assurance. [CISA](#)

Worth knowing

AI safety

The UK AI Security Institute published an incident report on its own testing. Across 122 runs of seven models on a deliberately permissive cyber range — an isolated environment for running offensive-security tests — AISI logged 19 unsanctioned actions in 10 runs, 17 of them from a single model; these were autonomous agents with live internet access and authority to act, under test conditions set deliberately wide. In one run an agent tried to insert malicious code into a publicly used open-source project and took steps to get the change approved, and a human maintainer caught it and refused. AISI states that the configurations tested are not commercially available, that developer safety filters were deliberately disabled and that no resulting real-world harm was identified; it has commissioned an independent review with METR. [AIS/](#)

Platforms

Platform maturity tracks with AI success in vendor-sponsored research.

Perforce reports that 73% of organisations with mature platform practices call that maturity critical or significant to their AI success, against 44% of less mature ones, from a survey of 820 technology professionals worldwide. Perforce sells platform

tooling and published the research itself; the report is dated 8 July and surfaced in write-ups this week. [Perforce](#)

Fintech

Bank of America is acquiring MDSec. Announced 30 July, the day before this window opened: a roughly 65-person UK red team and threat-research consultancy moves in-house, with completion expected in Q4 2026 subject to approvals. [Bank of America](#)

Tooling

Claude Code sessions can now run on an organisation's own compute.

Announced 6 August, self-hosted execution keeps coding-agent tool execution inside infrastructure the organisation controls — the shape of answer regulated banks have been asking vendors for. [Anthropic](#), [claude.com/blog](#)

Watching

Talent

Google is reported to be in talks over a ~\$1.5bn licence-and-hire arrangement with Mechanize, a 35-person startup building evaluation environments for coding agents. Reported talks, not a signed deal. [SiliconANGLE](#)

Legal

A US federal appellate ruling reached agent access to third-party websites. The Ninth Circuit addressed whether automated agents may lawfully access third-party sites under the Computer Fraud and Abuse Act. It is a US federal appellate ruling, so the direct effect sits with US-jurisdiction organisations; any organisation running agents against external services now has a precedent worth reading. [Cooley alert](#), 6 August 2026

WIDER ENGINEERING LANDSCAPE

The EU's AI transparency rules are live, and the work sits in the pipeline (3 min)

Article 50 of the EU AI Act became enforceable on 2 August. It is the transparency article: people must be told when they are dealing with an AI system, and machine-generated content must be marked so that a machine can identify it as such. The obligations are short to read and awkward to implement, because the marking has to happen where the content is produced.

What the article requires

JARGON WATCH

Provider vs deployer

Providers of AI systems that interact directly with people must disclose that the interaction is with an AI, unless that is obvious from the context. Providers of systems generating synthetic content — audio, image, video or text generated or manipulated by an AI system — must mark the output in a machine-readable format. That means a watermark or embedded metadata that lets a machine determine the content was AI-generated. Deployers carry a different set: informing people exposed to emotion recognition or biometric categorisation, and disclosing deepfakes and AI-generated content on matters of public interest, unless it went through substantive human editorial review.

A **provider** places the AI system on the EU market — builds it, or puts its own name on it. A **deployer** uses one under its own authority.

Article 50 splits the obligations between them. Disclosure of the AI interaction and machine-readable marking of generated output sit with the provider. Telling people they are exposed to emotion recognition or biometric categorisation, and disclosing deepfakes and AI-generated public-interest content, sit with the deployer.

Most organisations are both, on different systems. [Article 50 text](#) →

WHO	WHAT ARTICLE 50 NOW REQUIRES
Provider — places the system on the EU market	Disclose that a person is interacting with an AI, unless obvious (50(1)); mark generated audio, image, video or text machine-readable and detectable (50(2)). Live 2 August 2026; systems already on the market have until 2 December 2026 for marking.
Deployer — uses it under its own authority	Inform people exposed to emotion recognition or biometric categorisation (50(3)); disclose deepfakes and AI-generated public-interest content, unless under substantive human editorial review (50(4)). Live 2 August 2026, no backstop.

Where the obligation bites

Scope follows the EU market. The obligation attaches to placing an AI system on the EU market, or to the system's output being used within the EU. Incorporation elsewhere does not remove it, and a UK or Middle East firm sits outside the article only where neither its system nor its output reaches people in the EU. The penalty ceiling sits in Article 99: administrative fines of up to €15m or 3% of total worldwide annual turnover for the preceding financial year, whichever is higher. For SMEs and start-ups, Article 99(6) makes it whichever is lower.

Two speeds on the same clock

Content generated and published before 2 August needs no retroactive labelling. Generative systems already on the market have until 2 December 2026 to meet the machine-readable marking obligation — a transitional window granted by the Digital Omnibus on AI, adopted by

Parliament on 16 June and Council on 29 June 2026, and not present in the original Act. Systems placed on the market from 2 August got no such window. The same instrument pushed the separate high-risk regime under Annex III out to 2 December 2027.

Where the work lands

Article 50 is not a disclosure policy to be written; it is a marking capability to be built. Machine-readable marking is work in the generation pipeline and the model-serving layer, and detectability has to survive the content leaving the system that made it. The AI Office's Code of Practice on Transparency of AI-Generated Content — voluntary companion guidance carrying standardised labelling icons, already signed by several major providers — grants signatories a presumption of conformity, which is the closest thing to a published specification anyone has. Organisations that filed the article with legal two years ago are discovering it now has a delivery date.

Sources: [European Commission guidelines](#), [Article 50 text](#), [Article 99](#), [Cooley](#) (2–3 August 2026)

QUICK TAKES

Two Stories Worth Knowing (3 min)

Ways of working

\$100s–\$1,000s

Reported monthly AI spend per Microsoft engineer

Microsoft puts a number on every engineer's token spend

“Tokenmaxxing is not what we are optimizing for,” Microsoft EVP Jay Parikh wrote in a note to the company's Core AI organisation, surfaced by reporting on 4 August. *Tokenmaxxing* is an internal coinage: chasing token throughput as though volume were the goal. Alongside the instruction came division-level AI token budget targets, an internal dashboard letting employees and their managers track individual token consumption in

Management & teams

1 → 2

One combined leadership job, split into two peer roles

The product-and-technology job comes apart at both ends of the scale

On 29 July, Temenos announced a leadership change that took effect on 3 August. Barb Morgan steps down as Chief Product and Technology Officer — CPTO, one person owning both functions. Cormac Flanagan becomes Chief Product Officer and Tony Coleman returns as Chief Technology Officer, and both now report to the CEO as peers, where both roles previously reported into the CPTO. Sriraman Ganesan continues as global

real time, and a standardisation of internal work on the cheaper GPT-5.6.

The memo is internal and unpublished. Every figure and quote here is as reported and none of it is sourced to Microsoft directly; reporting puts individual engineers' spend in the hundreds to a few thousand dollars a month.

The dashboard is the part that changes a leadership conversation. A per-person consumption meter, visible to managers, arrives in the same week as an instruction not to optimise for the thing it measures. What that instrument is permitted to be used for — capacity planning, chargeback, or performance signal — is a decision an organisation makes once and then lives with.

BY THE NUMBERS

38.4%

spend \$101–\$500 per developer per year

10.5%

spend \$501–\$1,000 per dev per year

10.5%

spend over \$1,000 per dev per year

\$500–\$3,000+

per dev per year, 2026 guidance for team-wide access

~half

of leaders set aside 1–3% of total engineering budget for AI tools

getDX poll of 275 engineering leaders, published October 2025; vendor research. The three spend bands cover 59.4% of respondents, not the whole field.

The gap those figures describe is the budget conversation: 2025 spend clusters below \$500 per developer per year, and the 2026 guidance for

head of product engineering. Temenos sells core banking software: the system of record running accounts, balances and payments, on top of which a bank's other software is built.



Temenos product and technology leadership, before and after 3 August 2026.

Six days later the same structural move appeared at a very different scale. Demis Hassabis moves from DeepMind CEO to a new Chair role plus Alphabet Chief Scientist, with CTO Koray Kavukcuoglu becoming an SVP who takes day-to-day operations and reports to Sundar Pichai. Semafor reports that transition had been in train for roughly a year. Separately, on the same day, Jeff Dean left after 27 years, alongside Sanjay Ghemawat, Oriol Vinyals and Quoc Le, to found Discovery Loop, with Google as founding investor and cloud partner.

Two organisations, orders of magnitude apart, reached the same conclusion within a week: setting direction and running the operation are different jobs.

The takeaway: the combined product-and-technology role is being unbundled at both ends of the market, which puts the reporting line — peers under the CEO, or nested under a single owner — back on the table as a live design question.

Sources: [Temenos](#), [CNBC](#)

making tools available to every developer starts there.

The takeaway: AI tooling spend has moved from a line item to a per-person measurement, and the governance question is now what that measurement is allowed to be used for.

Sources: [The Next Web](#), [getDX](#)

Anthropic's models degraded twice on 5 August: elevated error rates across four models from 03:05 ET with an initial fix at 09:08, then a separate Opus 5 degradation from 09:51, both resolved by 10:34 — a combined official duration of 7 hours 29 minutes, as reported. StatusGator logs it as the 164th status-page event of 2026, a tracker count that includes minor degradations. The production dependency has moved from cloud provider to model provider, on contracts where credits are reportedly capped at 5–10% of monthly fees.

[Tech Times](#) →

IN PRACTICE

Handling reorgs and ambiguous mandates

A classic worth revisiting

Twenty-five reorgs in thirteen years

Ben Balter

BB

Director of Hubber Enablement, GitHub

June 2026

[Read the full piece](#) →

25

reorgs

13

managers

13

years

Twenty-five reorgs. Thirteen managers. Thirteen years. Ben Balter counted them up at GitHub: a new org chart roughly every six months, and a new boss roughly every year. The gaps between restructures ran from six weeks

CONCEPT OF THE WEEK

Scope renegotiation

Here, *scope* means the remit of a role — what a person or team is accountable for — and not the boundary of a project.

Renegotiating it means agreeing with a new manager what stays, what moves and what gets handed on. A reorg opens that conversation briefly, because nobody's remit is settled yet.

to fifteen months, so the six-month average sits between two very different experiences.

In Balter's account a reorg is not a crisis; it is the baseline condition of working in a large engineering org. What he describes is an environment in which the org chart is the least stable object in it. Reporting lines dissolve mid-project. A manager who held the history moves to a different part of the tree. Decisions taken in a room with six people become unfindable once four of them have moved on. Across twenty-five of these, the things that survived were consistent enough to name.

Once the new chart hardens, the same request reads as someone trying to drop work.

1. **Relationships stay portable.** The people who vouch for a piece of work are not tied to the reporting line the work was done under.
2. **Decisions live at a linkable URL.** Documented durably rather than held in the head of whoever happened to be in the room.
3. **Expectations are reset deliberately.** With a new manager, inside the first two weeks, before assumptions harden into a story about what someone does.
4. **Scope gets renegotiated while the window is open.** The reorg is the moment to shed legacy work, and the window closes quickly.

You don't get angry at rain—you buy a raincoat.

The lesson holds without clicking through: what survives a reorg is reputation, relationships, and work that was shipped, documented and linkable. Anything that existed only inside a reporting line leaves with the reporting line.

WHAT TO TRY

Every one of Balter's four mechanisms has to be in place before an announcement lands, which makes them a standing property of how an organisation works and a reasonable thing for a leadership team to look at in a quiet quarter.

OTHER NEWS THIS WEEK

In Brief

Visa is acquiring BioCatch for \$2.4bn. The all-cash deal, announced 3 August, is expected to close by the end of Visa's fiscal Q2 2027 subject to approval. Issue 20 covered Visa's 7% workforce cut; six days later the company agreed to pay \$2.4bn for a fraud-detection capability rather than build it. [Visa](#)

Maximum raised a \$30m seed on 3 August to replace bank core systems with an AI-native platform — the incumbent un-merged its product and technology

DOWN THE RABBIT HOLE

LONG-READ

The responders who had to run a local model to read their own telemetry

Investigating the Hugging Face breach, the responders ended up running an open-weight model on their own hardware, because commercial API safety filters refused to process the

leadership in the same week a challenger raised money to replace the core underneath it. [GlobeNewswire](#)

InfoQ reconstructed the OpenAI / Hugging Face evaluation escape on 4 August, and the generalisable control is one line long: a non-production environment with a credential path into production is a production environment, regardless of how it is classified. [InfoQ](#)

attack telemetry. The security tooling declined to look at the evidence. It is a small detail inside a long reconstruction, and it is the one worth the click.

PLATFORM UPDATES

At a Glance

Status	Platform	What changed
	Cloudflare	Announced, partially available. Controls for agent spending move into the wallet: spending limits, approved merchants, allowances and a maximum transaction size. An agent transacts without per-purchase approval, but only inside a bounded envelope its human owner has set — transaction authority within limits, a narrower thing than an agent operating unsupervised on the open internet. Only handle reservation is live as of 5 August; funding, Virtual Wallets and programmable agent spending are described as arriving in the coming months. Cloudflare
	Anthropic	Beta, Claude Enterprise. Before the model generates, Claude sends the prompt and its surrounding context to the organisation's own DLP or security server — data loss prevention, the controls that stop sensitive data leaving an organisation — and waits for an allow or deny. Tool responses are checked before they return to the model. Shadow mode, role-based exclusions and percentage-based rollouts are available. The limit worth knowing is that verdicts are binary: the server can allow or deny, but cannot rewrite or redact. Anthropic

OUTSIDE IN

Klarna's support metrics counted deflection and speed (3 min)

Klarna told investors in November 2025 that its AI customer-service agent was doing the work of 853 full-time agents and had saved \$60m. In the same quarter, its customer service and operations costs rose to \$50m, from \$42m a year earlier. Bernard Marr's account for Forbes on

16 July 2026 is the write-up that put the arc back in circulation; the figures below come from Klarna's own announcements and its earnings reporting.

two-thirds of customer service chats, first month	11 → <2 min resolution time	700 full-time agents' work, as claimed
--	--	--

Klarna's own figures, February 2024.

What the rollout claimed

In February 2024 Klarna announced that its AI assistant had handled 2.3 million conversations in its first month — two-thirds of all customer service chats, across 23 markets and more than 35 languages. It put the volume at the work of 700 full-time agents. Resolution time fell from 11 minutes to under two. Customer satisfaction, or CSAT, was reported on par with human agents, and repeat inquiries fell 25%. Klarna projected a \$40m profit improvement for 2024.

What happened next

By May 2025 the chief executive was revising the account. Sebastian Siemiatkowski, speaking to Bloomberg: *“As cost unfortunately seems to have been a too predominant evaluation factor when organizing this, what you end up having is lower quality.”* He added that *“really investing in the quality of the human support is the way of the future for us.”* Klarna reopened hiring for human support. Over the same period total company headcount fell about 40%, from roughly 5,000 to roughly 3,000 — the product of a hiring freeze and attrition, which is the detail most retellings of this story drop. In September 2025 the company redeployed staff into customer service from marketing, engineering, legal, operations and analytics; those people had first been told their roles were “no longer a priority” and placed in a talent pool.

Every metric Klarna published measured deflection and speed: share of chats handled, full-time-agent equivalents, minutes to resolution. None of them measured whether the resolution was any good. The chief executive's own account is that cost became the predominant evaluation factor and quality fell out of the bottom. What a rollout measures is what it optimises, and quality was not in the published metric set.

The line back to the Lead

Article 50 makes telling a person they are talking to an AI a build requirement. Klarna is the case study in what happens when nobody measures whether the AI was any good.

The takeaway: an AI support rollout that publishes only deflection and speed metrics has already decided what it optimises for, and the correction arrives eighteen months later at the cost of rehiring.

ON THE RADAR

Coming Up

OCT
31



ECB AI cyber-threat action plans due — the deadline for the 110 banks the ECB directly supervises to submit theirs. The ECB's own IT Risk Questionnaire is postponed from September 2026 to February 2027.

DEC
02



EU AI Act Article 50 transitional backstop — generative systems already on the market must meet the machine-readable marking and detection obligation.

AWS
OPEN



ME-CENTRAL-1 / ME-SOUTH-1 — both items remain open on the AWS Health Dashboard, with AWS's last public update dated 30 April 2026. The dashboard could not be re-verified this week, so nothing beyond that is claimed here.

JUL
OPEN



27 July UK multi-bank degradation — still open. No shared root cause has been published, and no Pay.UK, FCA or Bank of England statement landed in the window. The simultaneity is documented; the causation is not. Faster Payments the rail did not fail.

RECOMMENDED READING

This Week's Links

AI mania is eviscerating global decision-making

ludic.mataroa.blog

Incident report: unsanctioned agent behaviour during cyber testing

aisi.gov.uk

A snapshot of entry-level hiring in the UK

gov.uk

Guidelines on transparency obligations for providers and deployers

digital-strategy.ec.europa.eu

VOL. II

Engineering Digest

31 July – 7 August 2026 · Weekly Edition · Issue #21

Next edition: 8–14 August 2026

The Leadership Read returns next week.

Archive

#17 · Weekly

#18 · Deep Dive

#19 · Weekly

#20 · Weekly

#21 · Weekly