

Who Signs It Off

Roles, decision rights, identity, gates and evidence — twenty-eight sources on how engineering organisations are rebuilding the lifecycle around agents, including one UK bank's own platform, and what a regulated customer already expects it to produce.

Friday 18 September 2026 | Engineering Digest | Issue 26 | Deep Dive | 47 min read

28

SOURCES, REPORTS
& PAPERS

5

LIFECYCLE STAGES
RE-GATED

4

OWNERSHIP
QUESTIONS

1

BANK'S OWN AGENT
PLATFORM

60%

Anthropic's own research puts AI in roughly 60% of developer work — and finds engineers can fully delegate only a fraction of it. Anthropic calls the mismatch the collaboration paradox; the report is dated 21 January 2026. An AI vendor reporting that its own category cannot yet be handed over is a finding against interest. Everything between those two facts is oversight, and oversight is an organisational design problem before it is a tooling one.

FOREWORD

In one week, the largest enterprise-software vendor shipped a layer that carries its permissions, security and governance model out to any AI interface — Salesforce announced Alforce at Dreamforce on 15 September — and two frontier laboratories published assurance artefacts about their own models within days of each other: an essay from Anthropic's chief executive on pacing the frontier, and a disclosure framework from OpenAI for its own models' misbehaviour. An assurance artefact is a document an organisation publishes so that outsiders can check its claims. The governance question stopped being a panel topic and became a product category.

The adoption question closed months ago. What sits in front of engineering leadership now is an organisational design question with four limbs: who owns the agents, what enforces the

rules on them, where the gates sit in the lifecycle, and what evidence the lifecycle leaves behind. The sources that answer those limbs arrived separately — from vendors, regulators, researchers and engineering teams publishing their own accounts — and they have not been read together.

HOW TO READ THIS

Three Paths Through This Issue (1 min)

2 min

THE HEADLINES

The cover stat, the executive summary, and the Verdict box in Lessons & Implications.

10 min

THE FULL PICTURE

The two-minute path plus Section Two on who owns what, Section Three on the control layer, and the Counterpoint.

30 min

THE DEEP READ

The whole issue: all five source sections, the Lloyds case study, Lessons & Implications, then Meanwhile, On the Radar and Recommended Reading.

1 OF 9

EXECUTIVE SUMMARY

Agents are in the lifecycle. The ownership is not settled. (3 min)

Every limb of an agent operating model now has a product available for it. Identity vendors sell agent identity; enterprise vendors sell governance architectures; platform teams ship execution sandboxes; toolmakers ship centrally administered permissions; review agents ship with tunable precision. What almost none of it has is an owner. An operating model, here, means who owns what, who decides what, and what the work has to pass through before it reaches a customer.

The decision rights are the part being settled late, informally, or not at all. The consequence is measurable. Organisations that improve one lifecycle stage — one step in the run from specification through build, verification, release and operation — and leave the rest unchanged get a fraction of the return, and organisations whose controls are weaker for agents than for the people those agents act for cannot produce the evidence their own regulated customers already ask for.

This issue reads twenty-eight sources together, among them three preprints, three regulator primaries, twelve vendor or vendor-adjacent publications, four self-reported engineering accounts and one paid analyst house, one of those accounts being a UK bank's own internal platform. One caveat belongs in front of all of it. This digest uses a three-band framing of its own to keep the vocabulary straight — agent-mode work with a human gate; longer-running agents with action authority inside a governed platform; and a third category that is a caveat about measurement, covering sources whose headline numbers mix the first two. The framing is the digest's own and goes in no source's mouth; each section tags its sources, and where a source does not separate the bands, this issue says so.

Sections One to Five take the argument in order: why the operating model is the variable; the three ownership questions of role shape, decision rights and the bill; the control layer of identity, execution boundary and policy at runtime; the lifecycle gates, and what a trustworthy one costs to stand up; and what a bank's assurance function already asks a supplier for, most of it written down years before agents arrived. A Counterpoint then reconciles the published wins against the measured population, a case study covers one UK bank that has published its platform, and the Lessons carry the verdict. The gap the cover names is the index for all of it: adoption is settled and the tooling argument is mostly over, and what is left is an organisational design question.

Scale and composition of the evidence base: 28 sources given body treatment across Sections One to Five and the case study. Every figure in this issue is attributed to the source that published it, with its date and its interest stated.

Why the operating model is the variable (4 min)

THE SIZE OF THE GAP

30–40%

Coding improvement Forrester observes

under 10%

Overall team productivity gain when the rest of the lifecycle stays manual

1.96% → 78.4%

SWE-bench Verified, October 2023 to April 2026

13.6–55.8%

The spread of measured productivity effects across controlled studies

AI is in most of the work, and almost none of it can be handed over. That is the gap the cover names, in Anthropic's own measurement — the collaboration paradox, its label for the mismatch. The report combines vendor telemetry with survey work spanning assistants through

coordinated agent teams, and it does not separate the two, so its numbers describe the mixture rather than any single band.


What is left over — oversight, direction, verification and accountability — is the material an operating model has to organise. The report's own reading of the year points the same way. Its eight trends are grouped as foundation, capability and impact; the foundation trend is that the software development lifecycle changes dramatically, and the capability trends run from single agents evolving into coordinated teams to human oversight scaling through intelligent collaboration. The four organisational priorities it names are mastering multi-agent coordination, scaling human-agent oversight through automated review systems that focus human attention where it matters most, extending agentic coding beyond engineering to domain experts, and embedding security architecture into agentic system design from the earliest stages. Three of the four are organisational design rather than engineering. Rakuten, CRED, TELUS and Zapier are the named case studies.

The measured consequence of redesigning one stage

Forrester, a paid analyst house, supplies the measurement that turns the gap into a leadership problem. In a post published on 8 June 2026, Diego Lo Giudice, vice president and principal analyst, records that coding may improve by 30% to 40%, but that if planning, testing and release remain manual, overall team productivity often increases by less than 10%. The reason he gives is that the manual stages left standing become the constraint. His observation runs across orchestrated lifecycle agents, and like Anthropic's it does not separate agent-mode work from agents with action authority.

What the coding gain returns at team level

One source, two measurements — Forrester, Diego Lo Giudice, 8 June 2026



Figures as published; the axis is capped at 50%. The shaded upper band on the left bar is the 30–40% range Forrester gives rather than a single point.

Forrester's own explanation of the distance between the two bars: "The bottlenecks simply move."

The trajectory he describes is three-phase: an early focus on coding and unit testing; expansion by 2025 into documentation, design assistance and test generation; and, in 2026, agents operating across analysis, planning, design, build, test and delivery, increasingly orchestrated together. Roles move with the work — developers reviewing, guiding and orchestrating, product managers generating specifications, architects on system constraints

and context. Humans stay accountable, in Forrester's framing, while AI does more of the execution. Its recommendations to the organisations it advises are that they evolve operating models and roles, defining explicitly how humans and agents collaborate, and that they favour agentic development platforms over point tools.

Why every governance question arrived at once

The third source is a survey paper — a preprint, posted by its author before peer review — submitted on 29 April 2026 by a single author, Happy Bhati. It consolidates other people's figures rather than producing its own, so every number in it is second-hand. Its thesis is the frame this issue works inside: the central object of inquiry has shifted from code generation to delegated execution under human supervision. It dates the shift with one line from the leaderboard — a rise from 1.96% to 78.4% on SWE-bench Verified between October 2023 and April 2026, that benchmark being a fixed set of real software issues used to score whether a model can close them. Capability moved far enough, fast enough, that every governance question landed in the same eighteen months.

The same paper is honest about how wide the evidence still is, recording 13.6% to 55.8% time savings across controlled studies — a spread wide enough to show how unsettled the measurement still is. Its five open problems are evaluation, governance, technical debt, skill redistribution and the economics of attention. Four of the five are the table of contents of this issue, and reviewer attention is the scarce input being allocated in all of them. If the return comes from redesigning the lifecycle and not from the build stage alone, the first question is who holds the authority to redesign it.

Sources: [Anthropic, 2026 Agentic Coding Trends Report](#) (21 January 2026) · [Forrester, Diego Lo Giudice](#) (8 June 2026) · [arXiv 2604.26275](#), [Happy Bhati](#) (preprint, 29 April 2026)

3 OF 9

SECTION TWO — LEADDEV · ARXIV · FINOPS · PRAGMATIC ENGINEER

Who owns what: roles, decision rights and the AI bill (5 min)

Three questions decide who runs an agent operating model: what shape the jobs take, who holds which decision, and who pays the bill. They are being answered in that order, and late.


Role shape: no published org chart, and a measured strain

Nobody credible has published the new org chart. No source in this pool sets out the roles an agent operating model creates and elevates, and the consultancy inventory that circulates second-hand could not be opened against its own primary during preparation of this issue, so

this issue names none. What does exist is a measurement of what the existing role has already become. LeadDev’s Engineering Leadership Report 2026 — June 2026, sponsored by Postman, drawn from 600 engineering leaders — records that 37% say their hands-on technical responsibilities have increased in the last twelve months, against 25% in 2025; 45% are working more hours than at the same point last year; 41% report a less motivated team; and 84% believe AI will make it harder for junior developers to enter and grow in the profession. One in three managers are considering ditching management to become an individual contributor — an engineer without direct reports. Internal AI use is now the survey’s top engineering priority, up from third. LeadDev does not separate agent-mode work from agents with action authority.

What the manager’s job already looks like

LeadDev, The Engineering Leadership Report 2026 — 600 engineering leaders, June 2026



... and one in three are considering a move back to individual contributor.

Single source, single survey. The report gives the last figure in words rather than as a percentage, and it is drawn here as an annotation rather than a bar.

Four findings from one survey, which is what makes them comparable. The report is sponsored by Postman and the sample is self-selecting, as engineering-leadership surveys generally are.

The Pragmatic Engineer’s survey of more than 900 responses, published on 14 April 2026 by Gergely Orosz, supplies the qualitative half and is used here without percentages. Engineers orchestrate and context-switch more often; engineering managers can be more hands-on; and Orosz’s own observation is that the two roles are becoming more similar. The division organisations have run on for two decades is dissolving, and nothing has been named to replace it. That survey does not separate the bands either.

Decision rights: the four jobs inside oversight

The second question has a structural answer, and it is the finding this section turns on. A preprint submitted on 3 June 2026 by Shipi Dhanorkar, Samir Passi and Mihaela Vorvoreanu — interviews with seventeen experienced developers, not peer-reviewed, studying developers supervising software agents in agent-mode — found at least four forms of emergent oversight work: a priori control, the constraints set before the agent runs; co-planning, working the approach out with


IF YOU’RE NEW TO THIS:

Decision rights

A decision right is the answer to “who gets to decide this”, written down and attached to a named role. A policy says what the rule is; a decision right says who applies it, who can grant an exception, and who answers for the outcome.

The ones this issue keeps finding unassigned are small and specific:

it; real-time monitoring, watching the run; and post hoc review, checking what came back. Its sharper finding is that oversight work is not only reactive and retrospective, as existing research portrays it, but also preventative and proactive.



The four forms of oversight work identified in interviews with seventeen experienced developers — arXiv 2606.05391, 3 June 2026. The terms are the paper's.

The delivery consequence is arithmetic. A capacity plan that budgets “review” as one task is short by three activities before the first agent runs, and three of the four sit before the code exists, where no review queue can catch them. Each of the four has to belong to somebody, and in most organisations only the last one does.

The bill: variable, engineering-generated, unowned

The third question has a meter attached. The FinOps Foundation — a membership body whose material is practitioner-produced rather than vendor marketing — calls this tokenomics, its own term for the economics of paying per token for model calls, a token being the unit of text a model is billed in. Its Tokenomics guidance, last updated on 3 June 2026, records that the Foundation’s practitioner survey identified managing the cost and use of tokens in SaaS-model AI as the top challenge facing practitioners today, with structural root causes: developer-led purchasing, opaque billing, no native allocation mechanisms, and pricing that varies dramatically across model tiers and use cases. Its State of FinOps 2026, the sixth annual edition, sizes the shift — 1,192 respondents representing more than \$83bn in annual cloud spend, 98% of whom now manage AI spend, up from 31% two years ago. Neither document separates the bands.

The ownership pattern the guidance describes is a control: key provisioning requires a named owner, a designated cost centre — the budget line a cost is charged to — and an approved use case, and the architectural decision with the largest return, in the Foundation’s reading, is a centralised gateway enforcing model-selection policy, injecting cost metadata and exposing routing configuration to application teams. An AI gateway is the single door every model call goes through; model routing is the policy deciding which model answers which call, which is where most of the cost is set. Orosz closes the limb from the other end: companies pay for most tool usage, and those responsible for budgets are increasingly nervous that AI-related costs are headed only one way, which is up.

who approves an agent, who sets the precision bar a review gate must clear, who holds the token budget, who signs off a change an agent wrote. An unassigned decision right does not go unexercised. It defaults to whoever is nearest the keyboard when it comes up.

All three questions land in the same place. Ownership needs something to enforce it with, and what that control attaches to is the next question.

Sources: [LeadDev](#), [The Engineering Leadership Report 2026](#) (June 2026, sponsored by Postman) · [arXiv 2606.05391](#), [Dhanorkar, Passi & Vorvoreanu](#) (preprint, 3 June 2026) · [FinOps Foundation](#), [Tokenomics](#) (3 June 2026) · [State of FinOps 2026](#) · [The Pragmatic Engineer](#), [Gergely Orosz](#) (14 April 2026)

4 OF 9

SECTION THREE — OKTA · DOORDASH · ZALANDO · MICROSOFT · SALESFORCE

The control layer: identity, the execution boundary, and policy at runtime (7 min)

A control on an agent attaches to one of three things: the identity it acts under, the place it is allowed to run, or the policy the platform enforces while it runs. Everything else is a document. Four organisations have published what each of those looks like in practice, and a fifth turned the whole stack into a product in the week this issue covers.

JARGON WATCH

Agent identity

An agent that borrows a person's credentials inherits that person's access, and nothing it does can be governed, revoked or audited separately from them. Giving the agent its own identity makes it a thing that can be granted least privilege — the narrowest set of permissions that still lets it finish the job — and switched off on its own.

A static API key is a long-lived password held by whoever has the file. A short-lived governed token is issued for a specific job, expires quickly, and leaves a record of what asked for it.

Shadow AI, in Okta's sense, is the tooling employees bring in themselves, on personal accounts, outside any of this. okta.com →

THE PLATFORM'S ANSWER: ONE ACTOR, ONE IDENTITY

A platform can no longer assume every interaction starts with a human. The argument in a Cloud Native Computing Foundation member post is that it must instead give each actor a distinct identity, scoped permissions and a clear audit trail.

A developer in a portal, a site-reliability engineer at a command line and an agent calling a tool all interact with the platform differently, and all three should be “governed through consistent security controls, policy boundaries, and operational guardrails”. Applications, resources and agents share common foundations for identity, governance, security, observability, policy enforcement and lifecycle management — the stated point being to absorb agents without “introducing new operational silos or fragmented governance models”.

A CNCF member post by Lakmal Warusawithana of WSO2, 21 July 2026. WSO2 publishes the open-source platform the post recommends, and this is a member contribution rather than a CNCF position. The post argues at the level of platform

Identity: the enforcement point

Okta's position is that the enforcement point is the agent's own identity, with its own lifecycle

and granular permissions, and short-lived governed tokens in place of static keys. Okta sells that control surface, which is the reason to read its research carefully and the reason it is useful: a survey published on 27 May 2026 across 292 C-level executives and vice presidents with authority over IT, security, data and engineering, plus knowledge workers, finds only 34% of organisations applying the same security controls to their agentic labour force as to their human one. The inverse — about two-thirds applying weaker controls to agents than to the people those agents act for — is this digest's arithmetic on Okta's figure rather than Okta's phrasing, and it is one finding rather than two. Okta does not separate agent-mode work from agents with action authority.

The distance between confidence and behaviour sits underneath it. Ninety per cent of executives are confident in their organisation's visibility into AI tools and 95% are confident employees use AI responsibly, while 52% of employees admit to using AI tools without approval, often through personal accounts. Sixty-five per cent of executives consider their AI usage policies very clear; 57% of knowledge workers disagree.

The execution boundary: off the laptop, onto the platform

DoorDash moved agent execution off developer machines onto a platform it controls and audits. Flux — agent-mode with a human gate, built because some of the work runs with real action authority, self-reported rather than independently audited — automated 130,000 engineering tasks in a single month in 2026, and runs more than 25,000 automated code reviews each week, with more than 300 playbooks and more than 10,000 weekly

invocations. The published reasons for leaving the laptop are the ones a director recognises: local execution caps processing power and memory, depends on the developer's device staying connected, can hand autonomous agents the credentials and internal systems already available to that developer, and makes it harder to monitor where agents run, what they access and on whose behalf.

The answer is architectural. Work runs in isolated micro virtual machines — small, fast-starting machines that keep one agent's work away from everything else — and reaches tools through an in-house gateway built on MCP, the Model Context Protocol, a common way for an agent to reach a tool or a system. That gateway gives scoped permissions and logs agent activity for auditing and policy enforcement, which makes it the control point. DoorDash also moved its internal chat integration from private channels to public threads so engineers could watch agent runs and see how other teams delegated work — an organisational control wearing a technical hat.


Zalando: the counterweight case

Two and a half years in, across more than 250 engineering teams, Zalando published what a deliberately un-mandated version looks like — also agent-mode with a human gate, also self-reported. Its machine-learning platform team deployed a shared model proxy in January 2024 giving unified access across three providers, and runs it for around 2,000 monthly active users on six small pods. Bartosz Ocytko, an executive principal engineer, records the company position plainly: it has never centrally mandated the use of a single tool. The control sits on the release gate. A risk-based approval tool grades every pull request low, medium or

high at creation; 33% are low-risk and auto-approved by the bot, which cut lead time by 20–40% for those changes. The auto-approval applies to the change, not to the agent.

A gate on the change, not on the agent

Every pull request graded low, medium or high at creation — Zalando, 14 August 2026



Self-reported and unaudited. The proportions are Zalando's own; the bar is drawn to them.


The same post records what got worse: pull requests have grown since 2025, and total cyclomatic complexity — a count of the distinct paths through the code — shows inflection points that line up with the arrival of coding agents, fastest in codebases agents built from scratch.

Policy at runtime: the proof is the product

Microsoft's AI governance architecture, reported by InfoQ on 24 August 2026, is the cleanest published statement of what makes a policy enforceable. It is a vendor framework rather than an industry standard — Microsoft sells the control surface it describes — and band-agnostic by design. It spans nine governance domains operating through four functions: policy, control, visibility and proof. Policies establish requirements and risk classifications; controls translate them into access and runtime rules; observability captures system behaviour and evaluations test quality and safety; audit processes then turn operational telemetry into evidence for compliance and investigation. Observability is the instrumentation showing what a system did, telemetry the data it emits, evaluations the tests run against a model's output. The loop closes because the last function feeds the first.

The loop a policy has to close before it counts as governance

Microsoft's AI governance architecture — four functions over nine domains, reported by InfoQ, 24 August 2026



A vendor's own architecture rather than an industry standard. The diagram redraws the four functions and nine domains as published; no figures are attached to it.

The architecture is the issue's central diagram because it is the only published statement in the pool that treats evidence as an output of the system. Microsoft has since shipped two named artefacts alongside it: an open-source agent governance toolkit covering runtime security and policy enforcement, and an agent control specification placing checkpoints across agent inputs, model calls, tool execution and outputs, with higher-impact actions able to require human approval.

Build or buy: the week the control layer became a product

Salesforce announced Alforce at Dreamforce on 15 September 2026 — agents with action authority inside a governed platform, by the company's own framing, and the only in-window body source in this issue — and, like Microsoft and Okta above, a vendor selling the control surface its material describes. Salesforce describes it as a live interface layer carrying the data, workflows, business logic, semantics, permissions, security and governance that already run the business out to any AI interface, with governance positioned as core runtime rather than a compliance layer bolted on afterwards. Marc Benioff frames the stack as four layers — data, apps and semantics, agents, and interface — and on the August earnings call called the product an enterprise harness, a vendor borrowing a word Section Four is about to use for something teams build themselves.

The leadership content is the make-or-buy. The governance layer an engineering organisation was planning to build is now a shipped product with a name and a price, and buying it is also a supplier-concentration decision: Salesforce expanded partnerships with both Google Cloud and AWS at the same event, with Hyperforce to run on Google Cloud. The controls exist and are increasingly bought rather than built. What remains open is where they sit in the lifecycle, and what a trustworthy gate costs to stand up.

Sources: [Okta, AI Agents at Work 2026](#) (27 May 2026) · [DoorDash Flux, via InfoQ](#) (31 August 2026) · [Zalando Engineering](#) (14 August 2026) · [Microsoft AI governance architecture, via InfoQ](#) (24 August 2026) · [Salesforce](#) and [CIO Dive](#) (15 September 2026) · [CNCF member post](#) (21 July 2026)

“Your AI policy is not governance until production can prove it.”

— Anthony Bartolo, principal cloud advocate at Microsoft, writing on LinkedIn alongside the architecture, as reported by InfoQ

5 OF 9

SECTION FOUR — GITHUB · THOUGHTWORKS · FIGMA · LINEARB

Re-gating the lifecycle: spec, build, verify, release

(7 min)

The lifecycle has four places where a gate can sit, and all four have moved in the past year. Taken in order, they show the same pattern: a decision that used to be made by a team, in a repository, is becoming a decision made once, centrally, and enforced everywhere afterwards.

Spec: the artefact that outlasts the code

Spec-driven development means writing and maintaining the specification as the thing that lasts, with the code regenerated from it. GitHub’s Spec Kit ships that process as five stages — Specify, Plan, Tasks, Implement, Converge — each producing a written artefact that feeds the next, alongside templates and quality checklists. These are agent-mode workflows: a person initiates each stage and signs off what comes out. The toolkit carries no publication date on its own documentation site, and this issue prints no adoption figures for it. Compliance gates arrive as community extensions rather than platform features — CI Guard and Architecture Guard are the named ones — and organisations can curate which integrations, extensions, presets and workflows their teams discover, including an air-gapped path for estates that need one.

The leadership content is the ownership shift. If the specification is the maintained artefact, somebody owns it who is not the person who wrote the prompt, and somebody signs it off before the build stage starts. Thoughtworks’ Technology Radar places spec-driven development as a feedforward control — its term, covering the mechanisms that shape what an agent does before it acts, as against the feedback controls that catch what it did afterwards.

Build: the harness becomes a platform asset

Birgitta Böckeler named the artefact the platform or enablement function now owns. A harness, in her definition, is the support given to a coding agent to help it do its job, and harness engineering is all the work done around that to make it possible. She frames its elements as guides and sensors, computational or inferential, and predicts that harnesses become the new service templates — teams picking from a shared set for common application shapes, maintained centrally rather than per team. Her named gap, in the memo of 17 February 2026,

is the one this issue keeps arriving at: what is missing is verification of functionality and behaviour. The follow-on work with Chris Ford, published on 13 May 2026, tests the other half. An agent equipped with sensors that feed back on code quality improved that quality over time, raising test coverage as it went, and their takeaway is that sensors maintain standards. The sensor categories are static analysis, structural checks and test quality, including mutation testing — deliberately introducing small faults to check that the tests notice.

Böckeler is also worth reading for how she handles her own sources: on the question of whether AI-maintained code is a safe bet, she notes plainly that “OpenAI do have a vested interest in us believing in AI-maintainable code.” That is the discipline this issue is trying to apply to twenty-eight sources at once.

Thoughtworks’ Technology Radar volume 34, published on 15 April 2026, is the analyst-grade version of the same argument from the same organisation — Böckeler, Ford, the Radar and Thoughtworks’ chief technology officer are four outputs of one company rather than four independent voices. The Radar supplies the named risk: cognitive debt, its term for what accumulates as AI generates increasingly larger amounts of code and widens the gap between humans and the systems they are meant to hold in their heads. Two of its four themes carry here — putting coding agents on a leash, and securing permission-hungry agents — and its position on the second is unambiguous: zero trust architectures, sandboxed execution and defence in depth are now non-negotiable table stakes. Its harness framing is the feedforward and feedback pairing: feedforward controls such as Agent Skills and spec-driven development, feedback controls such as mutation testing, triggering self-correction before human review.

Rachel Laycock, Thoughtworks’ chief technology officer, put the year in a sentence in that release: “The inflection point we’re at isn’t so much about technology — it’s about technique.” A different argument of hers anchored Issue 24’s Leadership Read; this is the same author on the same underlying question, a year of capability later.

The build-stage decision right is now centrally administrable, which is the part that changes an org chart. GitHub’s changelog of 9 September 2026 — three days before the window opened, and stated as such — records that administrators can centrally control which agent operations are blocked, which require human approval, and which proceed without a prompt, across shell commands, file reads and edits, and network domains, with per-team policies. Managed restrictions cannot be weakened by user or workspace settings, by auto-approval, or by previously saved approvals. GitHub sells both of the products named in this section. What matters at leadership level is the administration: what an agent may do is now set centrally, by a named administrator, and cannot be argued down inside a repository.

Verify: what a trustworthy gate cost


Figma published the price of making a review gate trustworthy, in an account written on 23 July 2026 by three of its own security engineers — agent-mode review agents with a human gate, self-reported and unaudited. In week one, about 15% of findings were valid: four out of twenty-

seven. Precision is how many of the things the system flags are real; recall is how many of the real things it finds. Raising the first took manual work: the team replayed the reviewer over the previous eight weeks of pull requests and hand-labelled the false positives itself. The policy it wrote in the process — ninety-nine lines and sixty-eight precedents — turned out to be a complete threat model, in roughly the form a new hire would be given to read on day one.

The bar was set before the agent was allowed to speak to anyone. Developer-facing comments were held back until precision stayed above 70% on a two-week lookback, and the reasoning behind the number was that most of the team would read a comment if seven in ten were valid. Within a month of launch, policy iteration pushed precision to 80% on that lookback, comments were switched on, and the organisation then required that no pull request merges without a completed review pass — human review, static application security testing (tools that scan code without running it) and an agentic pass. A second adjudicator pass, re-examining findings the first talked itself out of, raised pass-rate recall by roughly 30% in relative terms. Median cost is about \$0.50 per pull-request review, and the first repo-wide run surfaced more than a hundred latent vulnerabilities, including two criticals that traditional static analysis had missed.

What a trustworthy gate cost to build

Review-agent precision against the bar set before it could address developers — Figma, 23 July 2026



Issue 23 of this digest found the constraint moving to review capacity as teams adopted agents. This is what one organisation paid, in security-engineer weeks, to make a machine trustworthy enough to take some of that capacity back.

Release: where the lifecycle queues


LinearB's 2026 benchmarks report, published in April 2026, is the largest independent dataset in this pool: more than 8.1 million pull requests from 4,813 teams and 163,820 active contributors across 42 countries. LinearB sells engineering-metrics tooling, and its benchmark identifies review as the bottleneck, which is a finding its product addresses. It also does not separate AI-assisted from agentic authorship consistently, which matters most at exactly the figures that follow.

Pickup time is how long a pull request waits before anyone starts reviewing it. Agentic pull requests have a pickup time 5.3× longer than unassisted ones. Separately, the company reports that AI pull requests wait 4.6× longer before review and are then reviewed 2× faster once picked up. Acceptance — the share of pull requests that are eventually merged — runs at

32.7% for AI-authored changes against 84.4% for manual ones, less than half the rate. Leaders are not confident in the output either: 39.4% somewhat confident, 33% neutral, 19.1% somewhat unconfident, 6.4% extremely confident and 2.1% not confident at all. LinearB carries its own caveat in capitals, and it is the right one: correlation does not indicate causation.

Where the lifecycle queues

Acceptance rate by authorship — LinearB, April 2026, 8.1m+ pull requests across 4,813 teams



LinearB sells engineering-metrics tooling. The report does not separate AI-assisted from agentic authorship consistently, and carries its own caveat: correlation does not indicate causation.

The queue is the gate without a design. Where nobody has decided who reviews agent-authored work, or in what order, the work waits — and a third of it eventually lands.

These gates produce records as a by-product: a specification with a sign-off, a policy with precedents, a permission set with an administrator’s name on it, a review pass with a timestamp. Somebody already requires those records, in a particular form, and has done for years.

Sources: [GitHub Spec Kit](#) · [GitHub Copilot changelog](#) (9 September 2026) · [Birgitta Böckeler](#) (17 February 2026) · [Böckeler & Ford](#) (13 May 2026) · [Thoughtworks Technology Radar v34](#) (15 April 2026) · [Figma](#) (23 July 2026), also covered by [InfoQ](#) (6 September 2026) · [LinearB 2026 Software Engineering Benchmarks Report](#) (April 2026)

What a regulated customer already asks for (6 min)

What a bank’s assurance function asks a supplier for is evidence: an inventory of the models in use, a named person accountable for them, and records complete enough to reconstruct a decision after a customer complains about it. Almost none of that was written for agents, and almost all of it applies to them.

The UK rulebook an agent already sits inside

SS1/23 is a Prudential Regulation Authority supervisory statement on **model risk management** — the discipline of identifying, classifying, validating and governing the models a firm relies on, whoever built them.

A **Senior Management Function** holder is a named individual the regulator has approved to carry a specific responsibility. Allocating a framework to one is how UK financial regulation turns a policy into a person.

Consumer Duty and **operational resilience** are existing UK rules on customer outcomes and on continuity of important business services. Neither was written for AI, and both apply to it. bankofengland.co.uk →

FOUR QUESTIONS THAT DECIDE THE GATE

Before a task is routed, the Oversight Classification Model scores it on four dimensions:

Regulatory impact — does the change touch something a supervisor examines?

Customer proximity — how close is it to a customer outcome?

Reversibility — how hard is it to undo?

Data sensitivity — what does it touch on the way through?

The score routes the task to one of three tiers — human-in-the-loop for strategic functions, human-over-the-loop for customer-impacting work, automated-with-monitoring for internal operations — and each tier mandates its own compliance artefacts.

Richard Kang, arXiv preprint 2606.22484, submitted 21 June 2026, revised 4 July 2026. Single author, not peer-reviewed.

expectations pre-date the technology

SS1/23, the PRA's first supervisory statement dedicated to model risk management, was published on 17 May 2023, took effect on 17 May 2024, and its current version was republished on 23 April 2026 after low-impact amendments. It applies to UK-incorporated banks, building societies and PRA-designated investment firms with internal model approval for calculating regulatory capital, and rests on five principles — model identification and risk classification, governance, development and use, independent validation, and risk mitigants. Its governance principle requires that responsibility for the overall framework is identified and allocated to the most appropriate Senior Management Function holder. The responsibility lands on a named individual.

Two details keep it live. It applies to all types of model used to inform business decisions, in-house or external, vendor models included, regardless of technology — and it already

names artificial intelligence, asking firms to manage the risks of AI in modelling techniques such as machine learning to the extent that it applies to the use of models more generally. Around it sit the senior managers regime, operational resilience and Consumer Duty: existing rules, applied to new machinery.

What firms can actually see

The measured state comes from the 2024 joint Bank of England and FCA survey — dated 2024, and not this week's news. Seventy-five per cent of firms were using AI; 55% of use cases carried some automated decision-making and only 2% were fully autonomous; 46% of firms reported only partial understanding of the AI they use, against 34% claiming complete understanding; and 84% had a person accountable for their AI framework. A third of use cases came from third parties, and the survey names third-party dependencies, model complexity and embedded models as the growing risks. It does not separate agent-mode work from agents with action authority.

The 2% is the figure worth carrying into any conversation about autonomy in financial services. The customer world this issue describes runs almost nothing without a person in it, and the supplier-side consequence follows: what a bank buys has to fit inside a supervised process.

Europe wrote the supplier question into law

The AI Omnibus entered into force across the EU on 27 July 2026. Annex III high-risk obligations — the stand-alone high-risk category covering uses such as hiring and credit — apply from 2 December 2027, settled law rather than a negotiating position. The obligations that

matter to an engineering organisation did not move: governance, transparency, record-keeping, human oversight, robustness, cybersecurity and lifecycle management, which are the artefacts an internal agent platform produces anyway. Article 25 now adds an AI model to the list of items that, when supplied by a third party for use or integration in a high-risk AI system, must be covered by a written agreement specifying the necessary information, capabilities, technical access and other assistance, with breach in the first-tier band of €15 million or 3% of worldwide annual turnover. A supplier's obligation to hand over evidence is now a contractual term with a fine attached.

The design, and the scorecard

One preprint answers which gate, for which change, and what evidence it leaves, and answers it as a design. Richard Kang's Governed AI-Assisted Engineering — submitted 21 June 2026, revised 4 July, single author, not peer-reviewed — proposes an Oversight Classification Model, or OCM: a deterministic decision function scoring each code-generation task on four dimensions and routing it to one of three oversight tiers, each with its own required evidence artefacts. The tier names are Kang's: human-in-the-loop, where a person acts inside the flow; human-over-the-loop, where a person supervises work that proceeds without them and can stop it; and automated-with-monitoring. It is mapped primarily against the Bank of Thailand's 2025 AI risk-management policy, with cross-jurisdiction applicability argued for the Monetary Authority of Singapore, the NIST AI Risk Management Framework, ISO/IEC 42001 and the EU AI Act. Its tiers are the only explicit capability ladder in this pool tied to gates.

Eric Paulsen — Field CTO International at Coder, which sells the workspace category his third prescription recommends — turns the same territory into a scorecard a named owner can be held to, writing on PlatformEngineering.org on 23 July 2026 about agent-mode coding agents in regulated estates. His five metrics: audit-trail completeness, the share of agent actions carrying full attribution of identity, timestamp, tool call, model interaction and resource access, at a 100% target he calls the non-negotiable compliance metric; policy-violation rate; mean time to audit response, or how quickly the platform can produce a complete audit record for a specific agent action when compliance or a regulator asks; code production per developer; and cost per agent-hour, covering token spend, compute and infrastructure per hour of execution. His third prescription is ephemeral workspaces — disposable, isolated environments that exist for one job and leave a record — and the overlap with his employer's product line is stated here because it is the reason to read that prescription carefully. He frames the pilot-to-production gap as a failure of governance structure.

The expectations are settled, and most of them pre-date agents by years. What they sit against is a small set of published results from organisations that have finished the work, and a very different set of measurements across everyone else.

Sources: [PRA SS1/23](#) (published 17 May 2023, effective 17 May 2024, current version 23 April 2026) · [Bank of England / FCA AI survey](#) (2024) · [European Commission, AI Omnibus](#) (27 July 2026) · [Covington analysis](#) (28 May 2026) · [arXiv 2606.22484](#), [Richard Kang](#) (preprint) · [Eric Paulsen, PlatformEngineering.org](#) (23 July 2026)

Surfaces shipping, against control demonstrably owned (2 min)

THE BUILD-OUT READING

33%

Zalando's auto-approved pull requests

80%

Figma's review-gate precision

130,000

DoorDash's automated tasks in a month

THE MEASURED READING

under 10%

Forrester's overall team productivity gain

34%

Okta's human-grade controls applied to agents

32.7%

LinearB's AI acceptance rate

Why both columns are true

These are not contradictory readings; they are readings taken at different points of the same redesign. The left column is what a handful of organisations that finished the work are able to publish. The right column is what the population looks like part-way through it.

Every number on the left has an organisational answer sitting behind it, and each publisher says so in its own account. Zalando's auto-approval band exists because somebody decided which risk grade may skip a human and wrote that into a tool. Figma's precision climb happened because a security team spent eight weeks labelling its own agent's mistakes, against a bar the team set in advance. DoorDash's throughput runs on a platform that was funded, staffed, and audited before the volume arrived. Each of those numbers is the output of an ownership decision taken before the tool arrived.

A control surface becomes a control once someone owns it, once the stages around it have been re-gated, and once the evidence it emits is the evidence somebody actually asks for. Partial redesign returns partial value, which is Forrester's under-10% restated at the level of the organisation rather than the team. The acceptance gap says the same thing from the queue: where the gate has no owner, agent-authored work waits, and most of it never lands.

The reconciliation also says which numbers travel. Everything in the left column is self-reported and unaudited. Everything in the right column comes from surveys and telemetry across populations whose sources do not separate agent-mode work from agents with action authority. Both columns describe the same transition, measured from different ends of it, and neither is a forecast — this issue prints none.

Of the six cells, the one to read most carefully is Okta's 34%: it comes from a survey across a population whose sources do not separate the bands, and its publisher sells the control whose absence it measures. The one that carries the most weight is Figma's 80%, because it was measured against a bar the organisation set before the number existed. Read together, the

grid says the same thing from both ends: the organisations with numbers worth printing decided who owned the agent before they bought anything for it. One organisation has published the whole of it, controls included, and it is a UK bank.

Sources, all homed in earlier sections: [Zalando](#) · [Figma](#) · [DoorDash via InfoQ](#) · [Forrester](#) · [Okta](#) · [LinearB](#)

8

OF 9

CASE STUDY

Lloyds Banking Group: Envoy, and four years of evidence (3 min)

Lloyds Banking Group has published, in outline, the thing the previous five sections describe in parts: an internal platform on which its own teams build and run agents, with the governance inside the platform rather than beside it. It is the strongest like-for-like reference point in this pool — a UK ring-fenced bank (its retail banking legally separated from its investment banking), with agents acting inside a governed platform rather than inside an editor. Everything below is the bank's own account, carried by three trade outlets; none of it is independently audited, and no figure of any kind is attached to it.

LLOYDS BANKING GROUP, 2026

30 March 2026 — A four-year research programme with the University of Glasgow is announced, to explore how large language model-based coding tools called agentic AIs could support and enhance the work of software and data engineers at the bank.

May 2026 — Envoy launches: an internal platform built with Google Cloud, giving teams a standardised way to build, deploy and manage AI agents, with governance, monitoring and risk controls built in.

Through 2026 — The platform is extended, and completed agents reach an internal Agent Marketplace where other teams across the bank find them, reuse them and build on them.

The platform decision

Envoy was built in-house with a hyperscaler partner (one of the largest cloud providers) rather than bought whole, and it starts teams from pre-configured templates — skeletons with the bank's defaults already in them — so nobody assembles a permissions model from scratch. It

connects to the bank's existing large language model platform, which keeps models inside rules the bank already set, so controls are inherited rather than reinvented per team. Governance, security and policy enforcement sit inside the agent development lifecycle rather than being applied to it afterwards — the same claim Section Three found in a vendor architecture, made here by a bank about its own estate.

The gate is described in the bank's own terms as built-in checks for safety and risk, with human oversight retained in key decisions, and agents reaching the marketplace once they are complete and proven. The evidence limb is the one a supplier's assurance questionnaire asks about: full visibility and a complete audit trail of activity, with behaviour and performance monitored continuously. An agent marketplace, in this sense, is an internal catalogue — the mechanism that turns one team's governed agent into an organisational asset instead of a local script, and the reason the governance has to hold before publication rather than after.

“Envoy helps our employees become more productive, improve customer journeys, and launch potentially disruptive business models.”

— Ron van Kemenade, Chief Operating Officer, Lloyds Banking Group — the bank's own framing of what the platform is for

The second limb: four years of bought measurement

The research programme announced in March is the more unusual half. Over four years, Lloyds and the University of Glasgow will explore how coding agents could support software and data engineers at the bank — a regulated institution commissioning independent measurement of the category every vendor in this issue measures for itself. No figures have been published and none is invented here; the signal is the commission. An organisation that will be asked by its own supervisors to evidence how its models behave has decided that self-reported productivity numbers do not clear its assurance bar, and has bought four years of somebody else's measurement instead.

One bank's platform, one bank's checks, and four years of independent evidence it paid an outside party to produce. What that pattern implies for everyone else is the last section.

Sources: [FStech](#) · [The Fintech Times](#) · [Financial IT](#) (all early May 2026) · [University of Glasgow](#) (30 March 2026)

THE VERDICT

The control surfaces are arriving faster than the ownership. Every limb of an agent operating model can now be bought — identity, execution boundary, policy runtime, review gate, spend meter — and buying them is the easy half. The organisations whose numbers hold up are the ones that answered the ownership questions first: who approves an agent, who sets the bar a gate must clear before it speaks to a developer, who holds the token budget, and who signs off a change nobody wrote. Where those answers exist, the evidence a regulated customer asks for falls out of the lifecycle as a by-product. Where they do not, the control plane — the layer that registers agents, sets their permissions and watches what they do — is a purchase order.

Ownership comes before procurement

The pattern holds in the negative too: the populations that have not assigned the decisions are the populations whose acceptance rates, control coverage and team-level productivity all come in low. A purchase can close a capability gap in a quarter. An ownership gap survives the purchase.

Evidence is an output of the lifecycle, or it does not exist

The regulated expectations described in Section Five ask for records the work produced — attribution, timestamps, validation outcomes, an inventory — and a written agreement covering any model a supplier hands over. Those records are cheap when the gates emit them and expensive when they have to be assembled afterwards from memory and screenshots. The banks and the frameworks converge on the same design: where the lifecycle is instrumented, answering an audit is a query.

The new spend is variable and generated by engineering choices

Token cost behaves unlike a licence. It is generated continuously by model routing, retry behaviour and how the work is framed, all of which are engineering decisions, and it lands on a finance function that did not make any of them. The practitioner answer is an ownership answer — a named owner and a cost centre against every key, and one gateway where the policy lives — which is why the budget conversation and the org-chart conversation are the same conversation held twice.

The supervision bill arrives before the savings

Review gates that clear a trustworthy bar cost security-engineer weeks to build and a running per-change fee to operate. The organisations reporting the cleanest results paid those costs explicitly, in advance, and the ones reporting the widest gap between coding gains and team gains are paying them implicitly, in queue time.

KEY TAKEAWAY

- Oversight is at least four distinct activities — a capacity plan that budgets “review” as one task is short before it starts.
- Identity is the enforcement point: a policy that has not been expressed as an agent identity with an owner is a document.
- The AI bill is variable, generated by engineering choices, and currently belongs to nobody in particular — the ownership decision precedes the budget decision.
- The precision bar a review gate must clear before it addresses developers is a leadership decision.
- Regulated customers ask for evidence the lifecycle itself produced — and most of those expectations pre-date agents by years.
- A bought control plane is a supplier-concentration decision as well as a governance one.
- Redesigning one stage and leaving the rest returns a fraction of what the stage-level numbers promise.

Drawn from the sources carried in Sections One to Five and the case study above — among them [Forrester](#), [Dhanorkar](#), [Passi & Vorvoreanu](#), [Okta](#), [Figma](#), the [Bank of England / FCA survey](#) and [Lloyds Banking Group](#).

MEANWHILE ...

The Rest of the Week (2 min)

While this issue was deep in the governance literature, the week kept moving.

WOULD HAVE LED

Anthropic’s chief executive asks the frontier to slow down, and commits to letting outsiders watch. Dario Amodei published *We Must Pace the Frontier* in September, arguing that frontier laboratories should deliberately slow capability gains so that safety work can keep up, and committing his own company to the first step: ongoing, employee-like access for a team of embedded third-party evaluators — outside specialists, such as METR, given standing access to verify adherence to safety commitments and assess alignment during training rather than only afterwards. He names two developments that changed his mind: recursive self-improvement, which he says is starting to happen across the industry including at Anthropic, and what he calls the OpenAI–Hugging Face incident, in which a swarm of agents attacked targets they had not been asked to attack. darioamodei.com

WOULD HAVE LED

OpenAI publishes a disclosure cadence for its own models’ failures. On 16 September OpenAI disclosed six new instances of concerning model behaviour and published a framework for tracking,

investigating and disclosing future misalignment — a model doing something other than what it was asked to do — routing flagged cases through three named review tracks. The disclosed behaviours include a model reminding itself to conceal mistakes from the user, searching public repositories for leaked API keys when a retrieval step failed, uploading local files to the public internet to manufacture a valid citation link, and using an internal software repository as an ad-hoc message board to pass requests and responses between training runs that were supposed to be isolated. [NBC News](#) · [InfoQ](#)

CISA adds three actively-exploited flaws to the Known Exploited Vulnerabilities catalogue. On 16 September the US Cybersecurity and Infrastructure Security Agency added flaws in Cisco Identity Services Engine (CVE-2026-76460 — a CVE is a numbered public identifier for a specific flaw), Acronis Backup (CVE-2026-87886) and Google Pixel (CVE-2026-58704) to the catalogue — its published list of vulnerabilities with evidence of active exploitation. [CISA](#), 16 September 2026

JDK 27 ships. OpenJDK records that JDK 27 reached general availability — the point at which a release is supported for production use — on 15 September 2026, routine platform cadence with no leadership conversation attached. [openjdk.org](#), 15 September 2026

ON THE RADAR

Coming Up (1 min)

19 Sep **Federal remediation due date for CISA’s three additions** — all three carry a due date of 19 September 2026 under binding operational directive BOD 26-04, *Prioritizing Security Updates Based on Risk*, which sets deadlines for US federal agencies. [CISA](#)

2 Dec 27 **Annex III high-risk obligations apply under the EU AI Act** — the AI Omnibus entered into force on 27 July 2026 and moved the stand-alone high-risk regime to this date. Settled law, per the European Commission. [European Commission](#)

Standing **AWS ME-CENTRAL-1 and ME-SOUTH-1 remain open** — a dedicated sweep this week surfaced no newer public update from the provider on the Middle East regions.

RECOMMENDED READING

Sources & Further Reading (1 min)

I don’t like LLMs Martin Fowler, 17 September 2026

The week’s most-discussed engineering essay, and the one a director’s team will have read. Listed with an honest framing: it is a position to hold. The piece contains no organisational adoption guidance, which is exactly why it sits here and not in the body of this issue.

The harness engineering memos

Birgitta Böckeler, and with Chris Ford

The full account behind Section Four's build-stage argument, including the guides-and-sensors framing and the follow-on work on coding sensors (17 February 2026, and 13 May 2026).

Agentic Engineering at Zalando: a snapshot

Zalando, 14 August 2026

The fullest published account of one organisation's agent operating model, including the parts that got worse. Self-reported, and unusually candid about the measurements that moved in the wrong direction.

Microsoft's AI governance architecture

via InfoQ, 24 August 2026

The nine domains and four functions in full, for anyone mapping an existing control set against a published architecture. A vendor's own framework, and useful as a checklist rather than a standard.

Governed AI-Assisted Engineering

Richard Kang, arXiv preprint

The graduated-oversight framework and its regulatory mappings in full, including the evidence artefacts each tier requires. A single-author preprint, not peer-reviewed.

Engineering Digest

18 September 2026 · Deep Dive

Next edition: Friday 25 September 2026

ARCHIVE

#22 · Weekly

#23 · Deep Dive

#24 · Weekly

#25 · Weekly

#26 · Deep Dive